

# 浅谈ASM

**Lunar**


**QQ: 5163721**

**[www.lunar2013.com](http://www.lunar2013.com)**

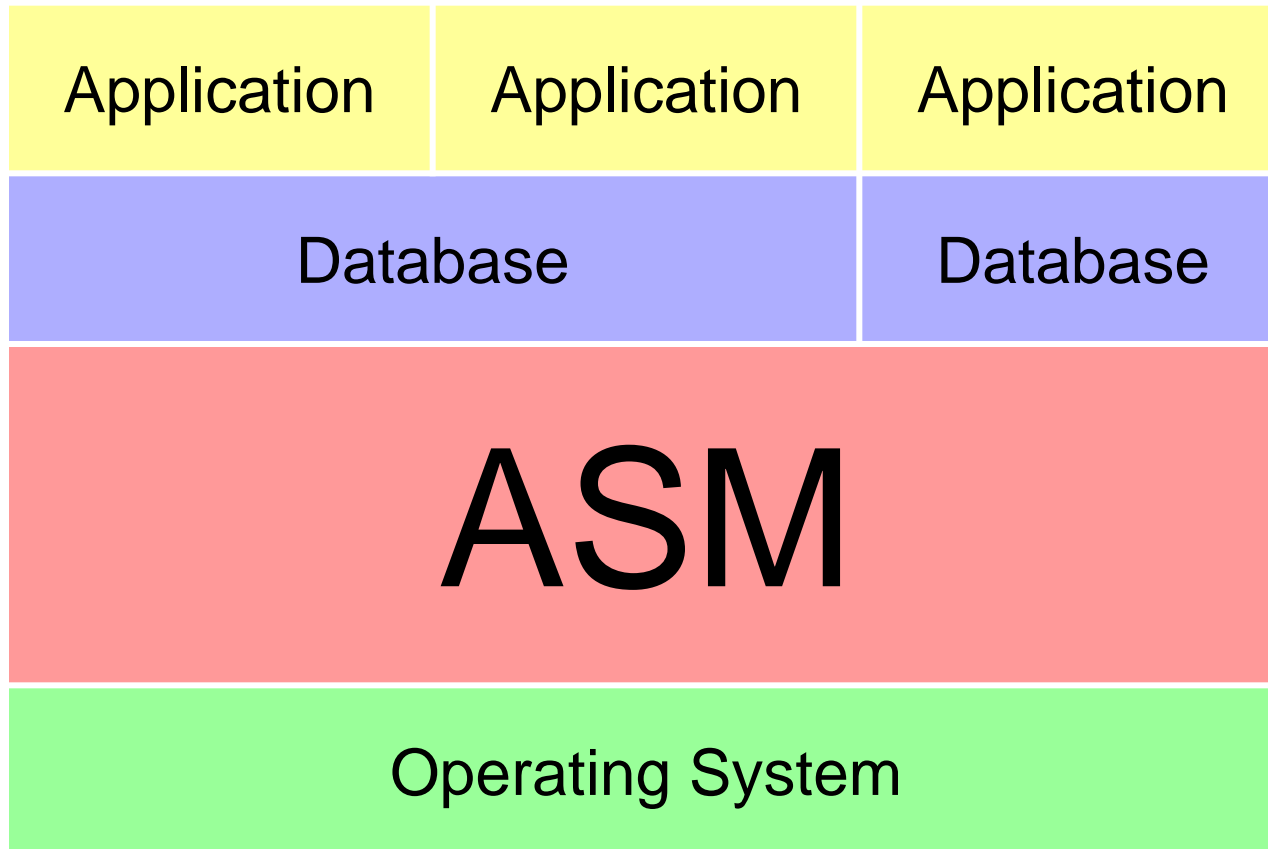
# Agend

- **ASM的历史和简介**
- **ASM体系架构**
- **ASM在不同版本的新特性**
- **ASM存储和分配机制**
- **ASM日常使用注意事项**

# Oracle 9i Software Stack



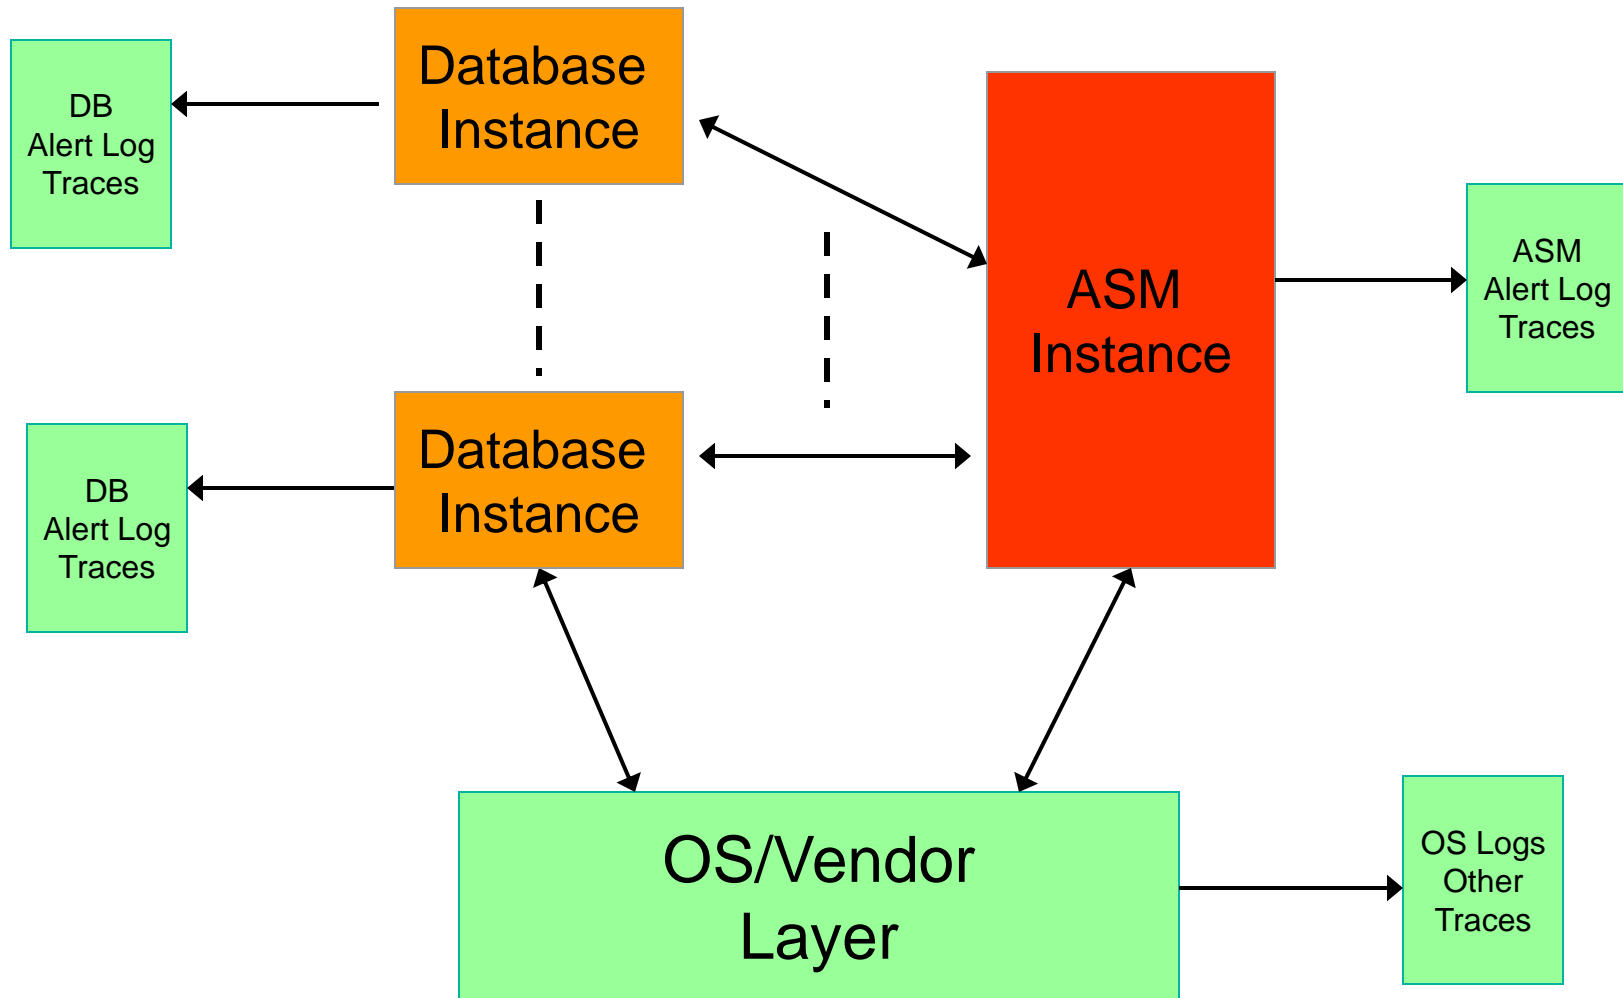
# Oracle 10g with ASM



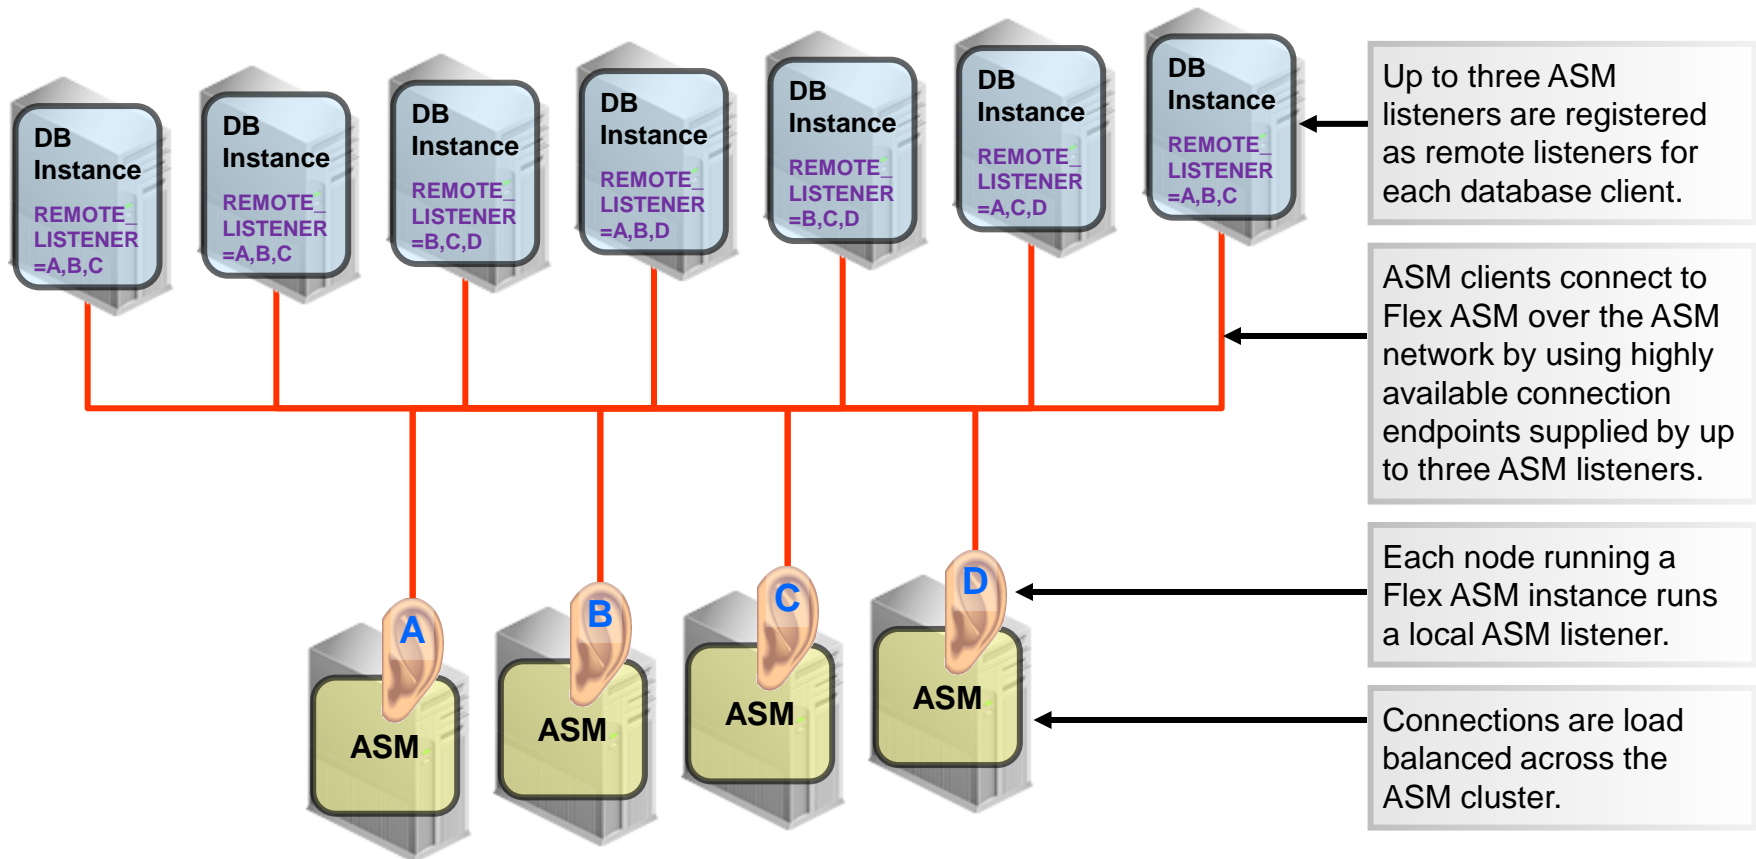
# Agend

- ASM的历史和简介
- **ASM体系架构**
- ASM在不同版本的新特性
- ASM存储存储和分配机制
- ASM日常使用注意事项

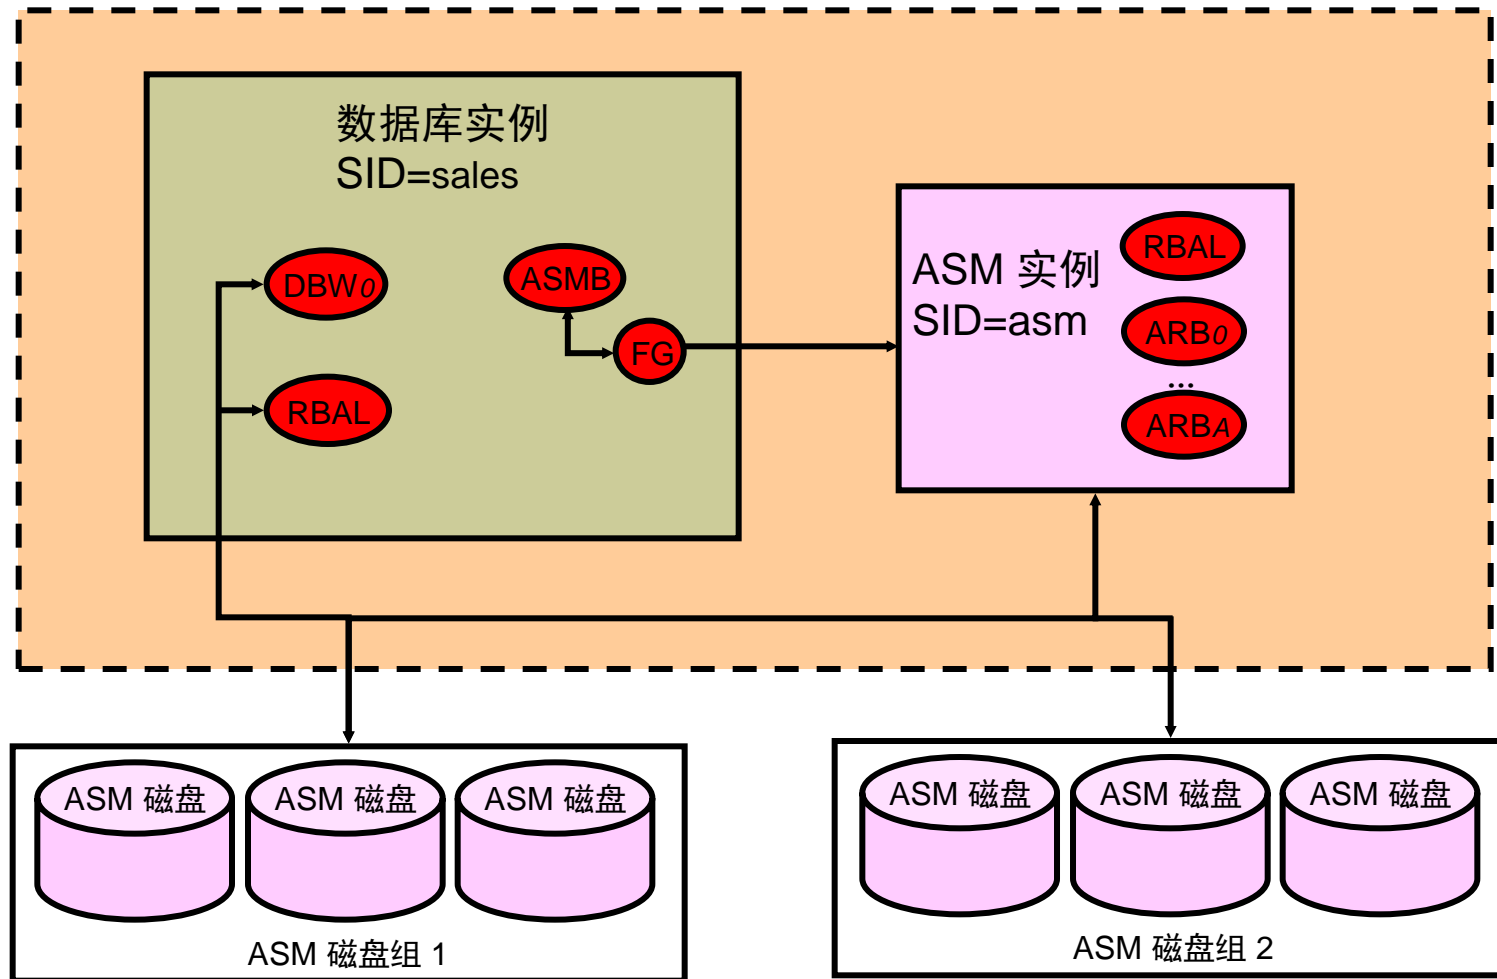
# 12.1 ASM以前的ASM架构概览



## 12.1 ASM以后的Flex ASM架构概览

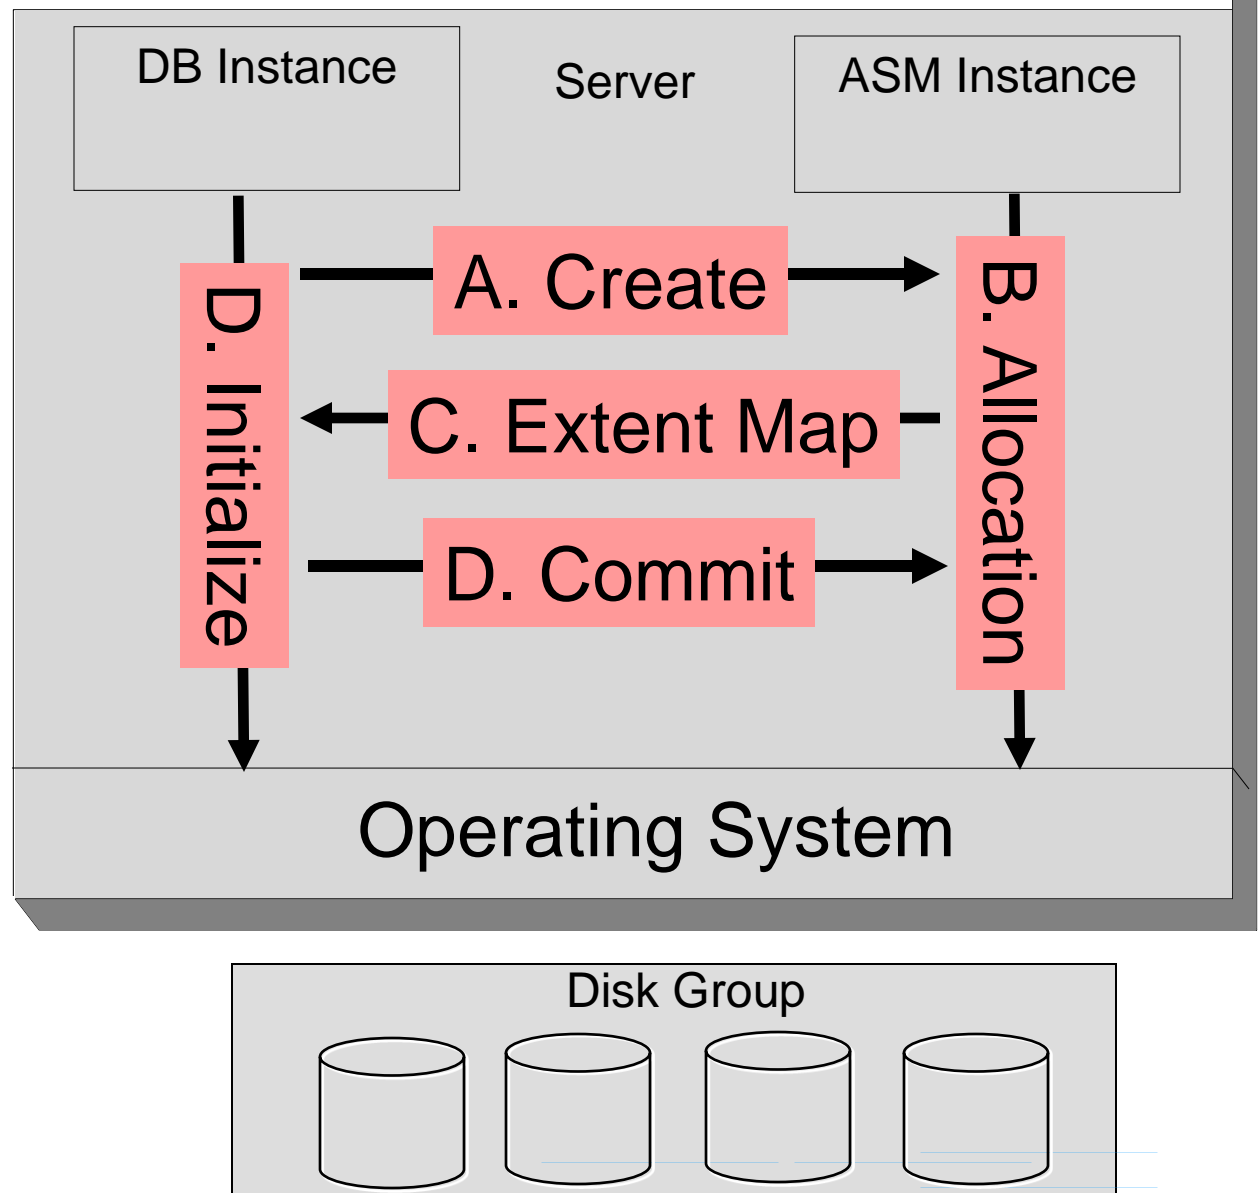


# ASM 常规体系结构






# Database file creation



# Database file open



# ASM的进程介绍--[图解ASMB](#)

## 1. PMON, SMON, LGWR, ARCH, CKPT, DBWR

## 2. ASM processes in database instance:

-----  
ASMB - ASM background process used for client extent management

RBAL - Background process that is used for diskgroup management

o???? - ASM slave processes

## 3. Processes that Run in ASM instance:

-----  
ASMB - Starts up only during DB startup using spfile

RBAL - Background process that is used for diskgroup management

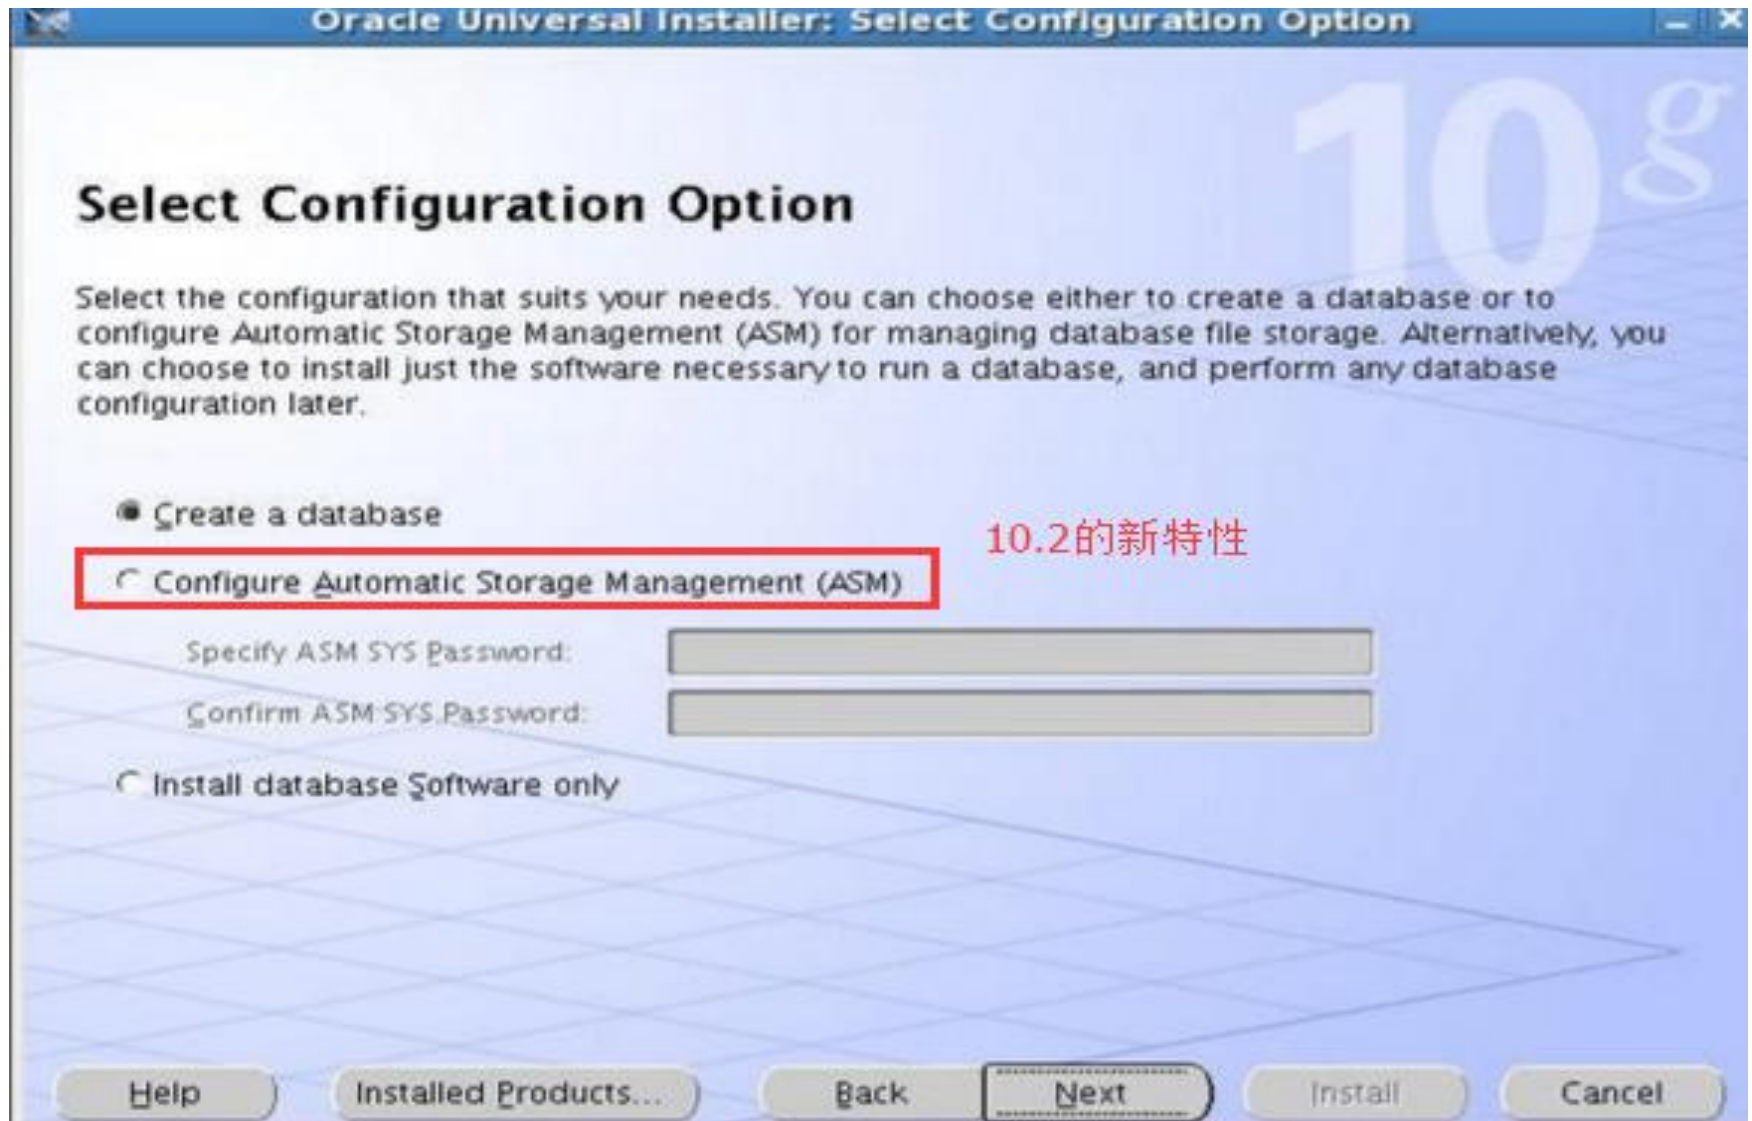
ARBn - ASM Rebalance Process. These processes

- DBW0 - DB writes, same as database DB writer, but deals with ASM cache
- SMON - Recovery process, Same as database SMON, but deals with diskgroup recovery
- CKPT - Checkpoint process, Similar to database CKPT
- PSP0 - Process that Starts other Processes, used to startup other backgrounds
- GMON - Group monitor, used for partner and status table, and node membership
- n00x - GMON Slave processes
- ora\_<procname>\_<dbsid> - Foregrounds servicing clients commands from client <procname> of database <dbsid>
- ora\_ASMB\_<dbsid>: Special ASM foreground that controls database <dbsid> ASMB
- KATE - Konductor of ASM Temporary Errands, used to process disk online
- VKTM - Process to maintain a fast timer, same as database
- PING - Process to measure network latency, same as database
- DIA? - Diag process, same as database
- DIAG - Diag process, same as database
- LGWR - Log writer, similar to database, but deals with diskgroups
- LMON - Lock monitor, Same as database
- LMS? - Lock monitor slaves, same as database
- MMAN - Autotune SGA process, Same as Database.
- b??? - Slave used to process offline disks.
- x??? - Slave used to expell disks after diskgroup reconfiguration
- pz?? - PQ slaves used for global Views

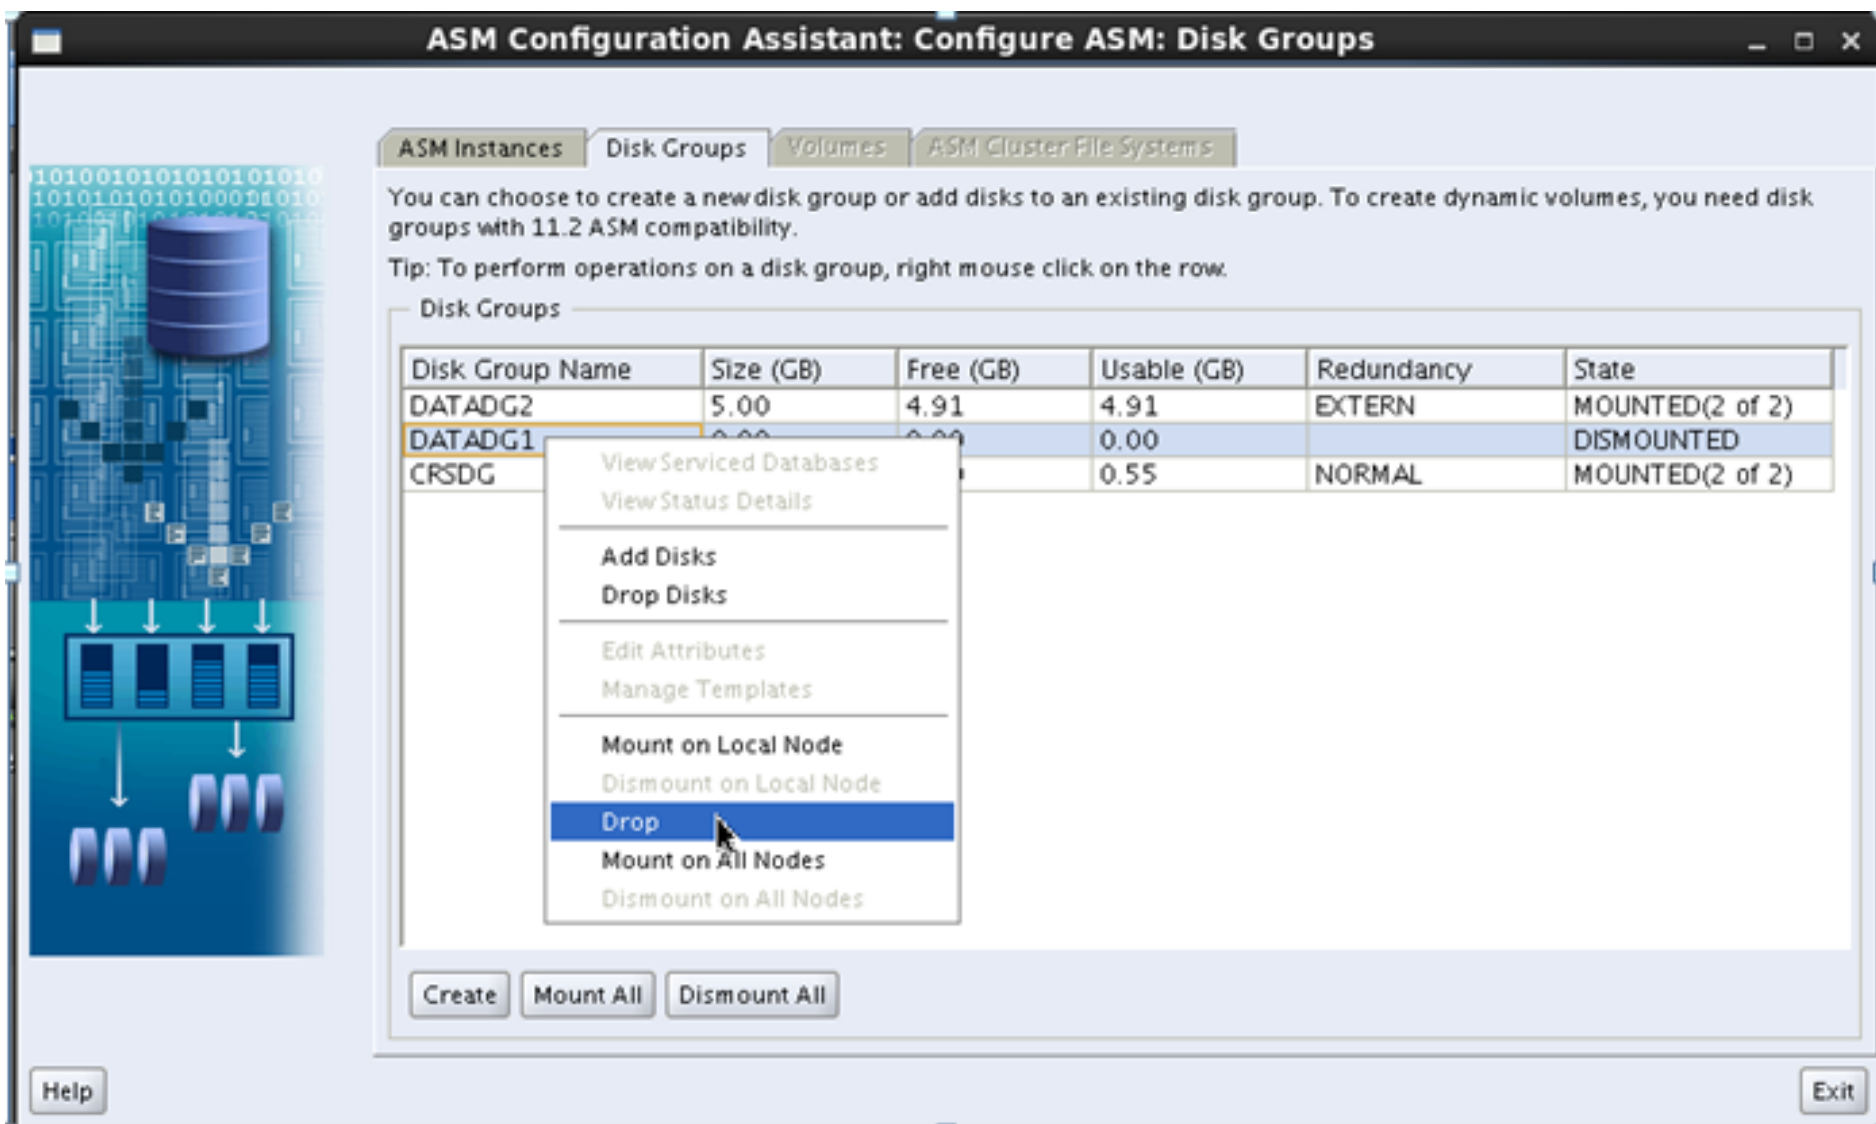
# Agend

- ASM的历史和简介
- ASM体系架构
- **ASM在不同版本的新特性**
- ASM存储和分配机制
- ASM日常使用注意事项

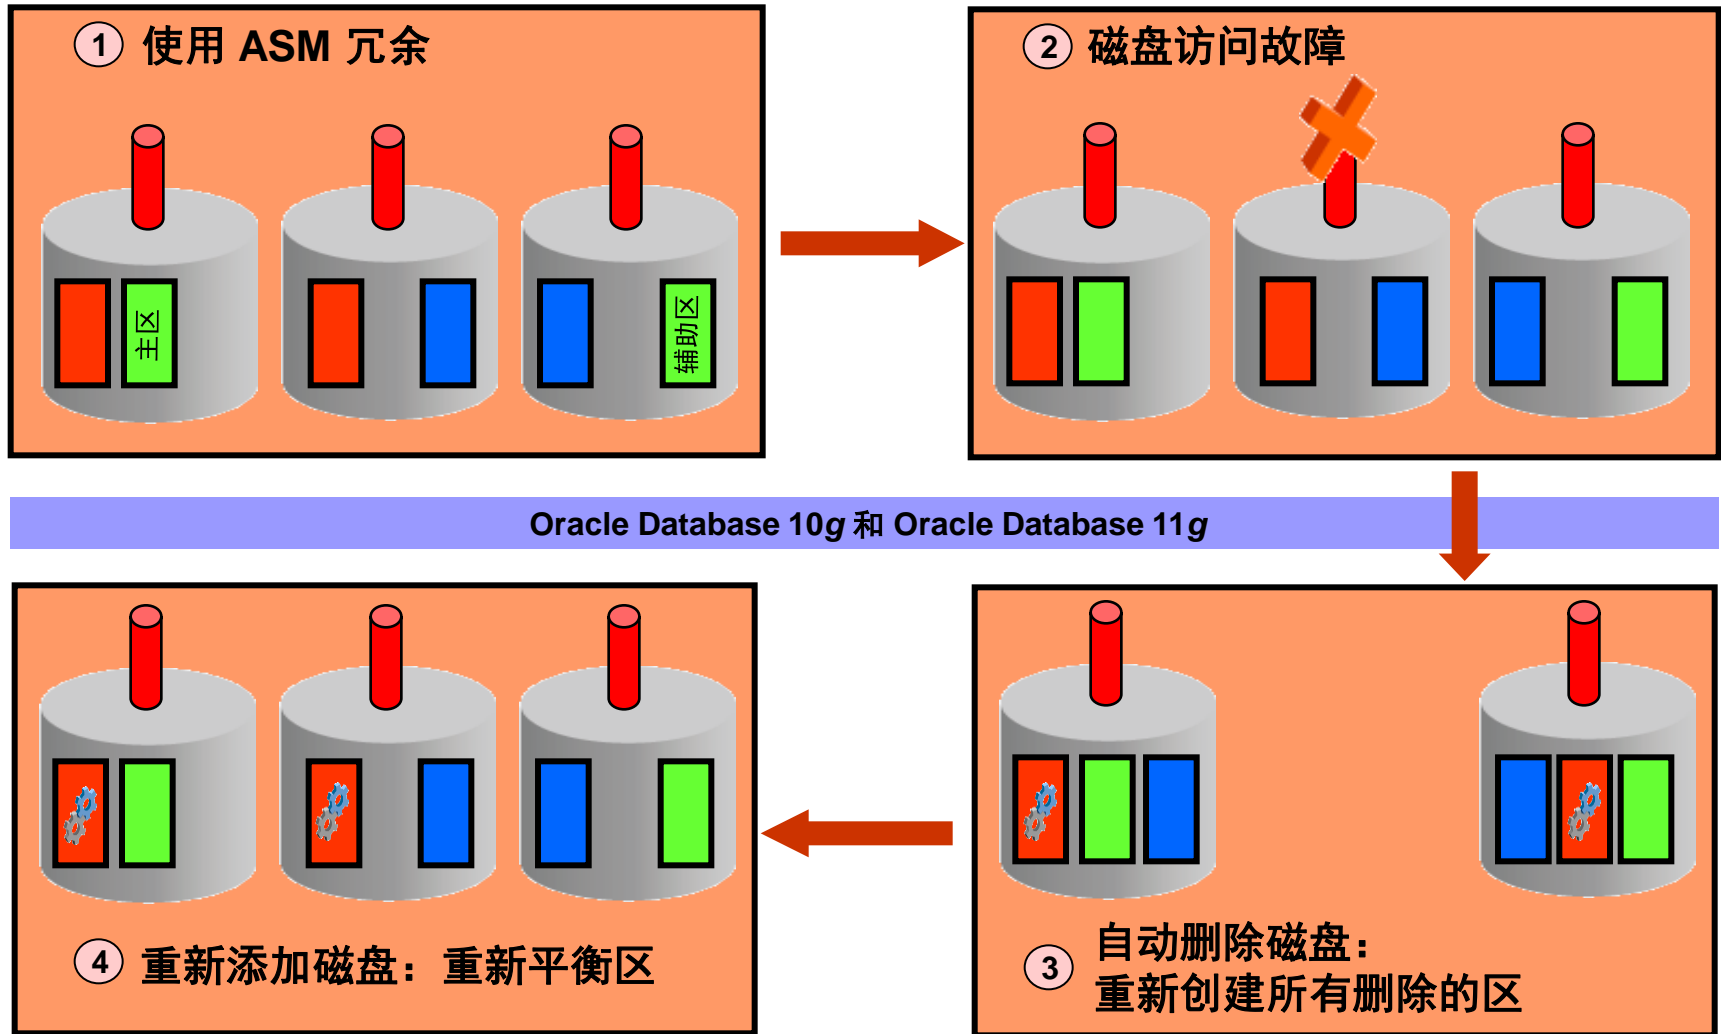
# ASM实例在10.2中可以单独配置和管理



# ASM实例在11.2中使用ASMCA配置和管理

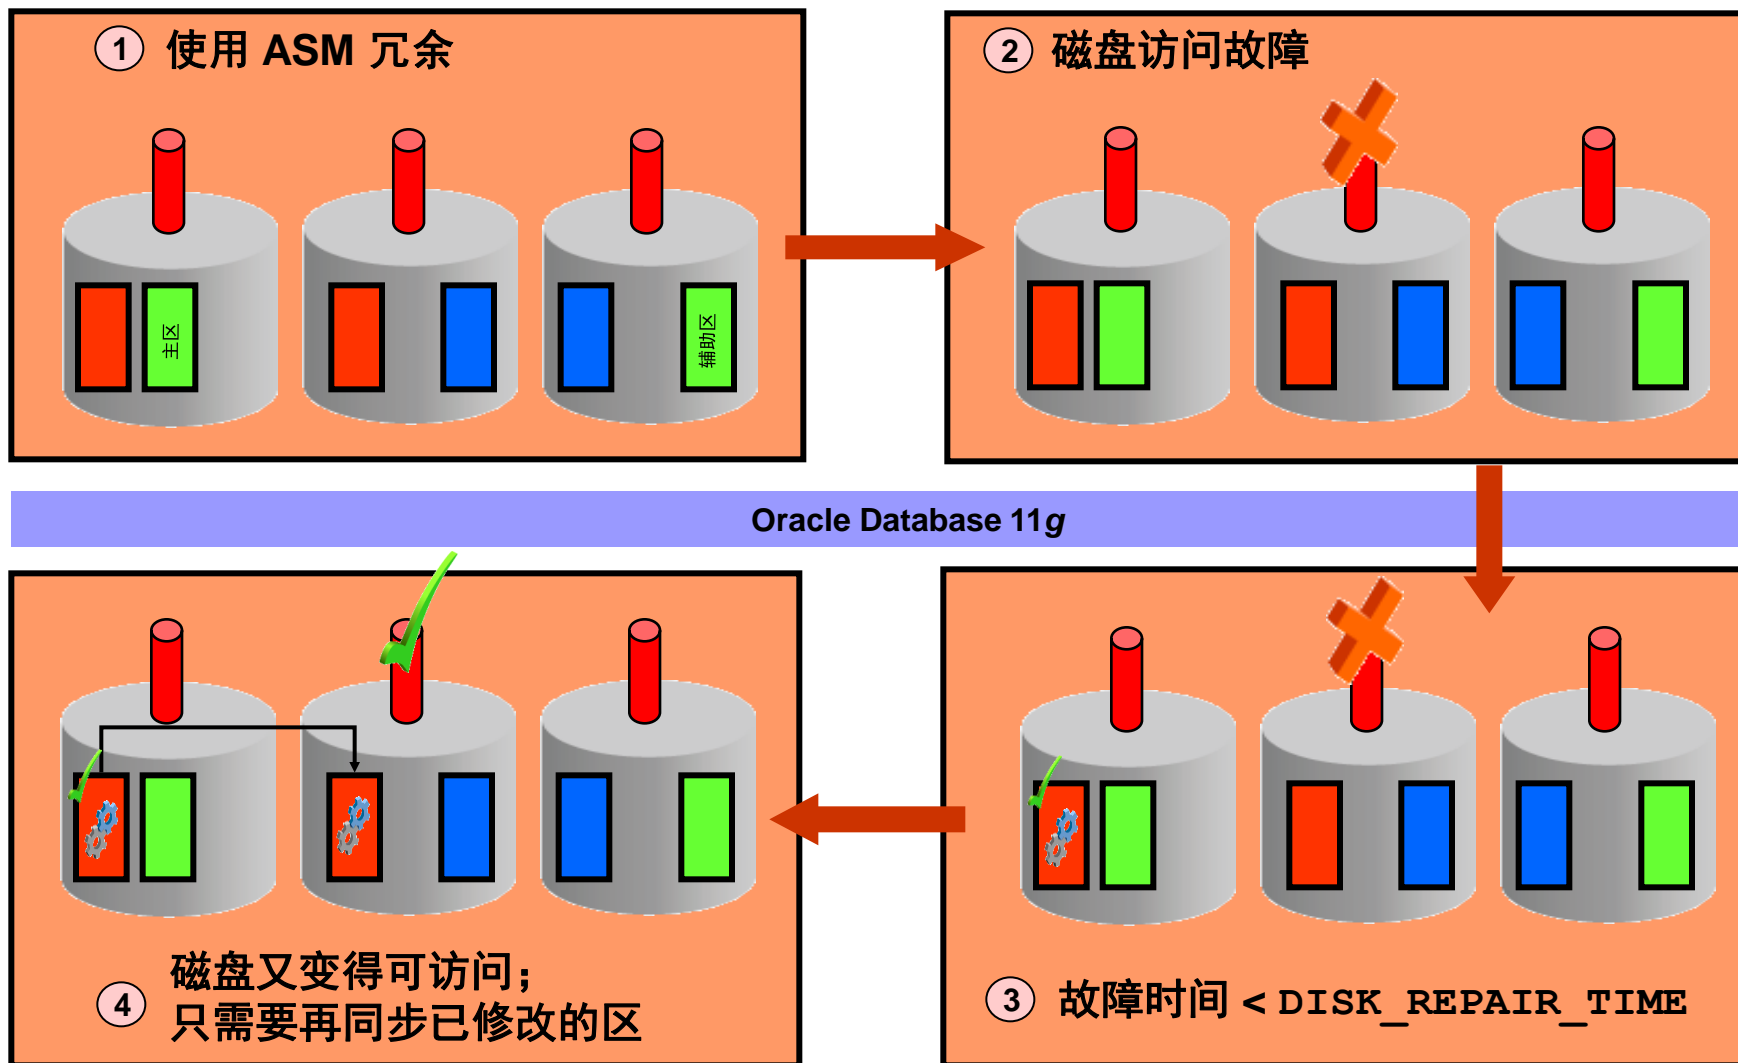


# 无 ASM 快速镜像再同步时



# Oracle 11.1中的增强

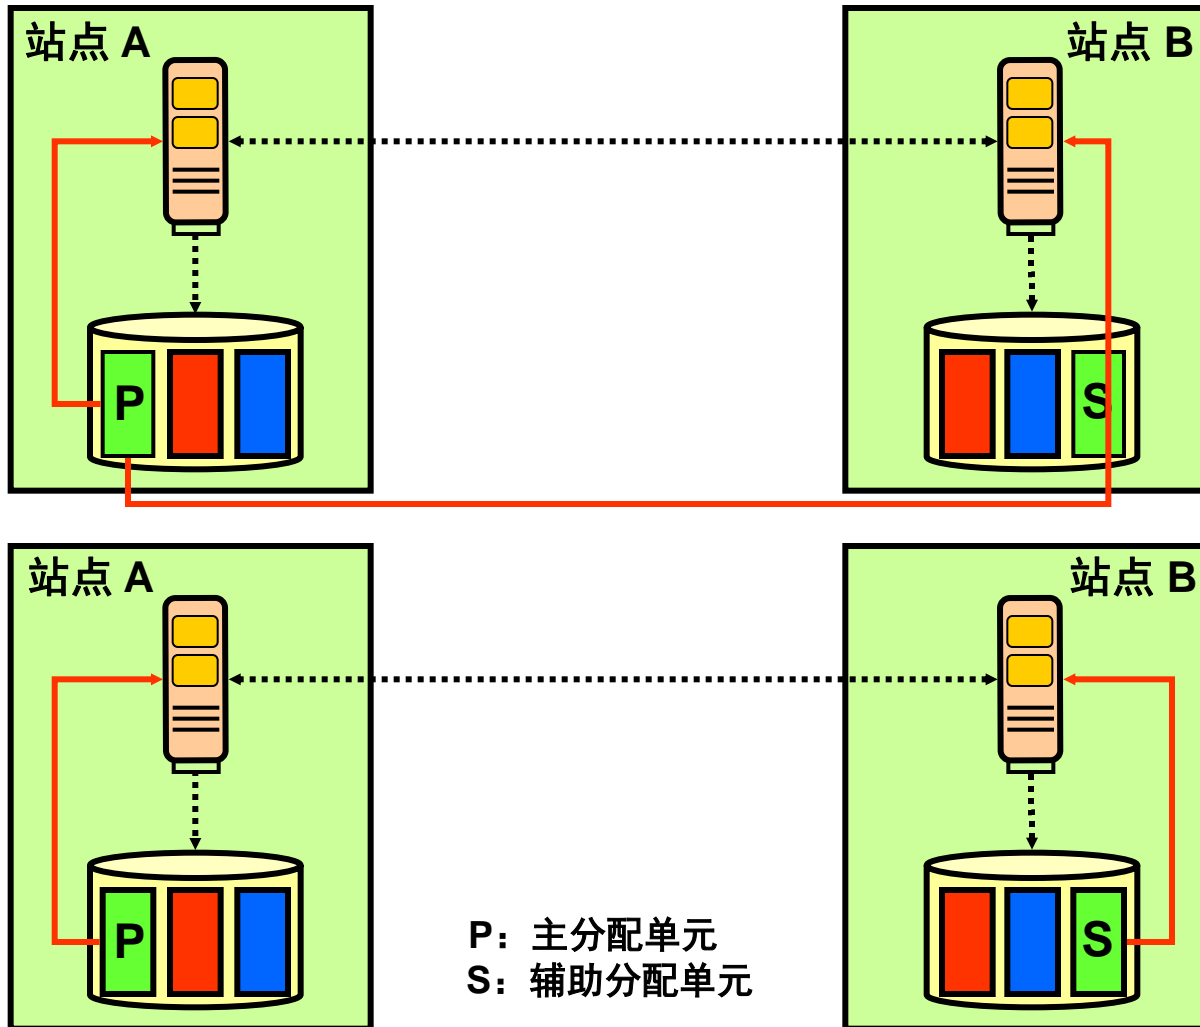
## ASM 快速镜像再同步：概览





# ASM 首选镜像读取：概览

镜像IO由谁完成？



# ASM 首选镜像读取：设置

## 设置

```
ASM_PREFERRED_READ_FAILURE_GROUPS=DATA.SITEA
```

在第一个实例上

```
ASM_PREFERRED_READ_FAILURE_GROUPS=DATA.SITEB
```

在第二个实例上


## 监视

```
SELECT preferred_read FROM v$asm_disk;
```

```
SELECT * FROM v$asm_disk_iostat;
```


# ASM 首选镜像读取：最佳方法

两个站点：正常冗余




仅两个故障组：每个实例一个

两个站点：高冗余



最多四个故障组：每个实例二个

三个站点：高冗余



仅三个故障组：每个实例一个

# Oracle Database 11.1 中的 ASM 可伸缩性

ASM 强制实施下列限制：

- 63 个磁盘组
- 10,000 个 ASM 磁盘
- 每个 ASM 磁盘 4 PB
- 40 EB 的存储空间
- 每个磁盘组 1 百万个文件
- 最大文件大小：
  - 外部冗余：140 PB
  - 正常冗余：42 PB
  - 高冗余：15 PB

• 可变extent：

- 前 20,000 个区 (0-19999) 等于 AU
- 接下来的 20,000 个区 (20000-39999) 等于 4 个 AU
- 40,000 个以上的区等于 16 个 AU

- 使用 SYSASM 角色管理 ASM 实例可以避免 DBA 与存储管理员之间出现重叠。

```
SQL> CONNECT / AS SYSASM
```

```
SQL> CREATE USER username IDENTIFIED by passwd;
```

```
SQL> GRANT SYSASM TO username;
```

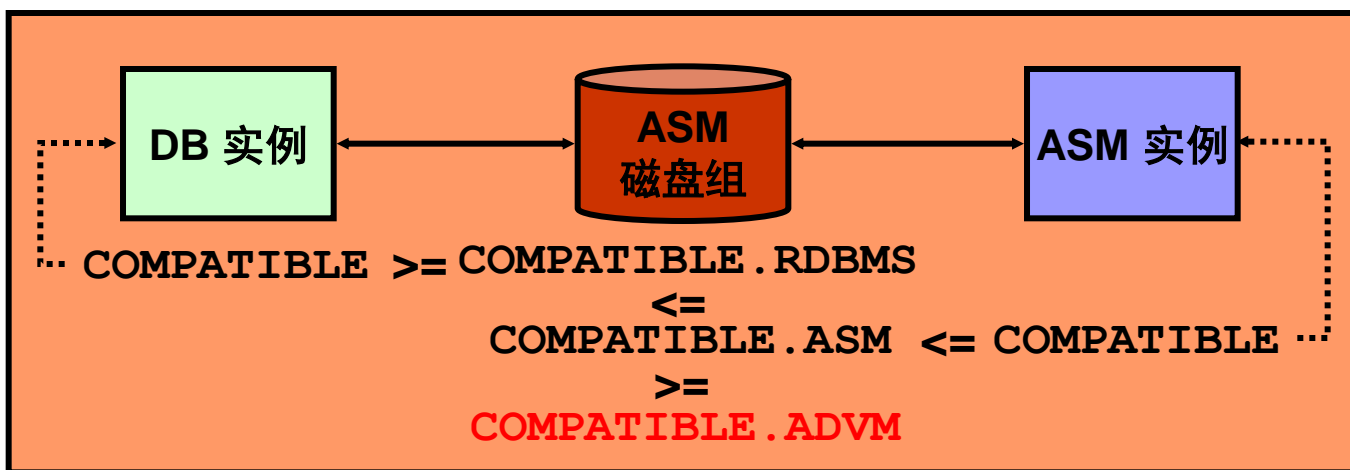
```
SQL> CONNECT username/passwd AS SYSASM;
```

```
SQL> DROP USER username;
```

## **SYSDBA 将被废弃：**

- Oracle Database 11g 版本 1 的行为与 Oracle Database 10g 中的相同。
- 在11.2的版本中，ASM 实例中的 SYSDBA 权限将受到限制。

- 可以分别控制每个磁盘组的兼容性：
  - ASM 兼容性控制 ASM 元数据的磁盘存储结构。
  - RDBMS 兼容性控制使用者客户机的最低级别。
  - **ADVM 兼容性确定是否可以使用 Oracle ASM 动态卷管理器。**
- 对磁盘组兼容性的设置是不可逆的。



# ASM 磁盘组属性

名称	属性	值	说明
au_size	C	1 2 4 8 16 32 64MB	磁盘组中分配单元的大小
compatible.rdbms	AC	有效的数据库版本	数据库与 ASM 之间交换的消息的格式
compatible.asm	AC	有效的 ASM 实例版本	磁盘上 ASM 元数据结构的格式
disk_repair_time	AC	0 M to 2 <sup>32</sup> D	磁盘脱机之后删除此磁盘之前的时间量
template.tname. redundancy	A	UNPROTECT MIRROR HIGH	指定模板的冗余
template.tname. stripe	A	COARSE FINE	指定模板的条带化属性

C: CREATE (创建)  
A: ALTER (变更)

```
CREATE DISKGROUP DATA NORMAL REDUNDANCY
DISK '/dev/raw/raw1', '/dev/raw/raw2'
ATTRIBUTE 'compatible.asm'='11.1';
```

# 适用于快速重新平衡的受限装载磁盘组

- 磁盘组只能装载在单个实例上。
- 任何数据库客户机或其它 ASM 实例都不能获得访问权限。
- 重新平衡可以继续进行，而无锁定开销。

① `ALTER DISKGROUP data DISMOUNT;`

② `ALTER DISKGROUP data MOUNT RESTRICT;`

③ 维护任务：添加/删除磁盘...

④ `ALTER DISKGROUP data DISMOUNT;`

⑤ `ALTER DISKGROUP data MOUNT;`



- 默认情况下，MOUNT 使用 NOFORCE 选项：
  - 所有磁盘必须可用
- 带有 FORCE 选项的 MOUNT：
  - 如果存在仲裁磁盘，则会使不可用的磁盘脱机
  - 如果所有磁盘都可用，则操作会失败

```
ALTER DISKGROUP data MOUNT [FORCE | NOFORCE];
```

Automatic Storage Management: +ASM\_edcdr12p1

**Update Message**  
Disk Group DATA has been dismounted successfully

Home Performance **Disk Groups** Configuration Users

Create Mount All Dismount

Mount Dismount Rebalance Check Delete

Select All | Select None

Select	Name	State	Redundancy	Size (GB)	Used (GB)	Used (%)	Usable Free (GB)	Mem D
<input checked="" type="checkbox"/>	DATA	DISMOUNTED	n/a	0.00	0.00	n/a	0.00	

**TIP** The usable free space specifies the amount of space that can be safely used for data. A value above zero means that redundancy can be properly restored after a disk failure.

**TIP** Mount All and Dismount All operation will only mount and dismount the disk groups specified in Auto Mount Disk Groups parameter.

**Confirmation**

Are you sure you want to mount disk group DATA?

**Hide Advanced Options**

☐ Do Not Force  
All disks belonging to a disk group must be accessible for the mount or dismount operation to succeed.

☒ Force  
Mount the disk group even if some disks belonging to the disk group are not accessible.

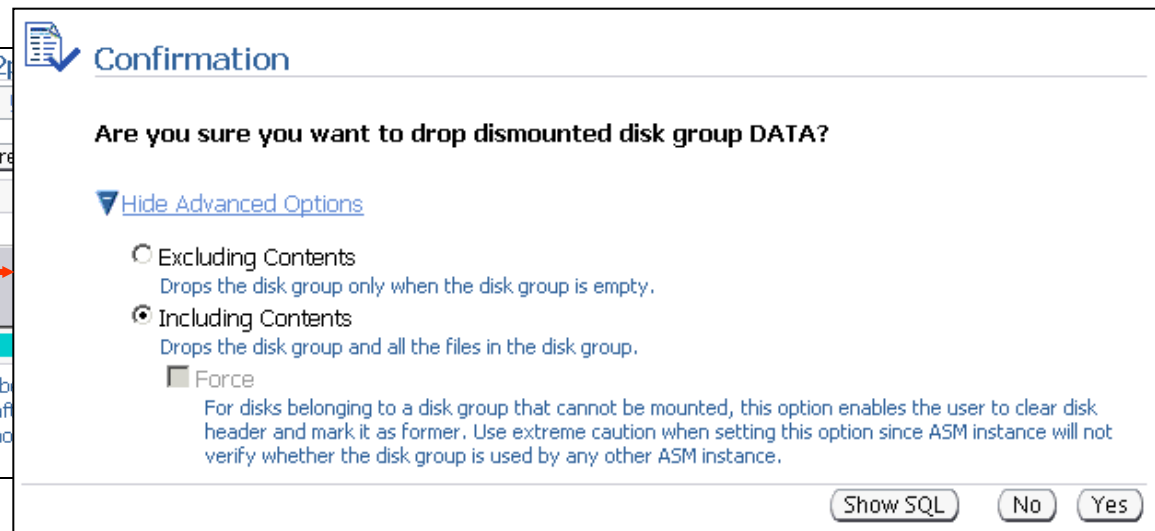
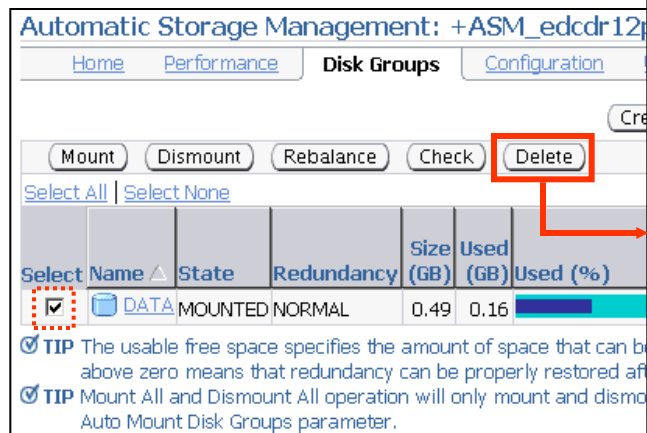
☐ Restricted  
When mounted in this mode, clients cannot access files in the disk group and no other ASM instance in the cluster can mount this disk group.

Show SQL No Yes

# 强制删除磁盘组

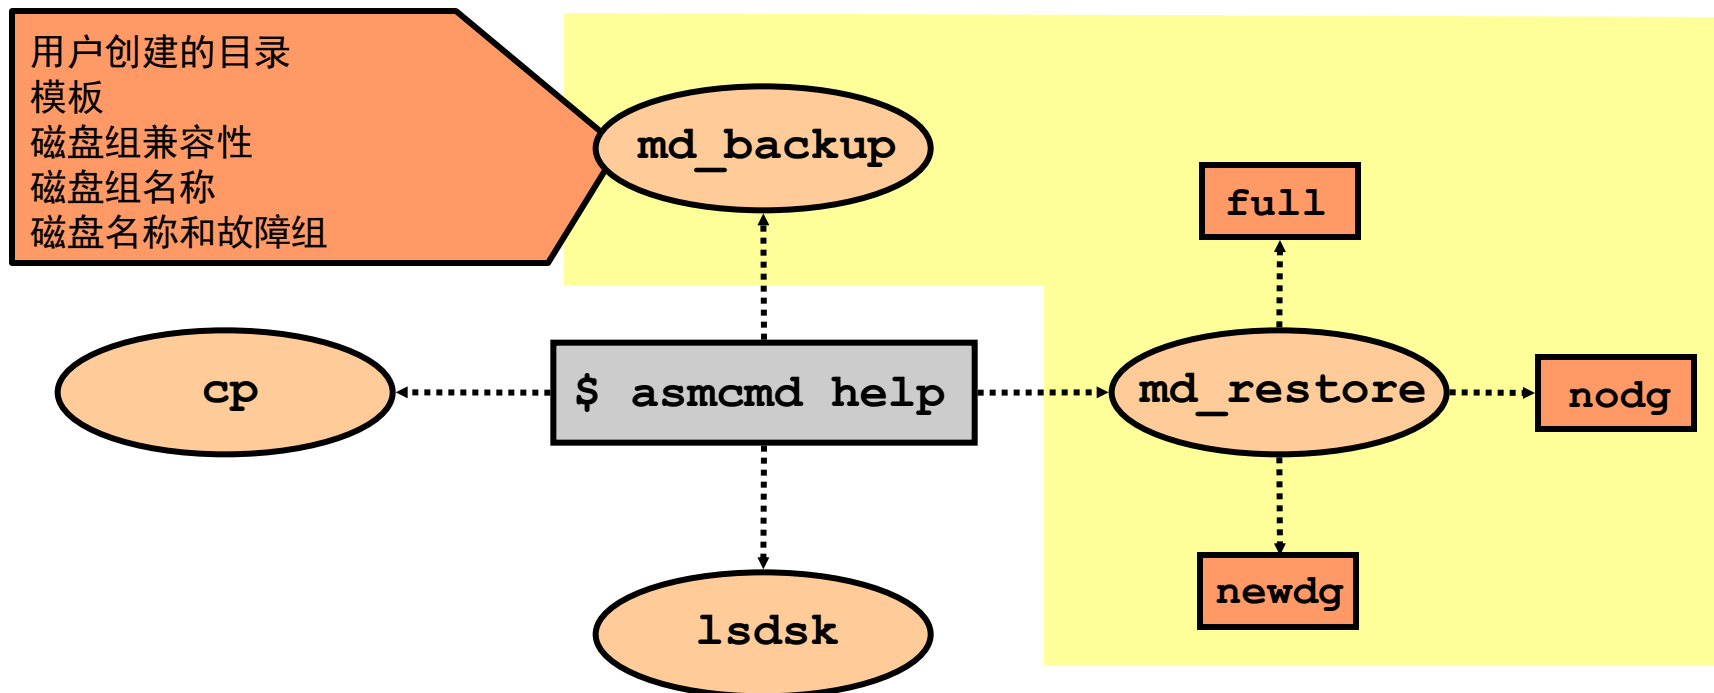
- 允许用户删除无法装载的磁盘组
- 如果在任何位置装载了磁盘组，则操作会失败

**DROP DISKGROUP data FORCE INCLUDING CONTENTS;**



# ASMCMD 扩展

11.1

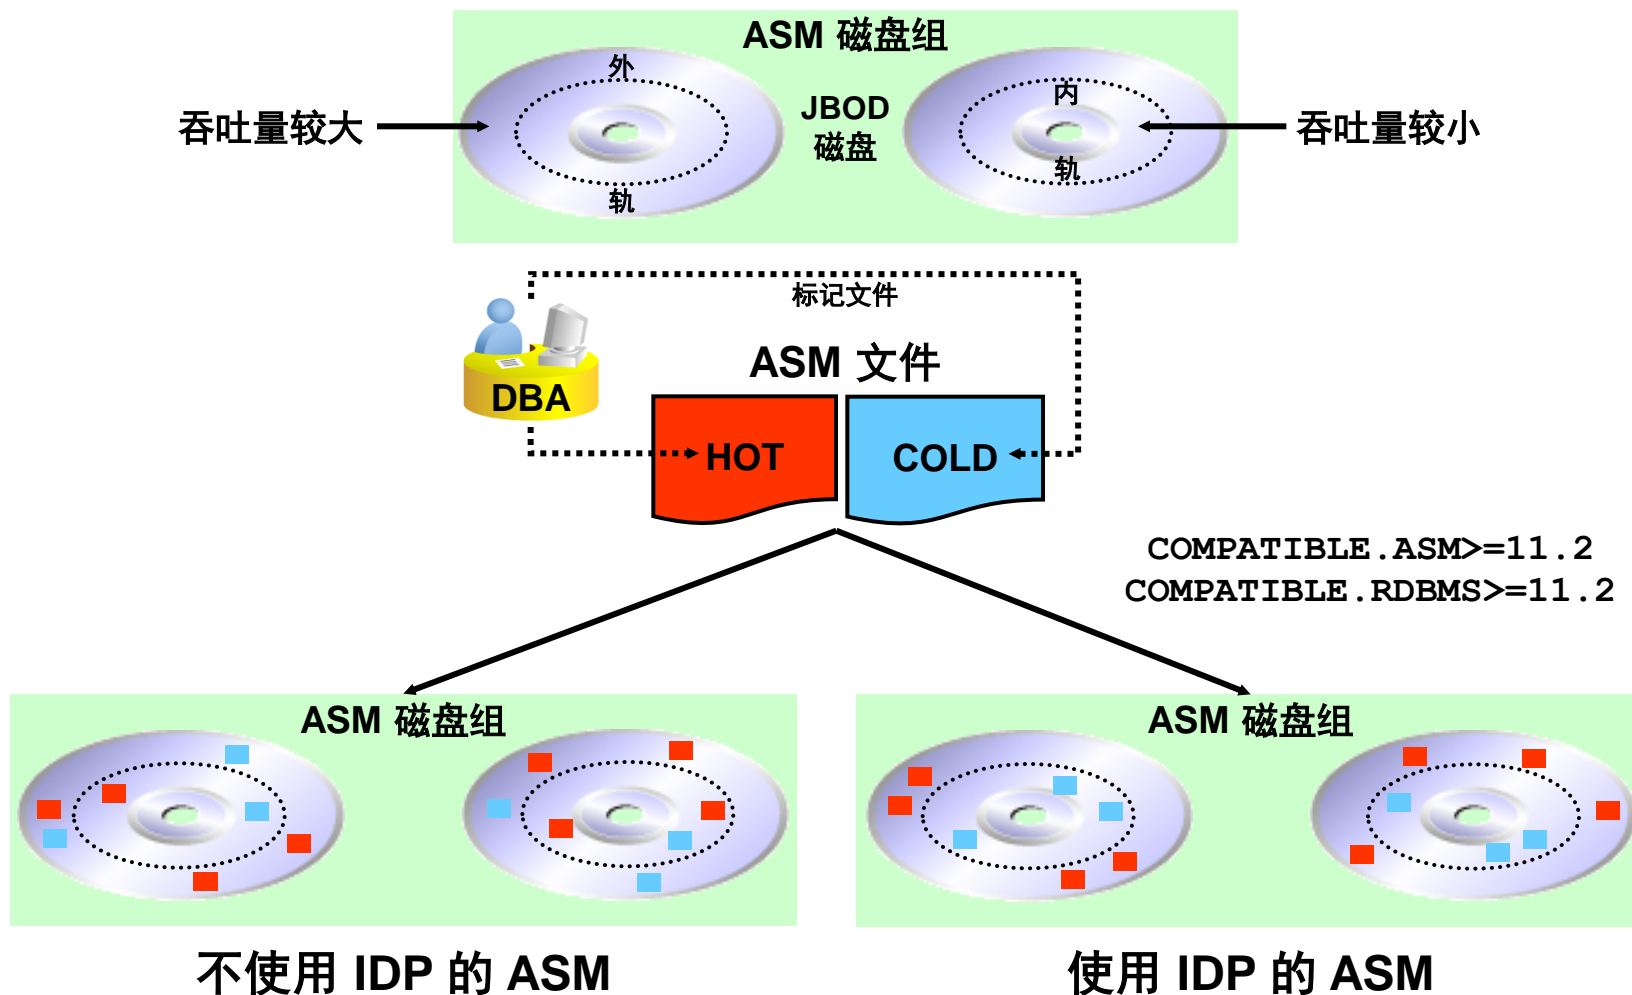


# OCR 和Voting Disk放在ASM上

11.2 更新

- **OCR 是常规ASM文件**
  - 一个新的ASM文件类型
- **Voting Disk**
  - 存储在选定的ASM磁盘上
  - 使用‘crsctl’ 为Voting disk指定一个磁盘组
  - 基于 Ext/Normal/High 冗余级别，ASM 自动创建1/3/5 个Voting Disk
  - Quorum failure group
- ASM 基于磁盘组的冗余级别决定OCR和Voting Disk 的冗余度
- ASM SPFILE参数文件也可以被ASM支持

# ASM 智能数据放置



# ASM 智能数据放置的最佳实践

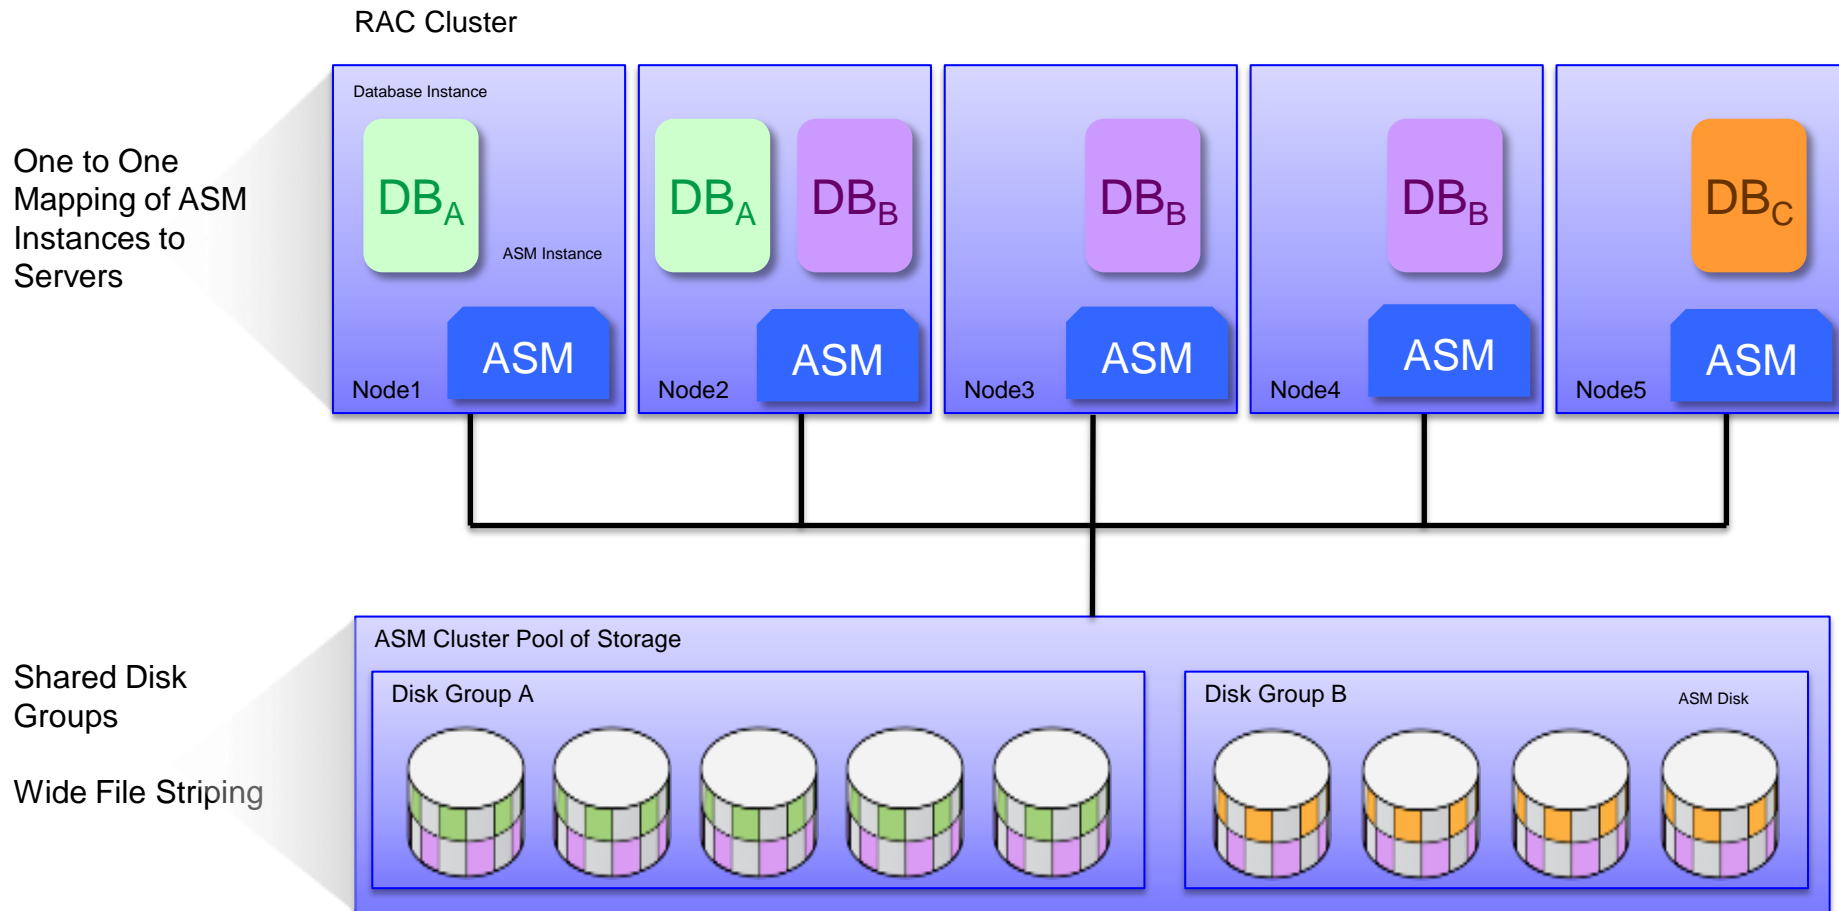
读取频率	写入频率	主要区域	辅助区域
高	高	HOT	MIRRORHOT
高	低	HOT	MIRRORCOLD
低	高	COLD	MIRRORHOT
低	低	COLD	MIRRORCOLD

# 12.1 ASM - 物理元数据冗余

```
NOTE: client cdblunar:cdblunar:ASM mounted group 3 (DATADG1)↓
Sun Nov 08 20:57:42 2015↓
WARNING: cache read a corrupt block: group=3(DATADG1) dsk=0 blk=1 disk=0 (DATADG1_0000) incarn=3916348983 au=0 blk=1 count=1↓
Sun Nov 08 20:57:42 2015↓
Errors in file /u01/app/grid/diag/asm/+asm/+ASM/trace/+ASM_ora_19789.trc:↓
ORA-15196: invalid ASM block header [kfc.c:29297] [endian_kfbh] [2147483648] [1] [0 != 1]↓
NOTE: a corrupted block from group DATADG1 was dumped to /u01/app/grid/diag/asm/+asm/+ASM/trace/+ASM_ora_19789.trc↓
WARNING: cache read (retry) a corrupt block: group=3(DATADG1) dsk=0 blk=1 disk=0 (DATADG1_0000) incarn=3916348983 au=0 blk=1 count=1↓
Sun Nov 08 20:57:42 2015↓
Errors in file /u01/app/grid/diag/asm/+asm/+ASM/trace/+ASM_ora_19789.trc:↓
ORA-15196: invalid ASM block header [kfc.c:29297] [endian_kfbh] [2147483648] [1] [0 != 1]↓
ORA-15196: invalid ASM block header [kfc.c:29297] [endian_kfbh] [2147483648] [1] [0 != 1]↓
WARNING: Failed to verify disk 0 (DATADG1_0000) of group 3 (DATADG1) path /dev/sdb1 reason: endian_kfbh 0 != 1↓
NOTE: corrupt disk header dumped to /u01/app/grid/diag/asm/+asm/+ASM/trace/+ASM_ora_19789.trc↓
NOTE: cache repaired a corrupt block: group=3(DATADG1) dsk=0 blk=1 on disk 0 from disk=0 (DATADG1_0000) incarn=3916348983 au=11 blk=1 count=1↓
WARNING: cache read a corrupt block: group=3(DATADG1) dsk=0 blk=3 disk=0 (DATADG1_0000) incarn=3916348983 au=0 blk=3 count=1↓
Sun Nov 08 20:57:42 2015↓
Errors in file /u01/app/grid/diag/asm/+asm/+ASM/trace/+ASM_ora_19789.trc:↓
ORA-15196: invalid ASM block header [kfc.c:29297] [endian_kfbh] [2147483648] [3] [0 != 1]↓
NOTE: a corrupted block from group DATADG1 was dumped to /u01/app/grid/diag/asm/+asm/+ASM/trace/+ASM_ora_19789.trc↓
WARNING: cache read (retry) a corrupt block: group=3(DATADG1) dsk=0 blk=3 disk=0 (DATADG1_0000) incarn=3916348983 au=0 blk=3 count=1↓
Sun Nov 08 20:57:42 2015↓
Errors in file /u01/app/grid/diag/asm/+asm/+ASM/trace/+ASM_ora_19789.trc:↓
ORA-15196: invalid ASM block header [kfc.c:29297] [endian_kfbh] [2147483648] [3] [0 != 1]↓
ORA-15196: invalid ASM block header [kfc.c:29297] [endian_kfbh] [2147483648] [3] [0 != 1]↓
WARNING: Failed to verify disk 0 (DATADG1_0000) of group 3 (DATADG1) path /dev/sdb1 reason: endian_kfbh 0 != 1↓
NOTE: corrupt disk header dumped to /u01/app/grid/diag/asm/+asm/+ASM/trace/+ASM_ora_19789.trc↓
NOTE: cache repaired a corrupt block: group=3(DATADG1) dsk=0 blk=3 on disk 0 from disk=0 (DATADG1_0000) incarn=3916348983 au=11 blk=3 count=1↓
SYS@cdblunar>↓
```

# Automatic Storage Management (ASM) Overview

## Current State



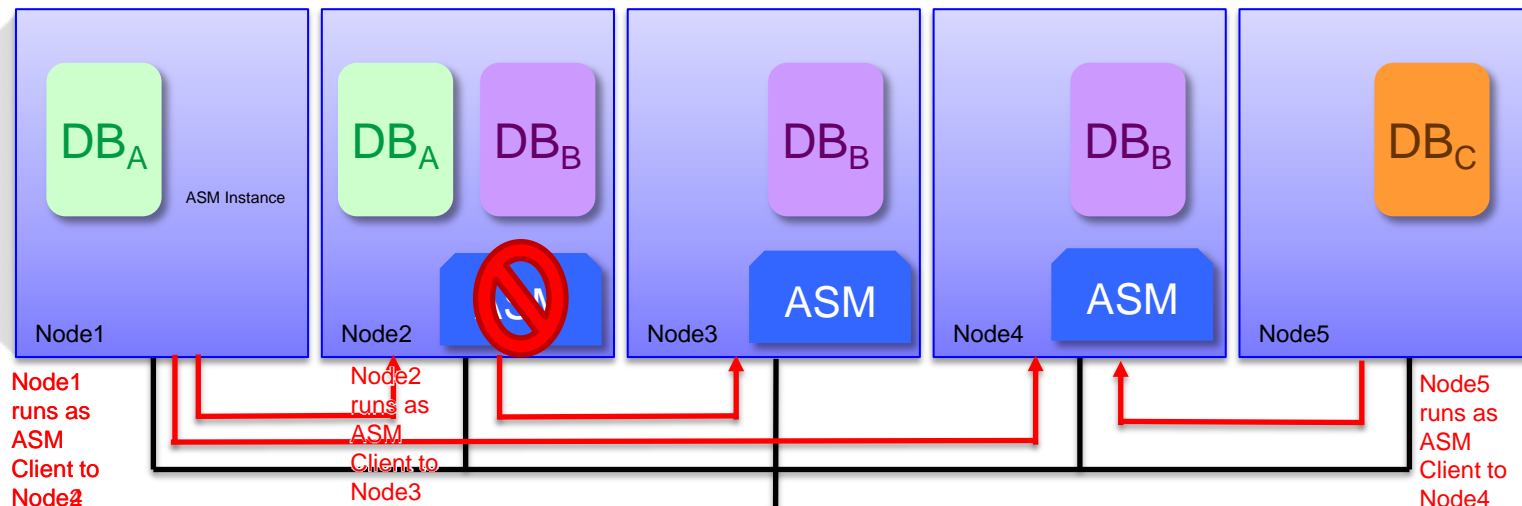


# Flex ASM: Eliminate 1:1 Server Mapping

## New: ASM Storage Consolidation in Oracle Database 12c


RAC Cluster

Databases share  
ASM instances



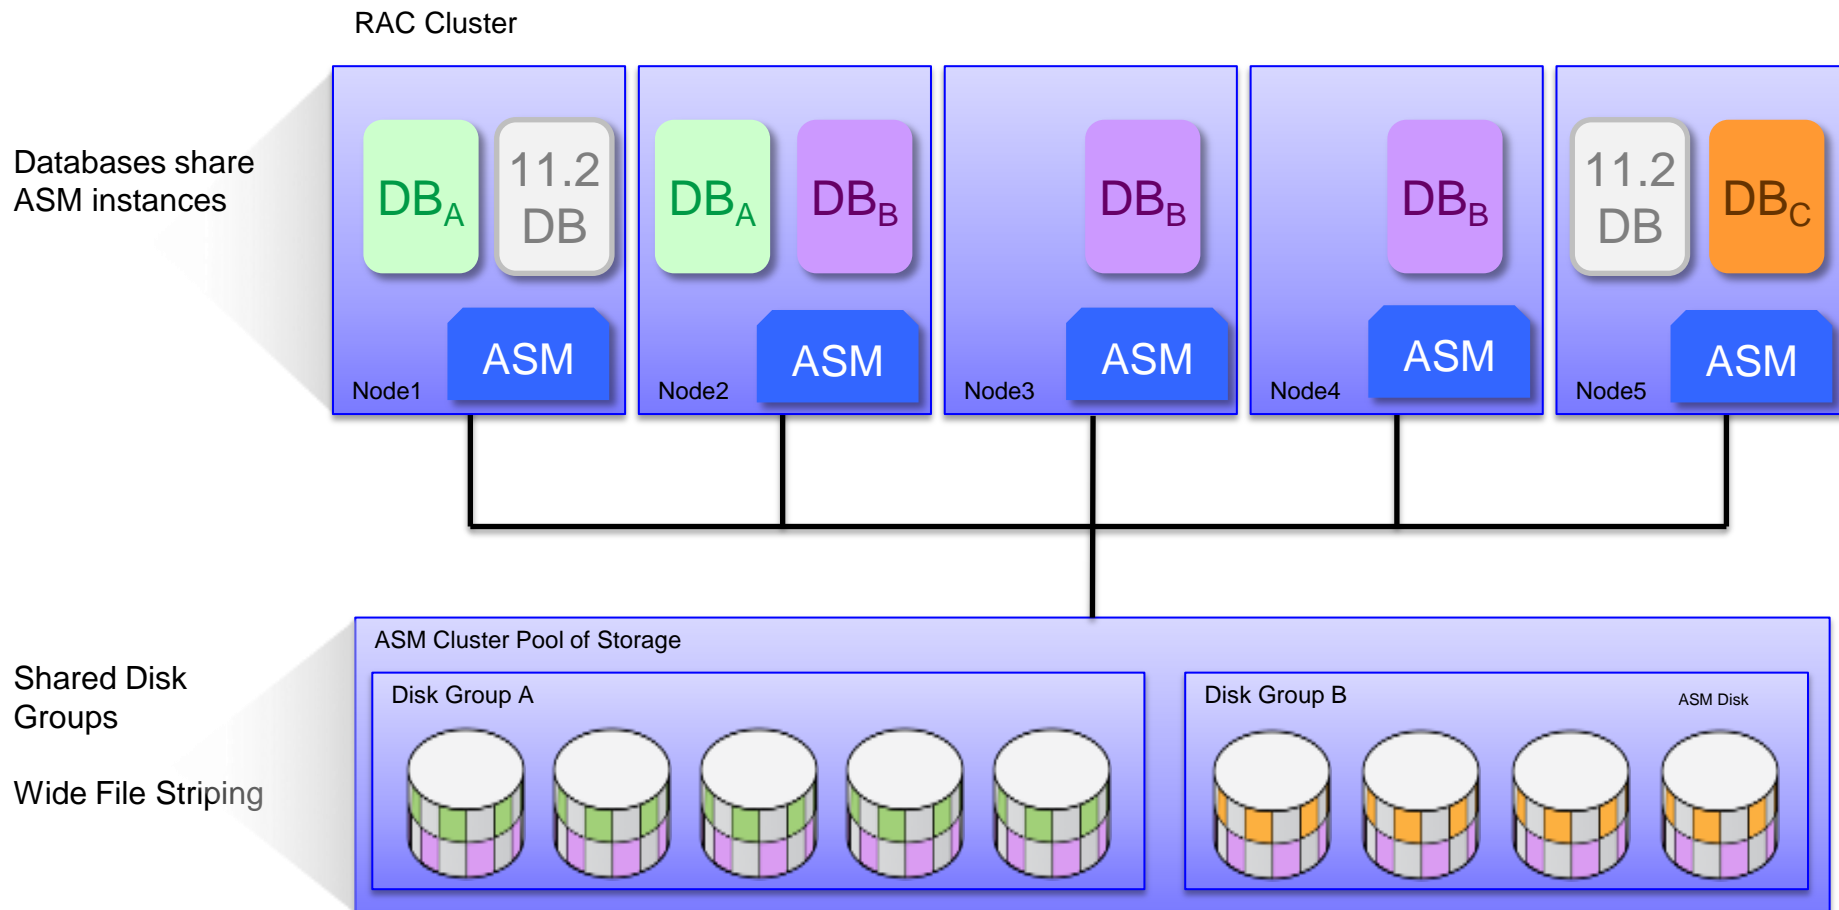
Shared Disk  
Groups

Wide File Striping



# Flex ASM: Supporting Oracle Database 11g

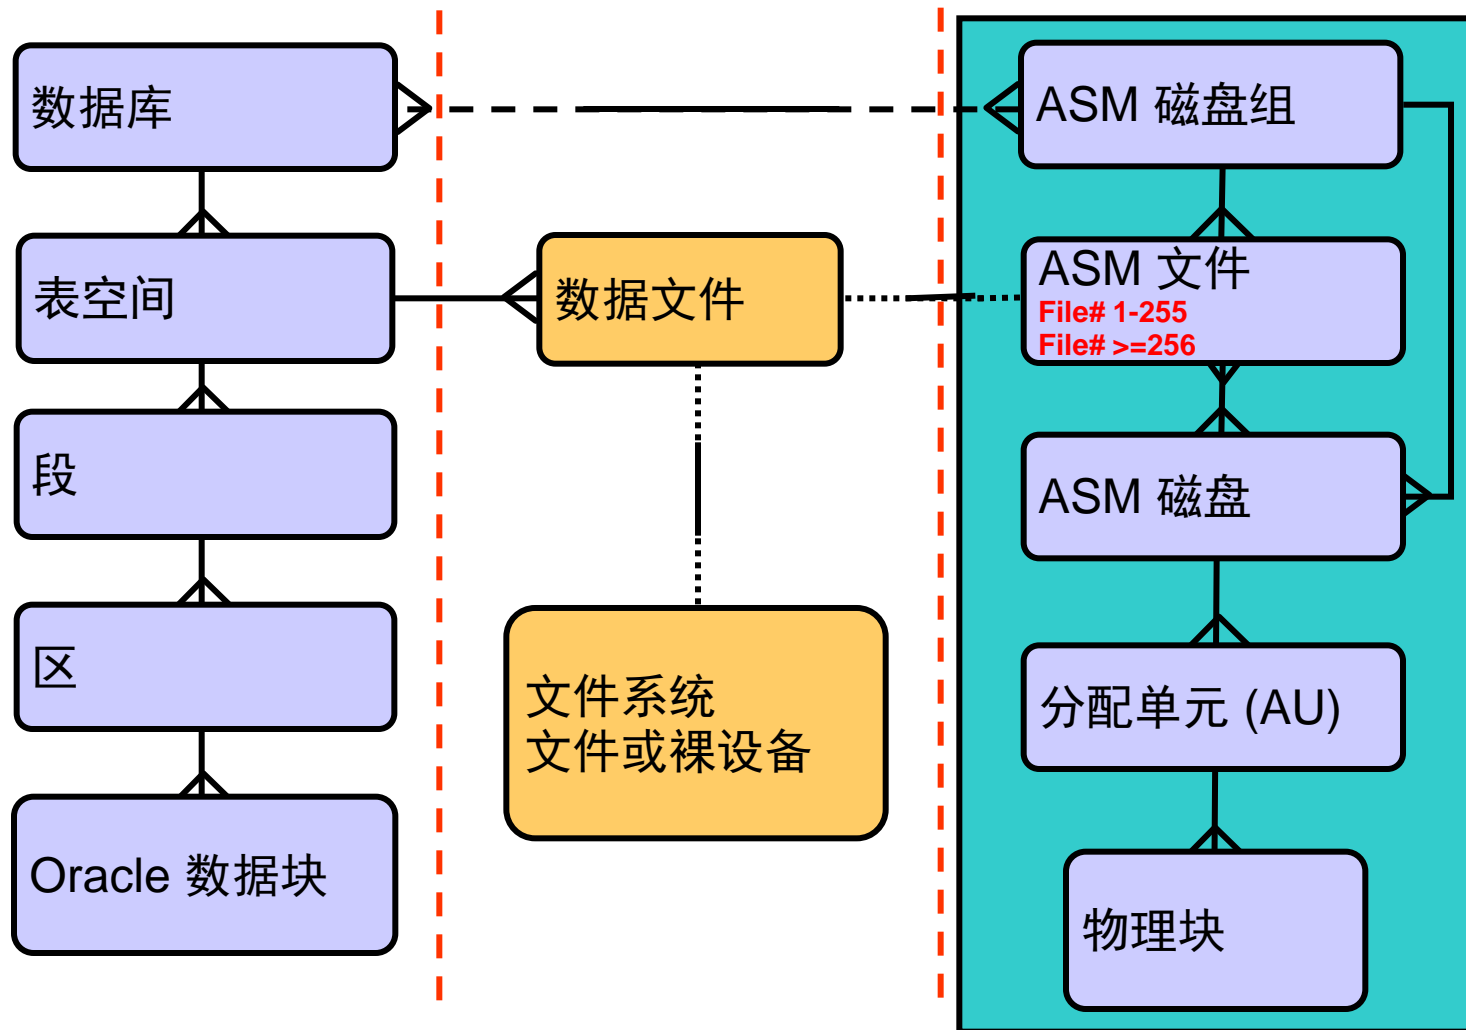
Previous Database Versions Will Host Local ASM Instance



# Agend

- ASM的历史和简介
- ASM体系架构
- ASM在不同版本的新特性
- **ASM存储和分配机制**
- ASM相关产品介绍


# ASM 存储：概念



# ASM 磁盘组

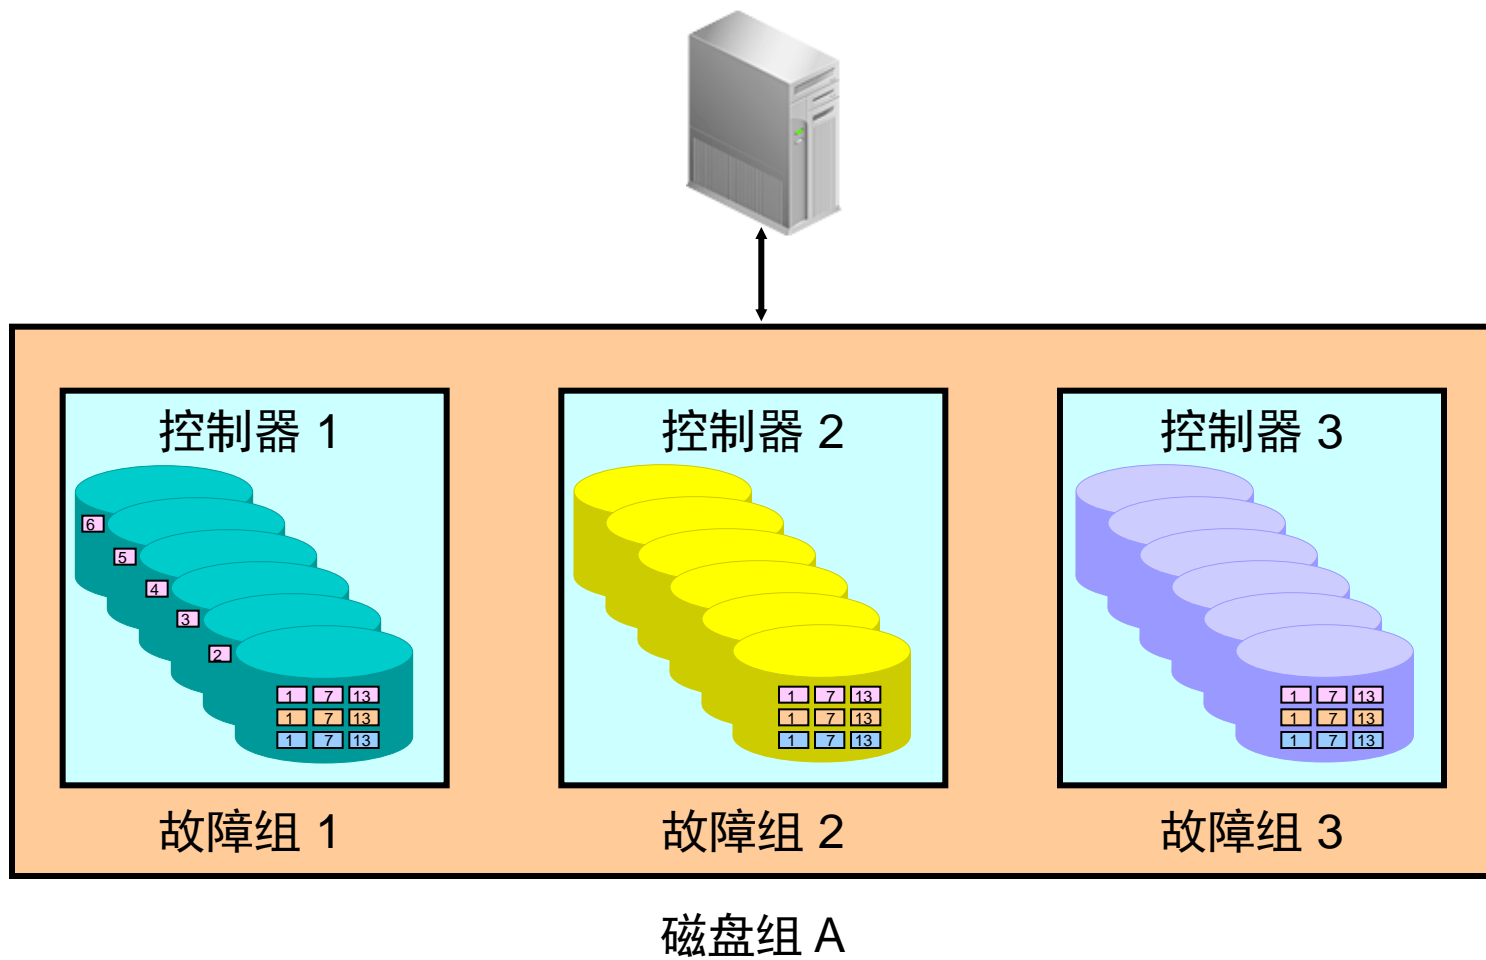
- 作为逻辑单元管理的磁盘组
- 将磁盘总空间划分为统一大小的单元
- 将各个文件平均分配到所有磁盘中
- 根据文件类型使用粗粒度或细粒度的条带化
- 管理的是磁盘组而非文件

ASM 实例



磁盘组

# 故障组



# ASM磁盘和磁盘组（Disk Groups）

**ASM 磁盘 = 磁盘分区 或 LUN**


**ASM 磁盘组 = 作为一个逻辑单元进行管理的ASM 磁盘池**

把每一个文件平均散布到跨**ASM**磁盘组的所有**ASM**磁盘


- 使用的分配单元（**AU**）大小为 1 MB

条带化

- 基于文件的类型 (**V\$ASM\_TEMPLATE**)
- Coarse (1 MB)
- Fine (128 KB)

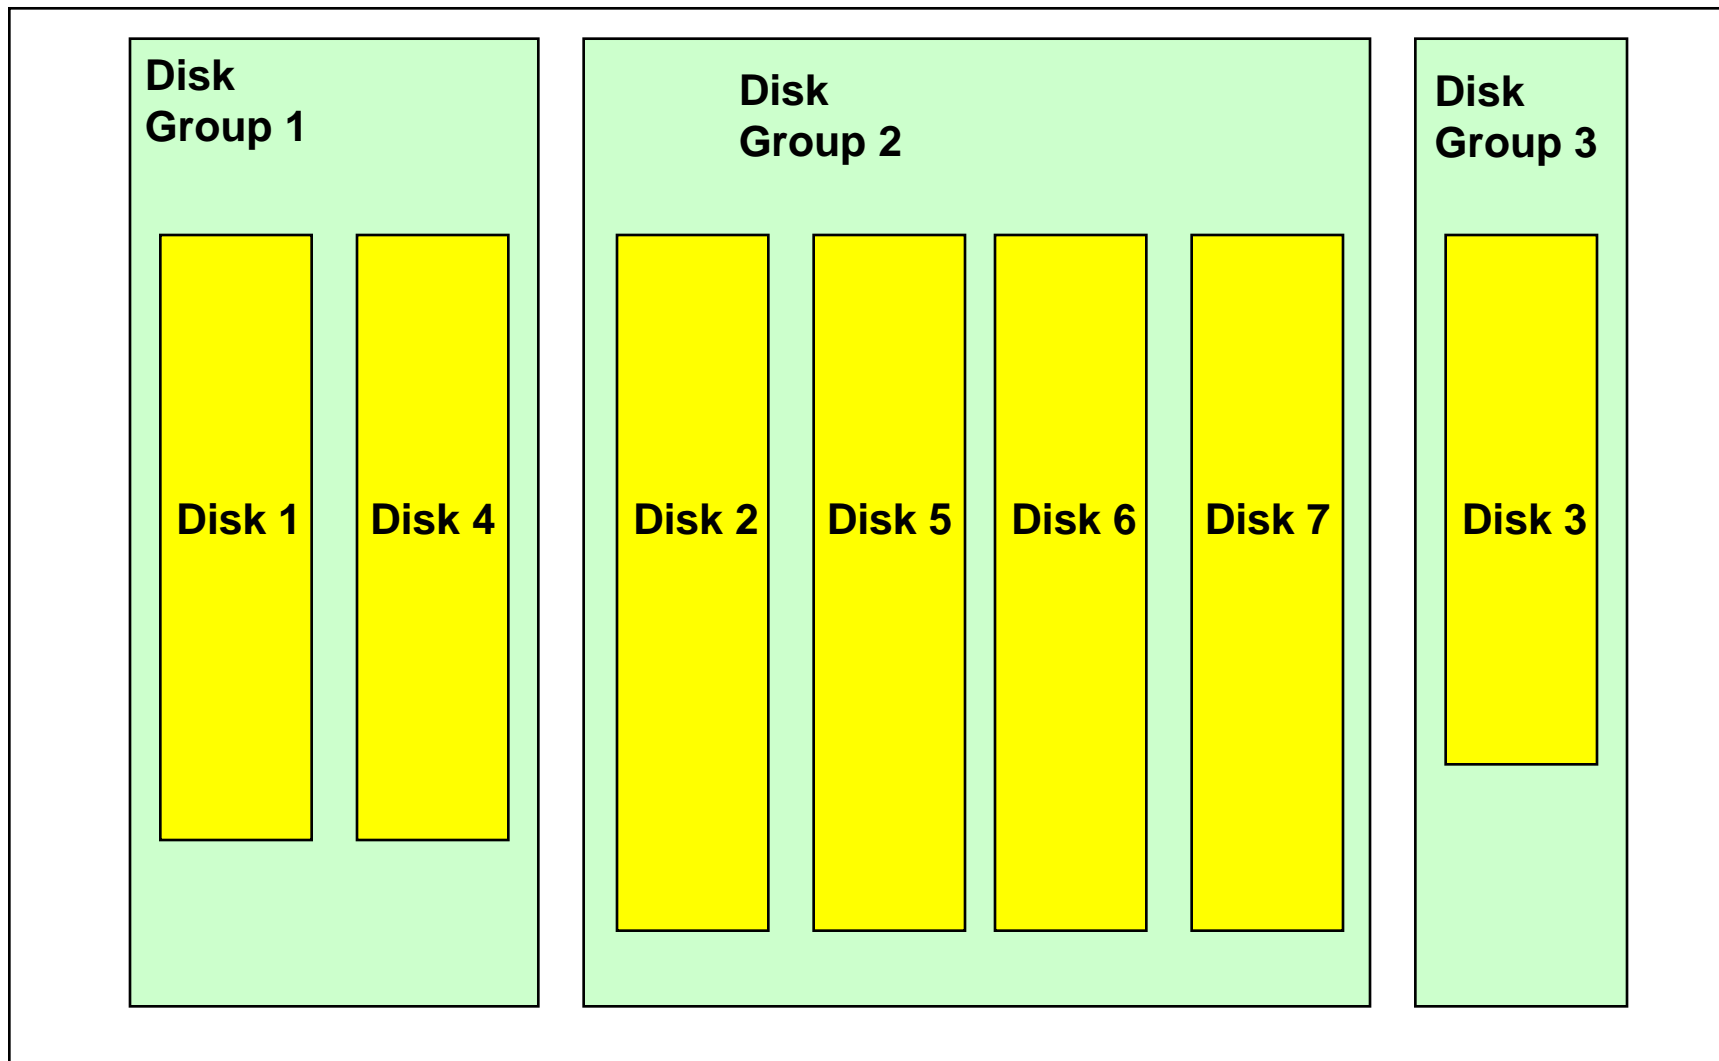


# ASM 磁盘组和磁盘

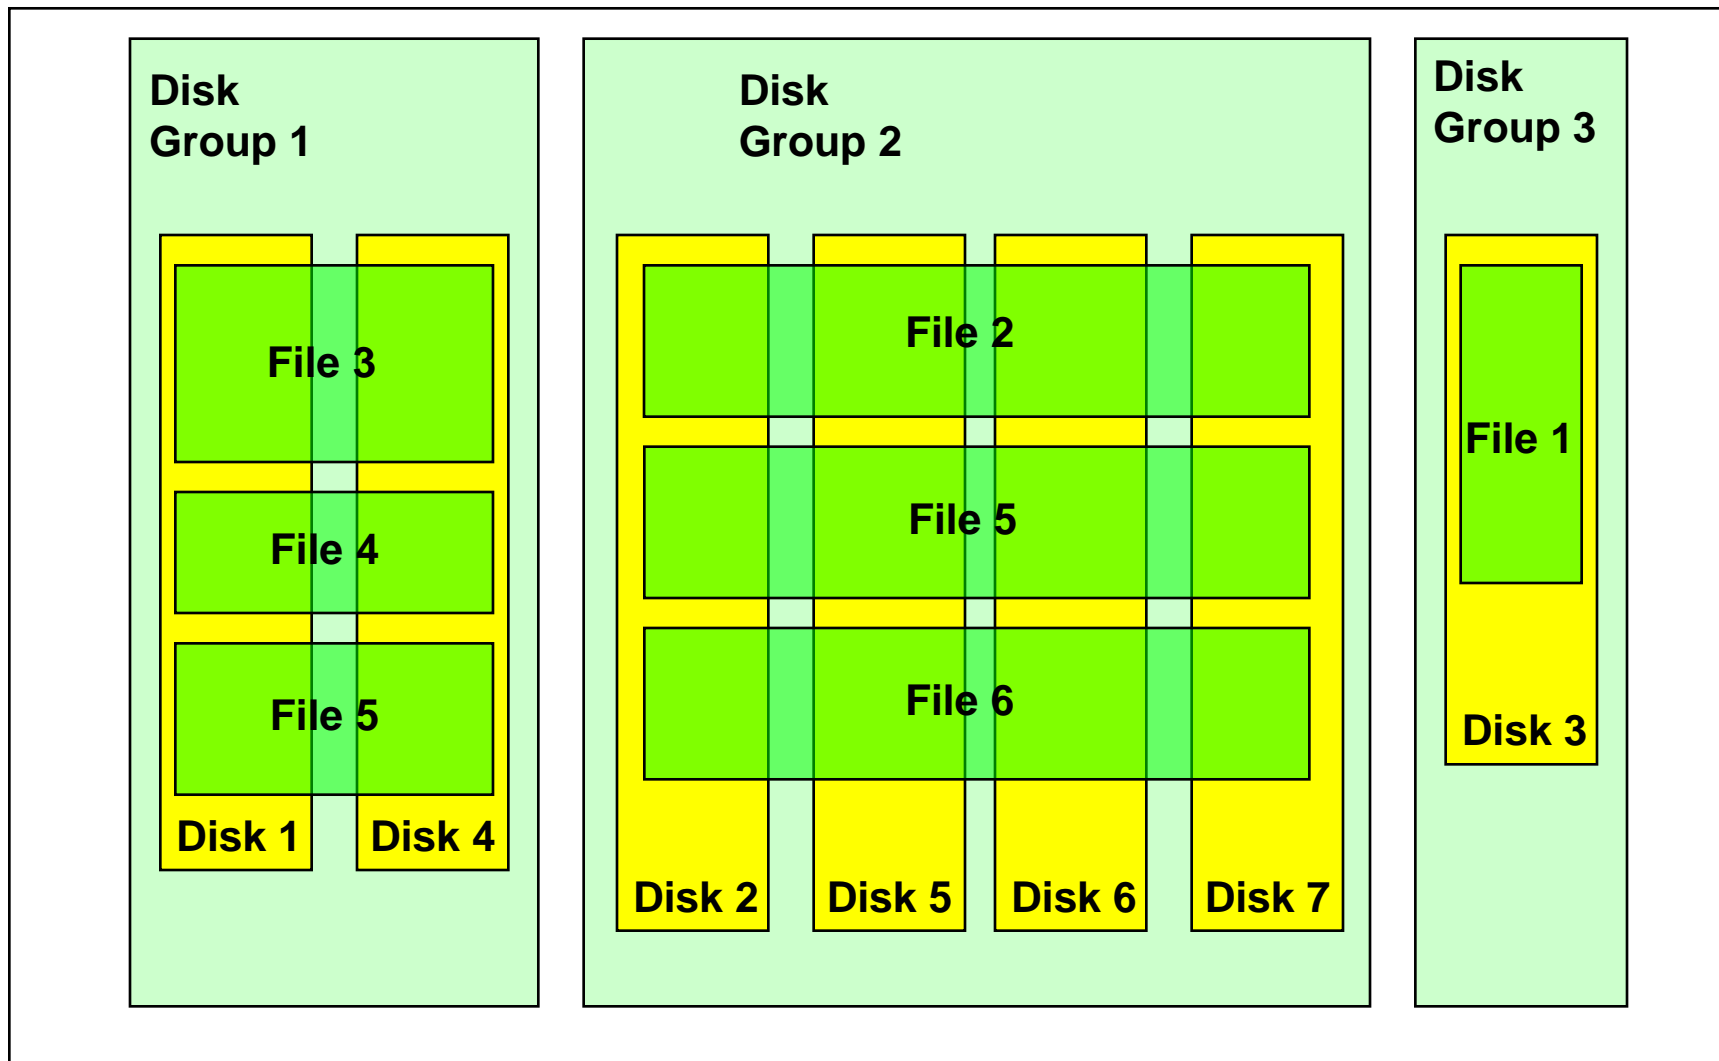




# ASM 磁盘组和磁盘




# ASM 磁盘组、磁盘和数据库文件




# Mirroring & Stripping

- *Mirroring* : Redundant copies of data over multiple disks




■ *Primary data extents*  
■ *Secondary data extents*

- *Stripping* : Spreading data over multiple disks



# Allocation Unit

**ASM Disk**



**1 MB**

**ASM disks are divided into allocation units**

**Unit of allocating storage to a file**

**One megabyte in size**

**Small enough not to be hot**


**Large enough for efficient sequential access**

**Many allocation units in a single file**

**`_ASM_AUSIZE`**

# Multiple Allocation Unit Size support

- Higher performance for large seq I/O (DW)
- Better leverage of Hardware RAID read-ahead
- Set Oracle MAXIO = AU\_size



## Allocation Unit (AU)

Selected at disk group creation time  
and may be 1,2,4,8,16,32,64 MB

## Variable size ASM file extents

Extent size = AU size up to 20,000 extents

Extent size = 8\*AU up to 40,000 extents

Extent size = 64\*AU beyond 40,001 extents


## Striping

Coarse Stripe size always = one AU

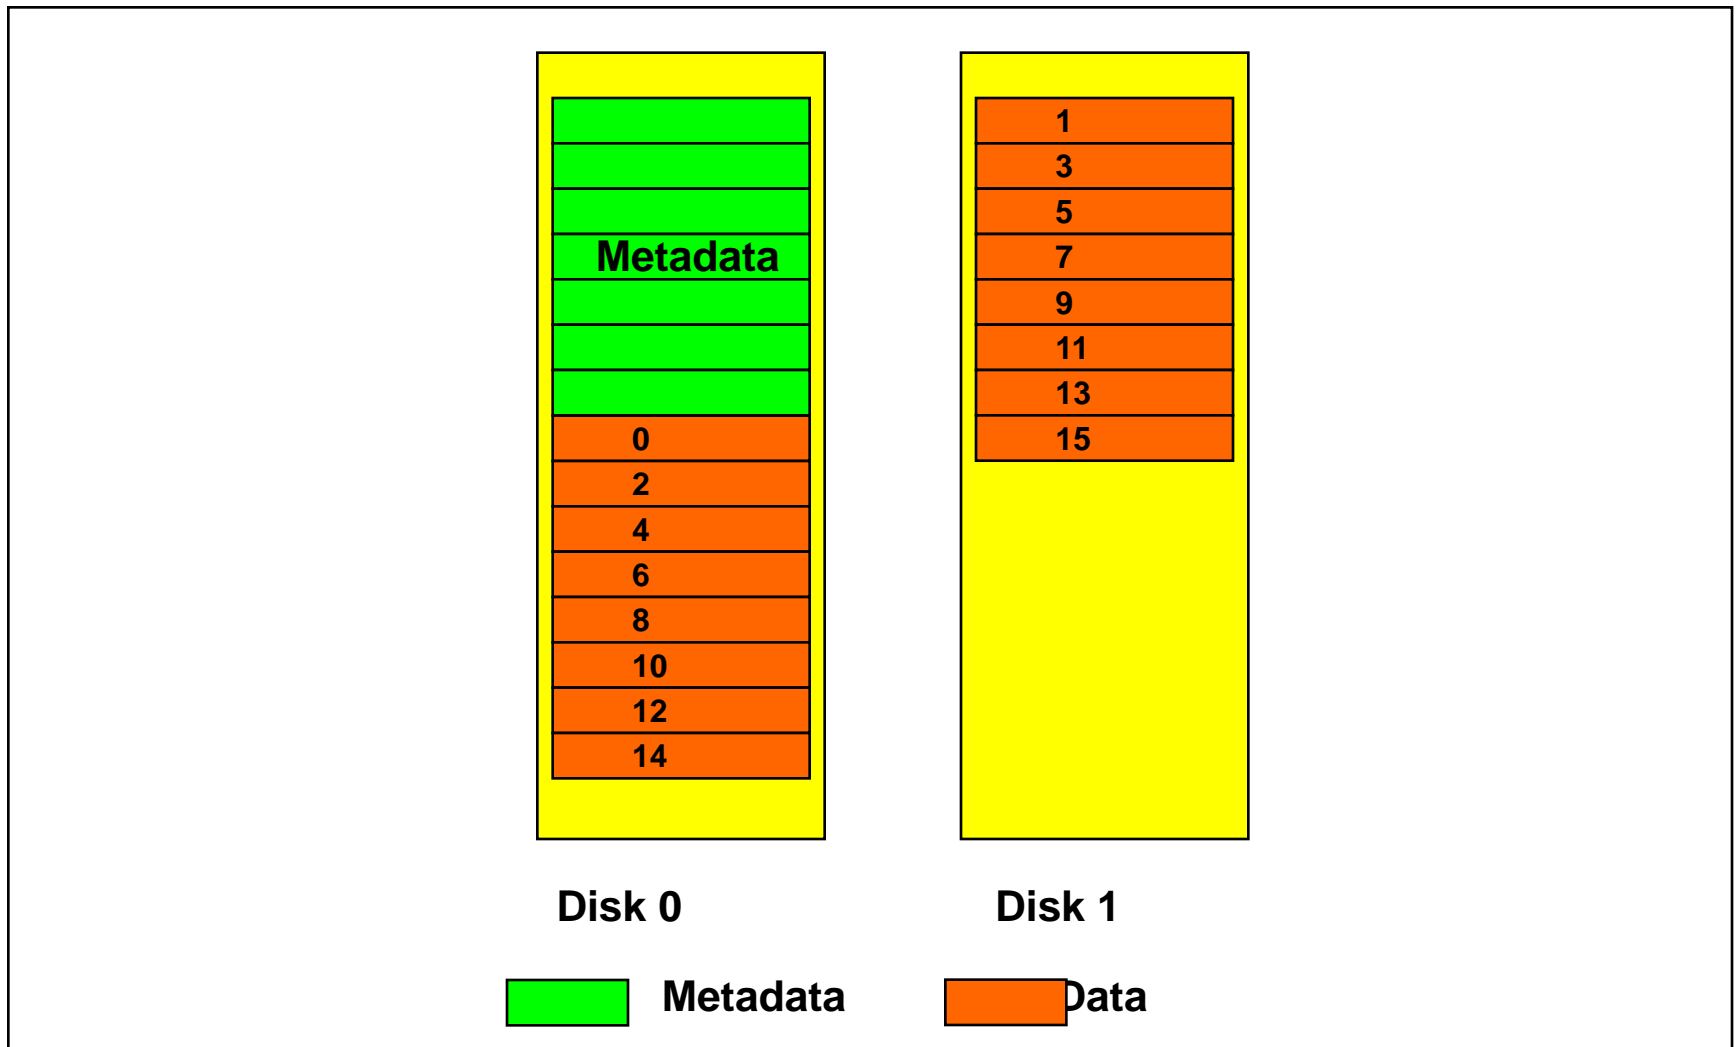
Fine Stripe size always = 128 KB

# ASM中Extent 的分布-1块磁盘成1个磁盘组

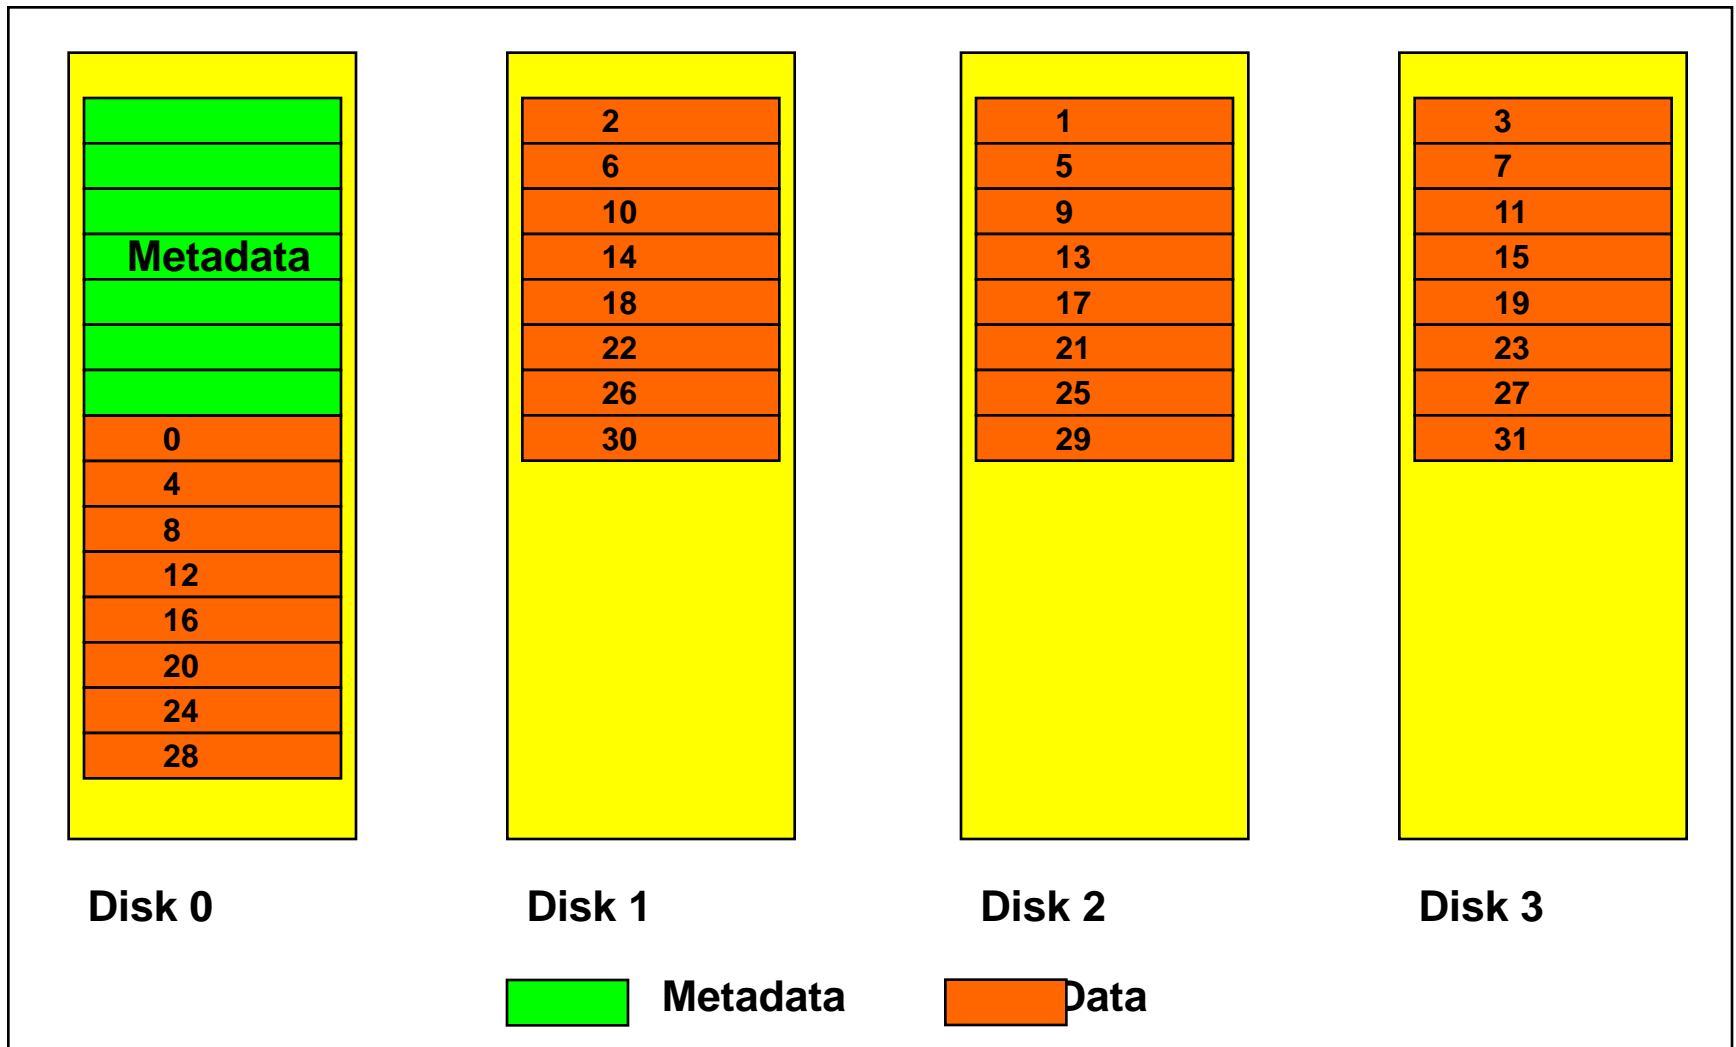
<http://juliandyke.com/>



# ASM中Extent 的分布-2块磁盘成1个磁盘组

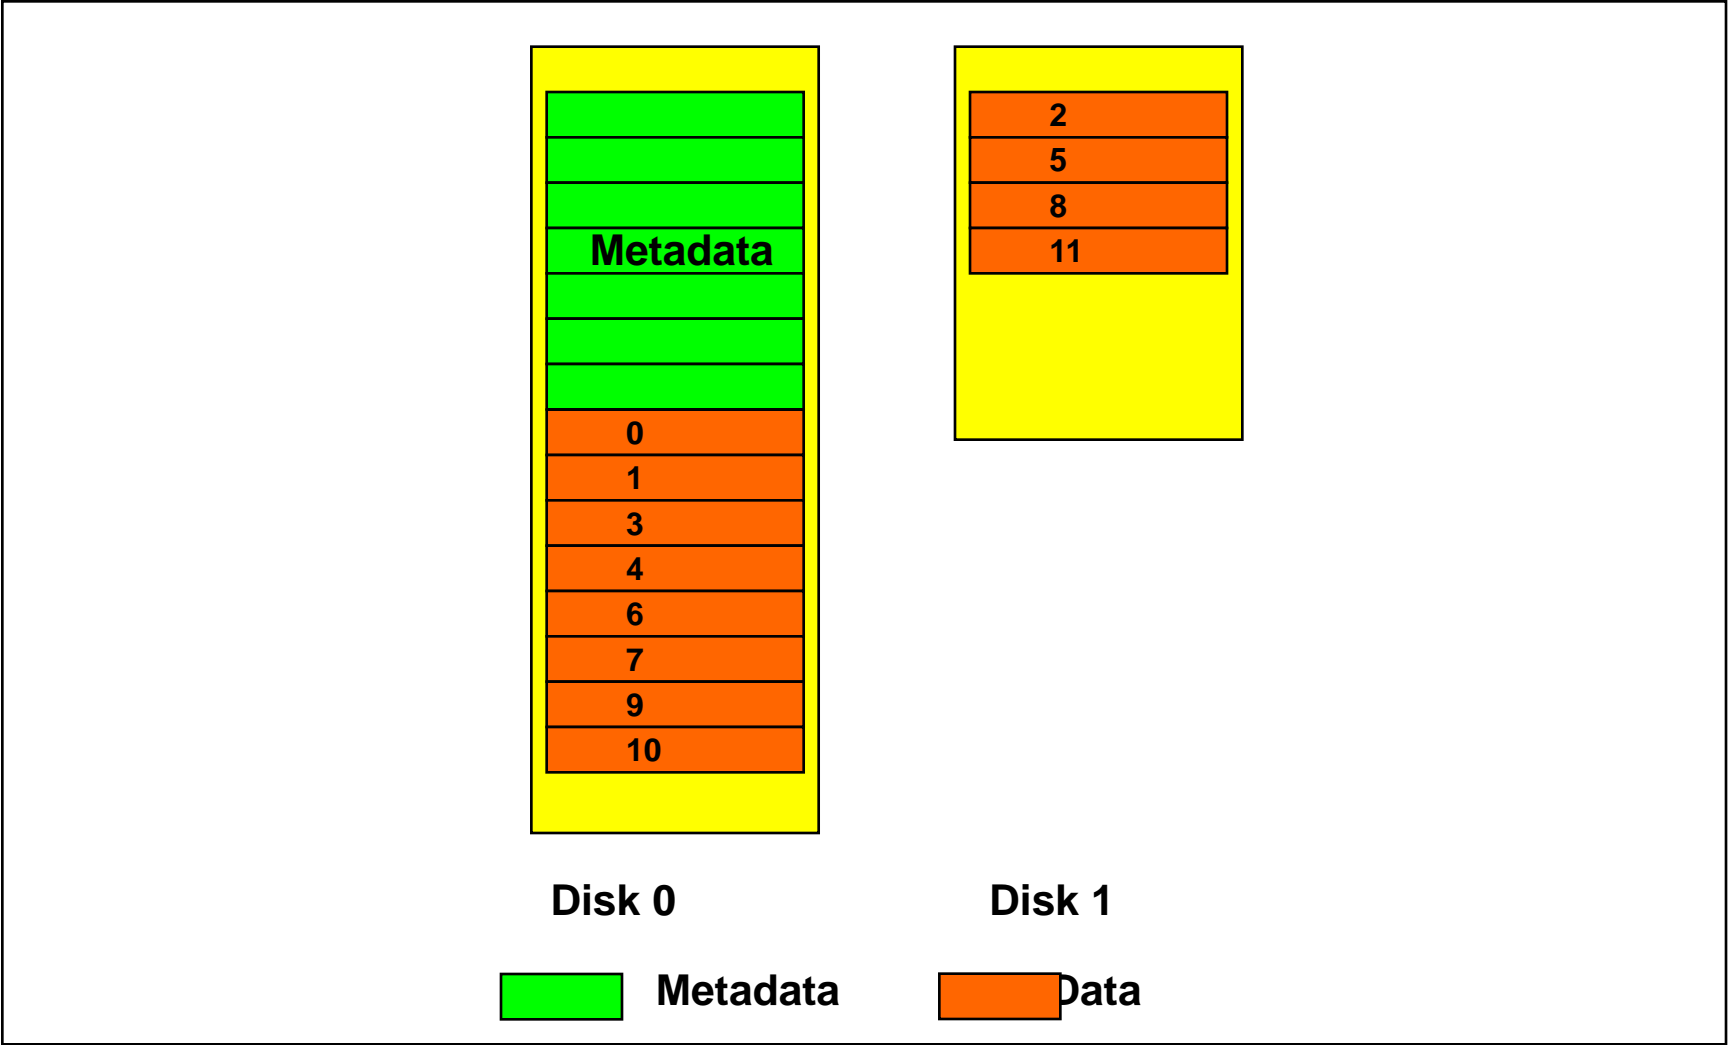


# ASM中Extent 的分布-4块磁盘成1个磁盘组

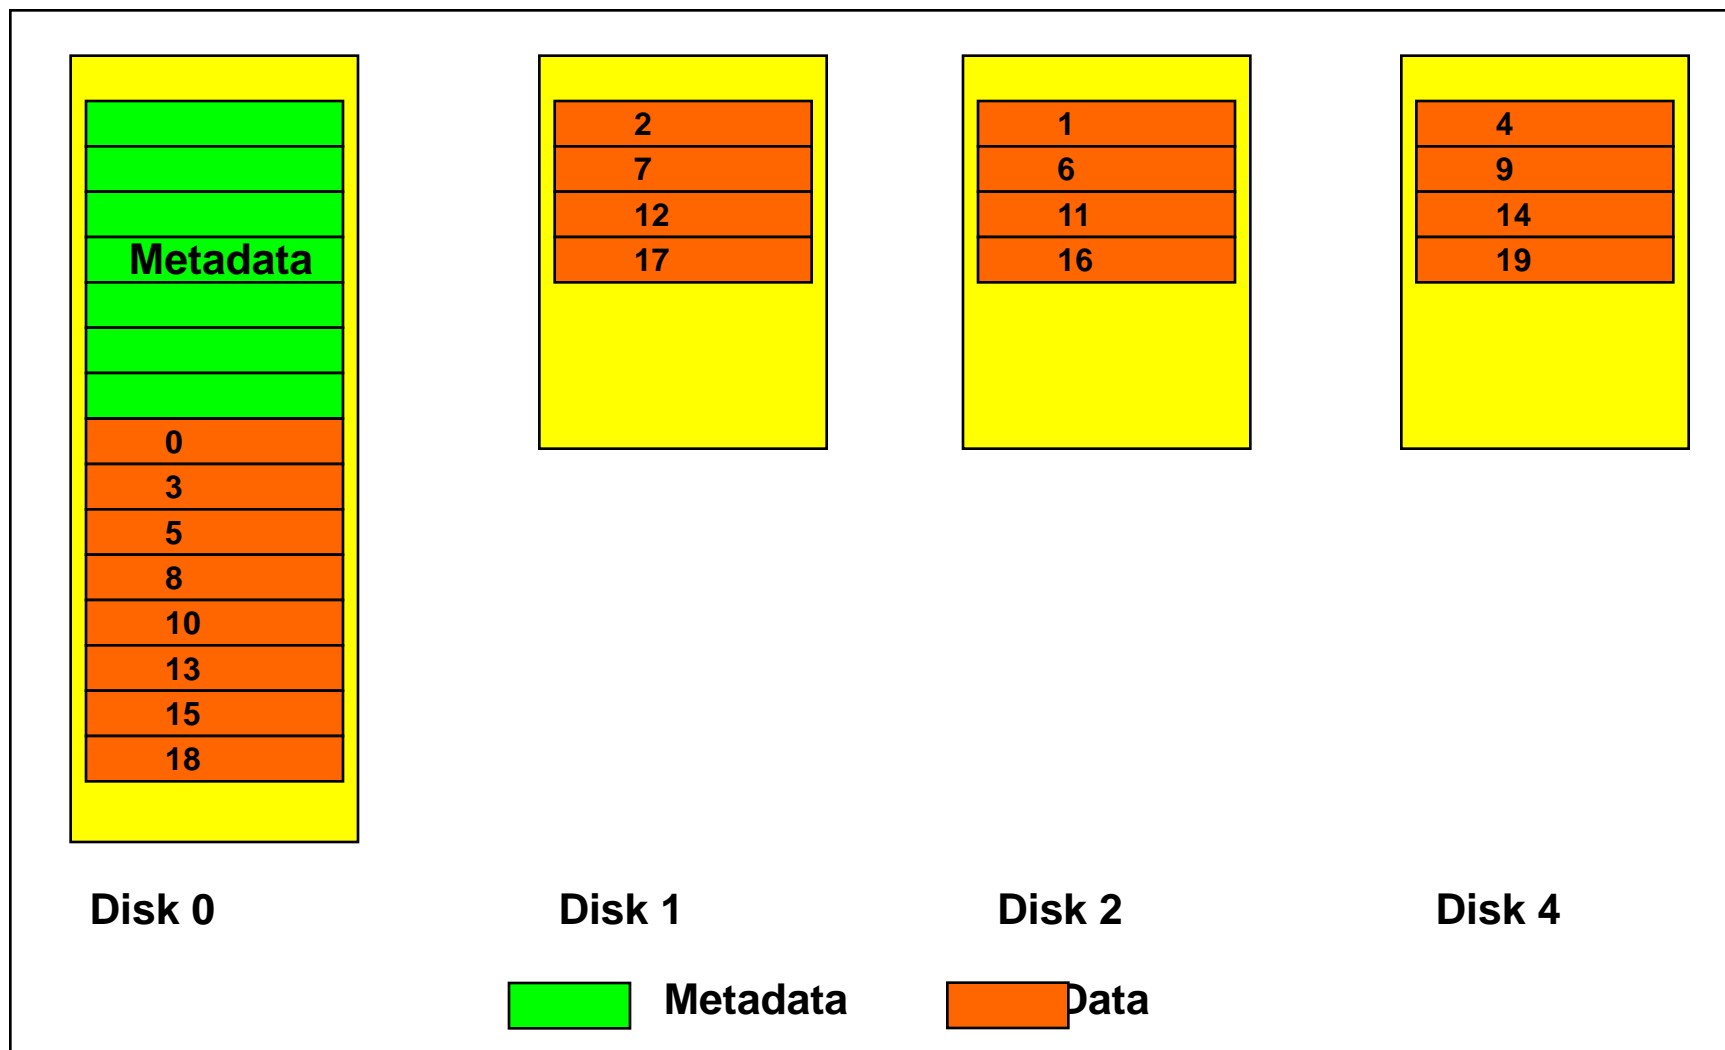




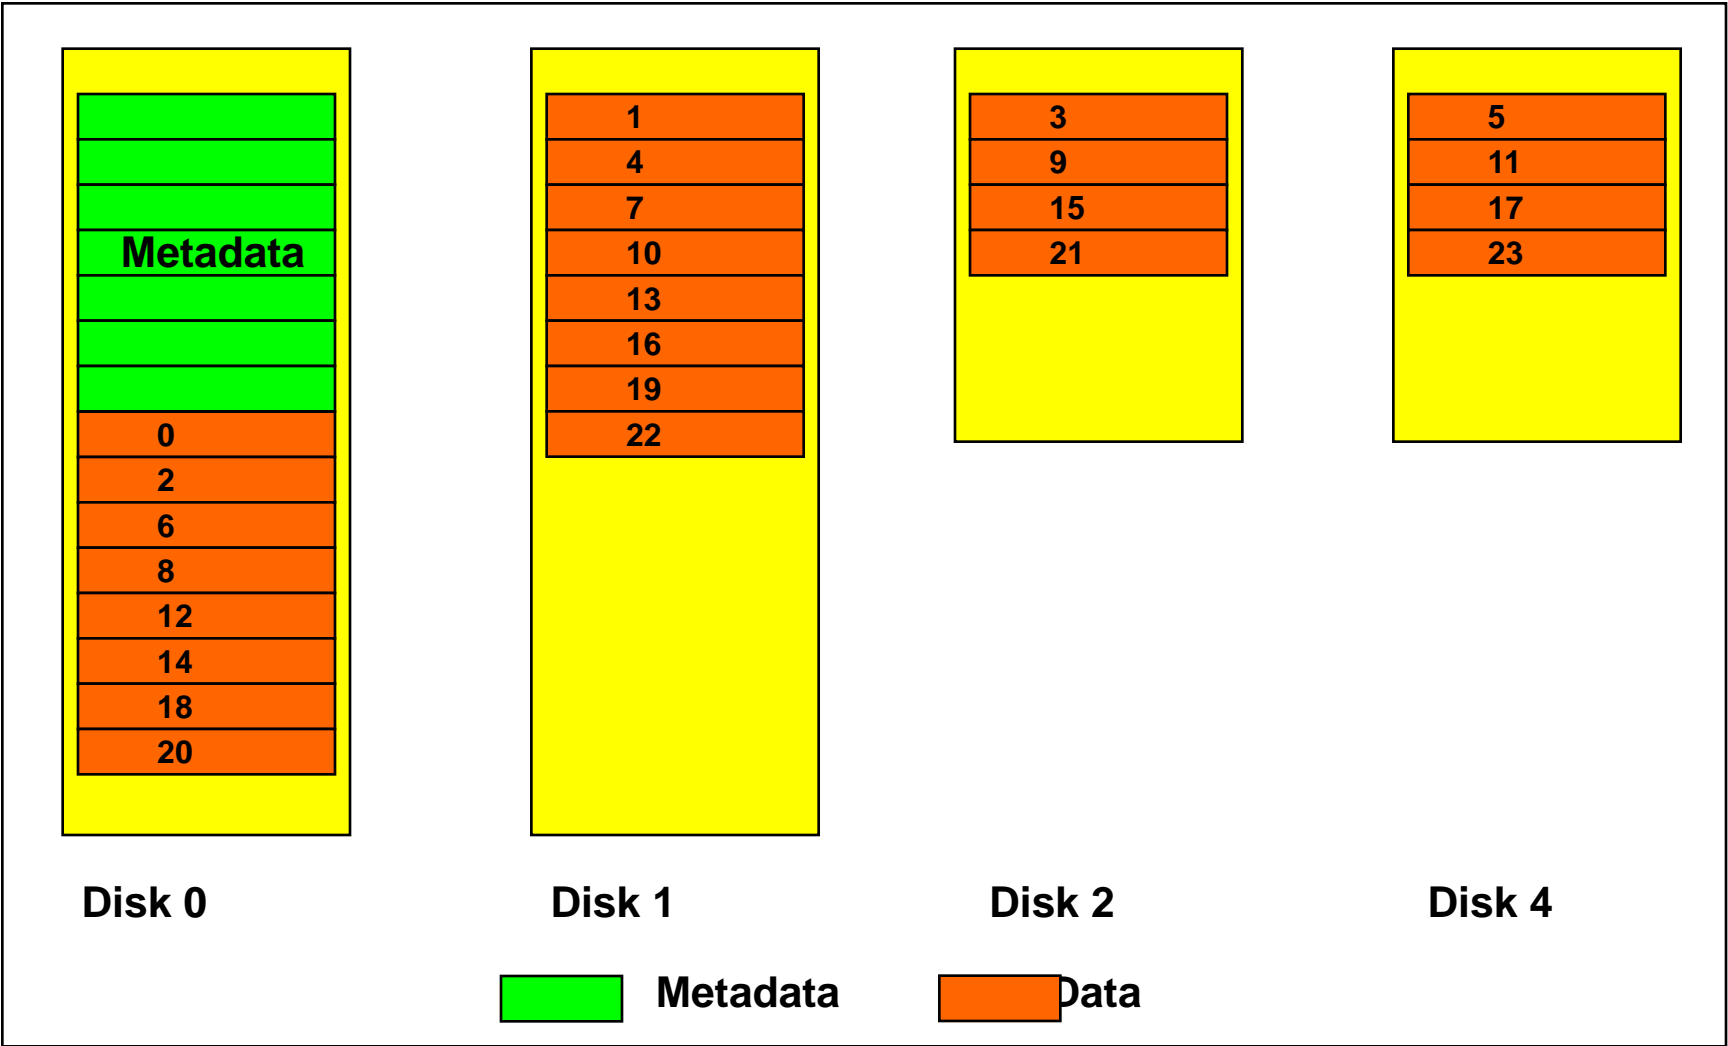
# ASM中Extent 的分布-1块大磁和1块小盘组成一个磁盘组



# ASM中Extent 的分布-1块大磁和3块小盘组成一个磁盘组




# ASM中Extent 的分布-2块大磁和2块小盘组成一个磁盘组




# 磁盘组的镜像

- 在extent级别镜像
- **Primary** 和 **mirror extent** 不在相同的**ASM**磁盘
- 失败组的概念（**Failure group**）
  - **Default failure group = Disk**
- **External redundancy**
  - 由底层硬件实现mirroring
- **Normal redundancy**
  - 2-way mirroring
- **High redundancy**
  - 3-way mirroring



# Normal Redundancy




Failure group 1

Failure group 2

Five megabyte normal redundancy file

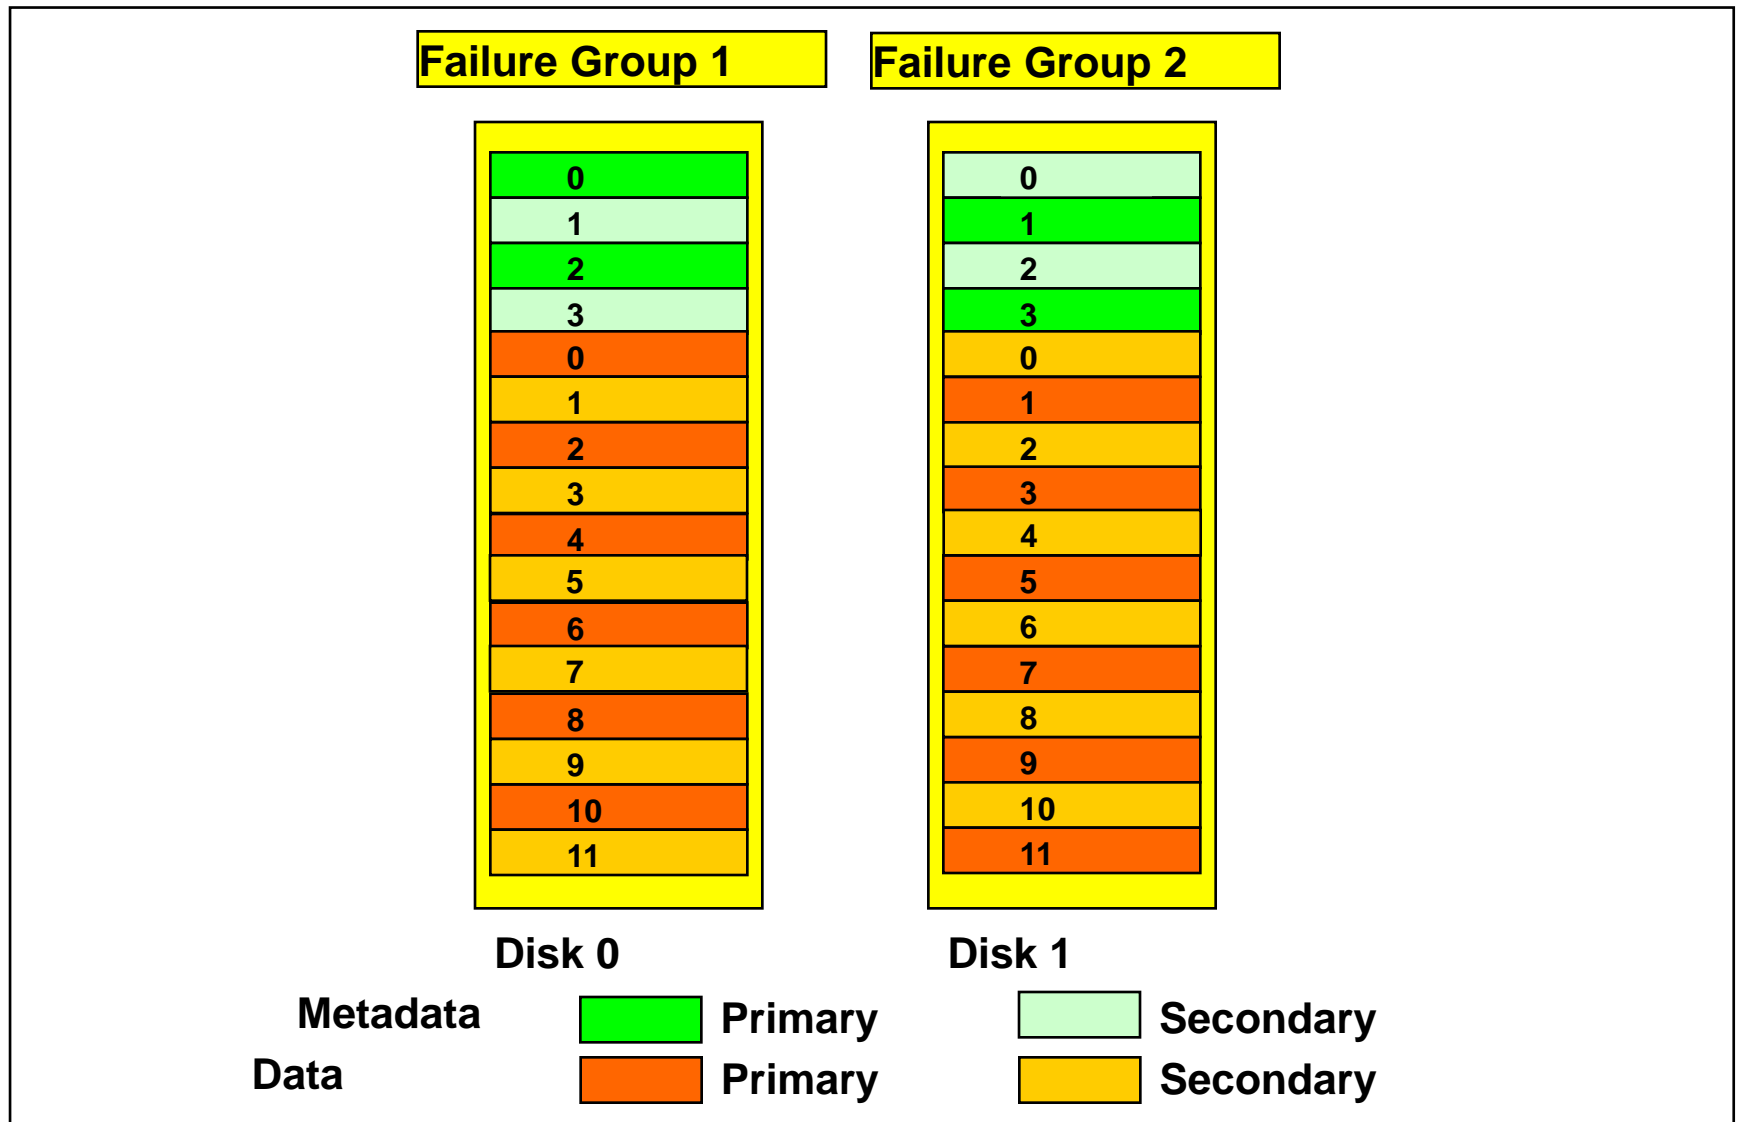
# Primary Extents    # Secondary Extents

# Virtual Extents (Extent Sets)



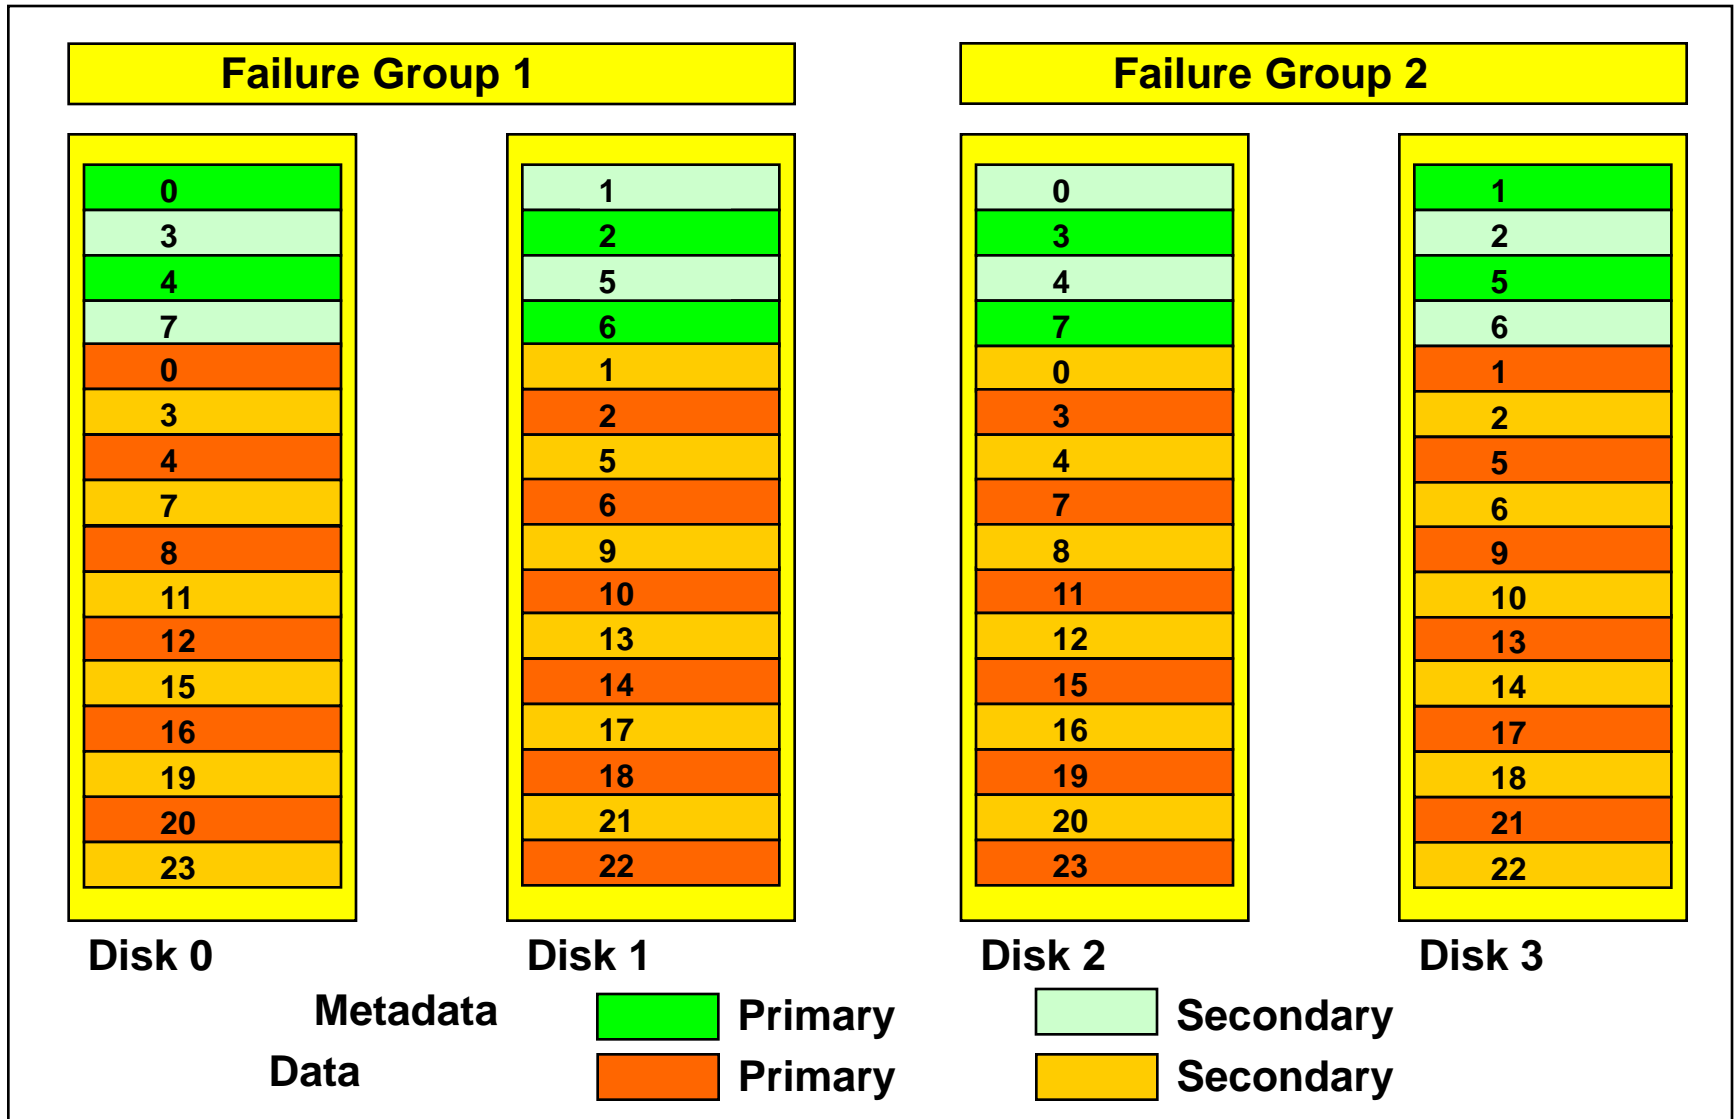
# ASM故障组-Normal Redundancy

每个故障组包含1块磁盘



# ASM故障组-Normal Redundancy

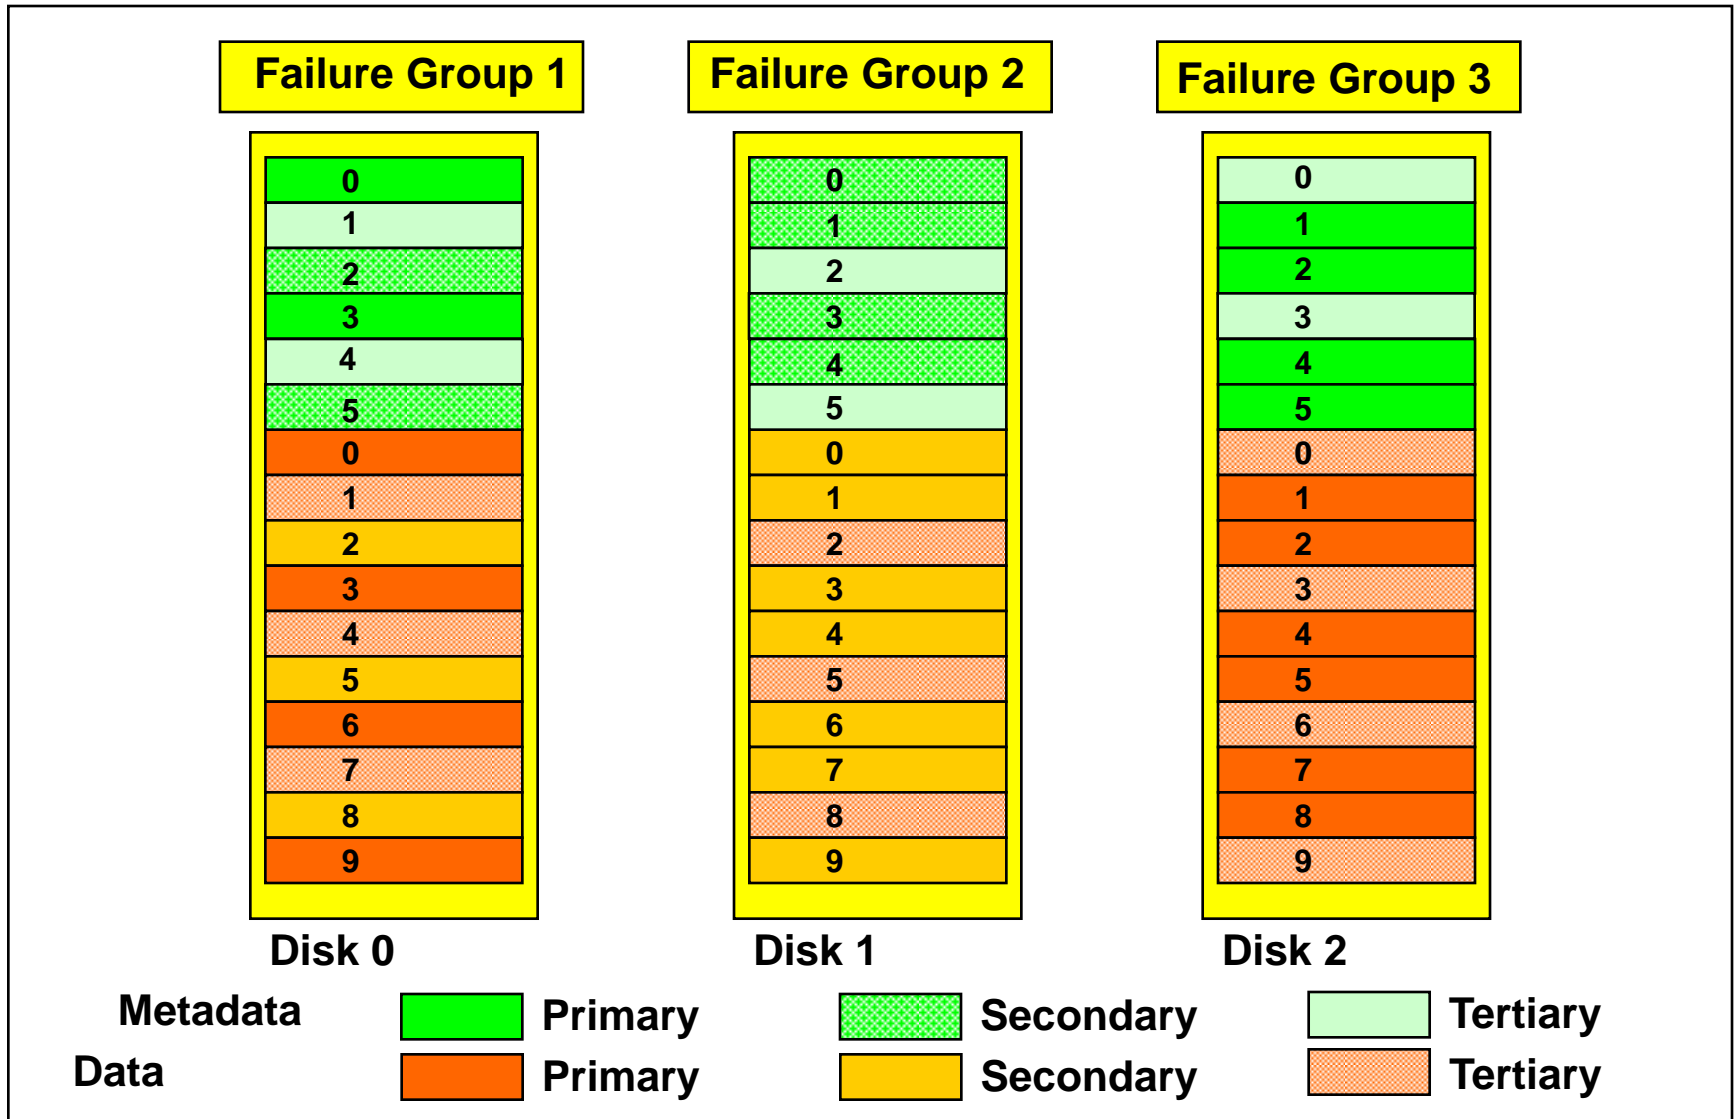
每个故障组包含2块磁盘








# ASM故障组-High Redundancy

每个故障组包含1块磁盘




# 伙伴的概念（Partnering）






# How ASM allocates redundant copies?


FG1



FG2

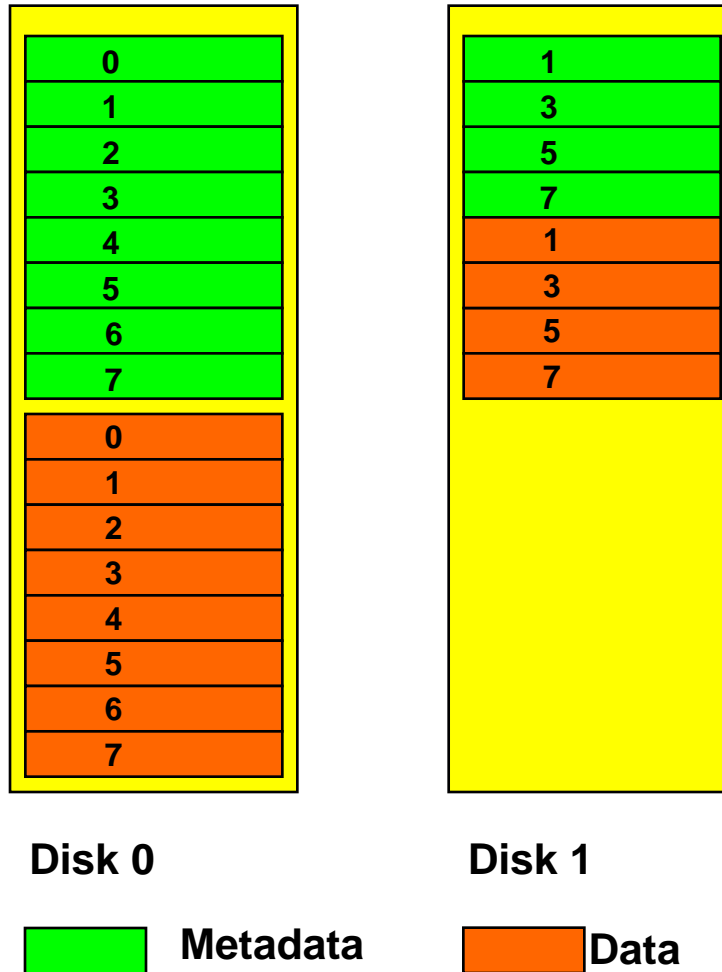


# ASM partnering concept



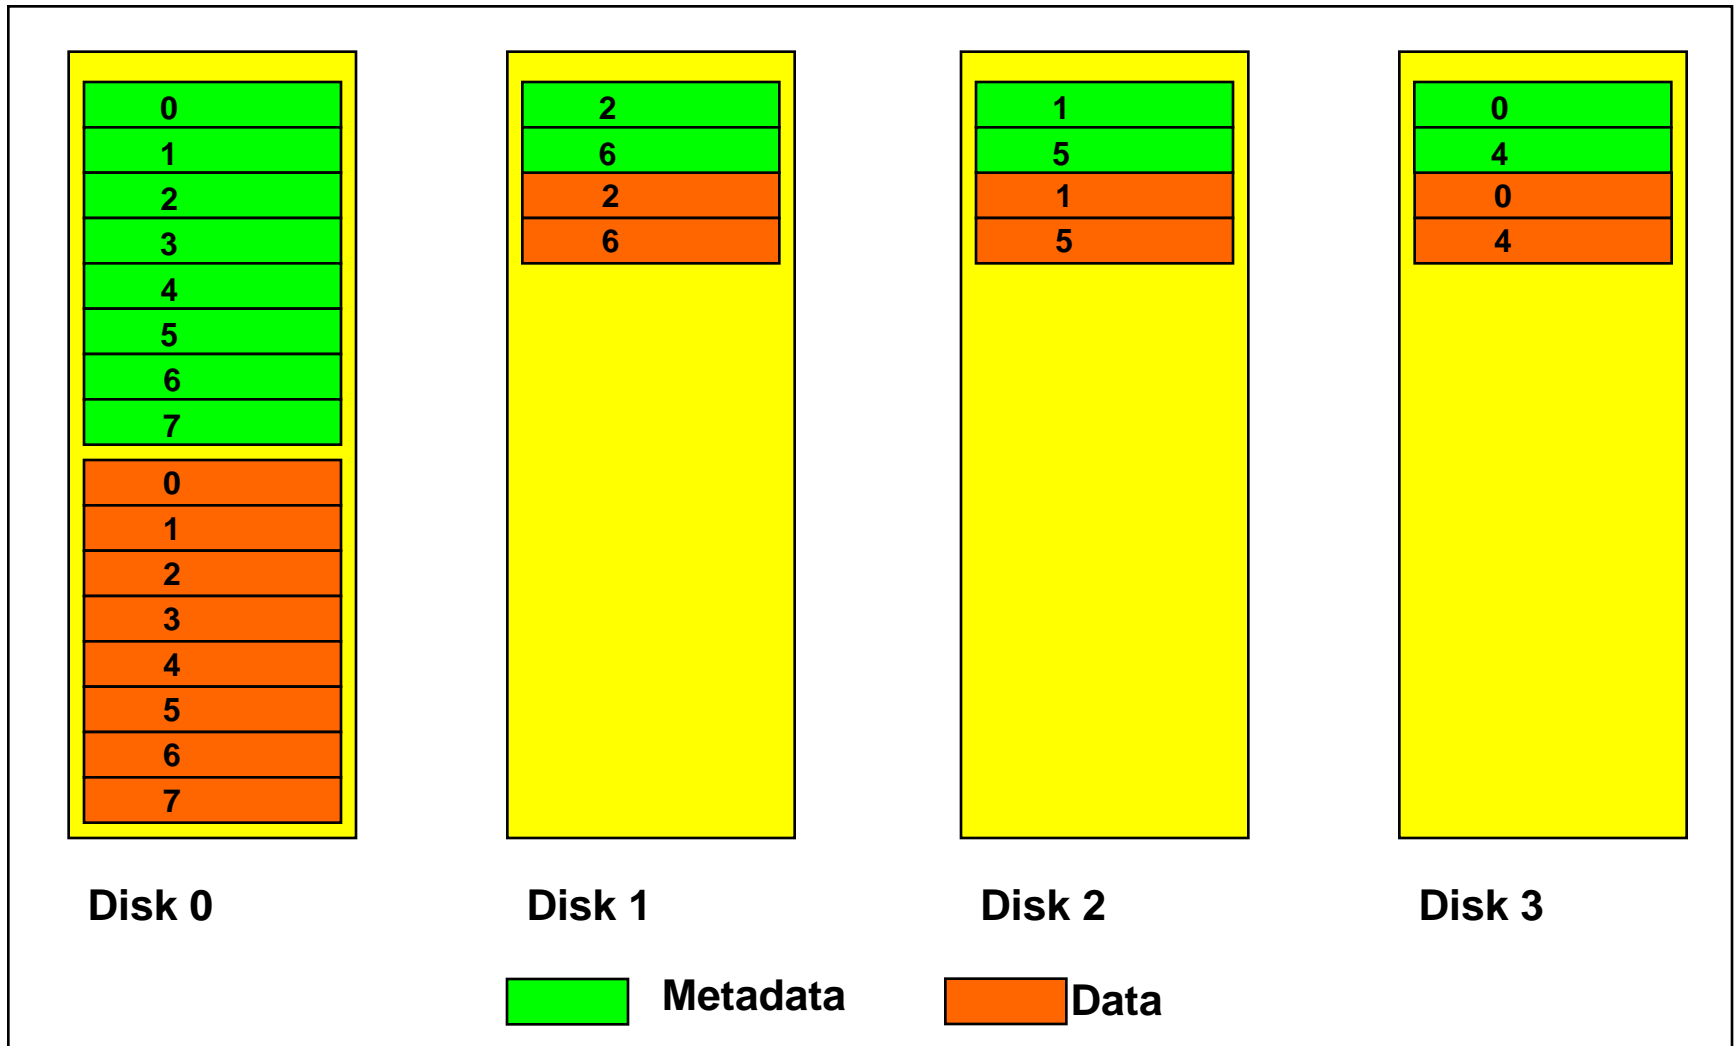
# ASM重平衡 (Rebalancing)

## Adding disks - 1 disks to 2 disks




# ASM重平衡 (Rebalancing)

## Adding disks - 1 disks to 4 disks



# ASM重平衡 (Rebalancing)

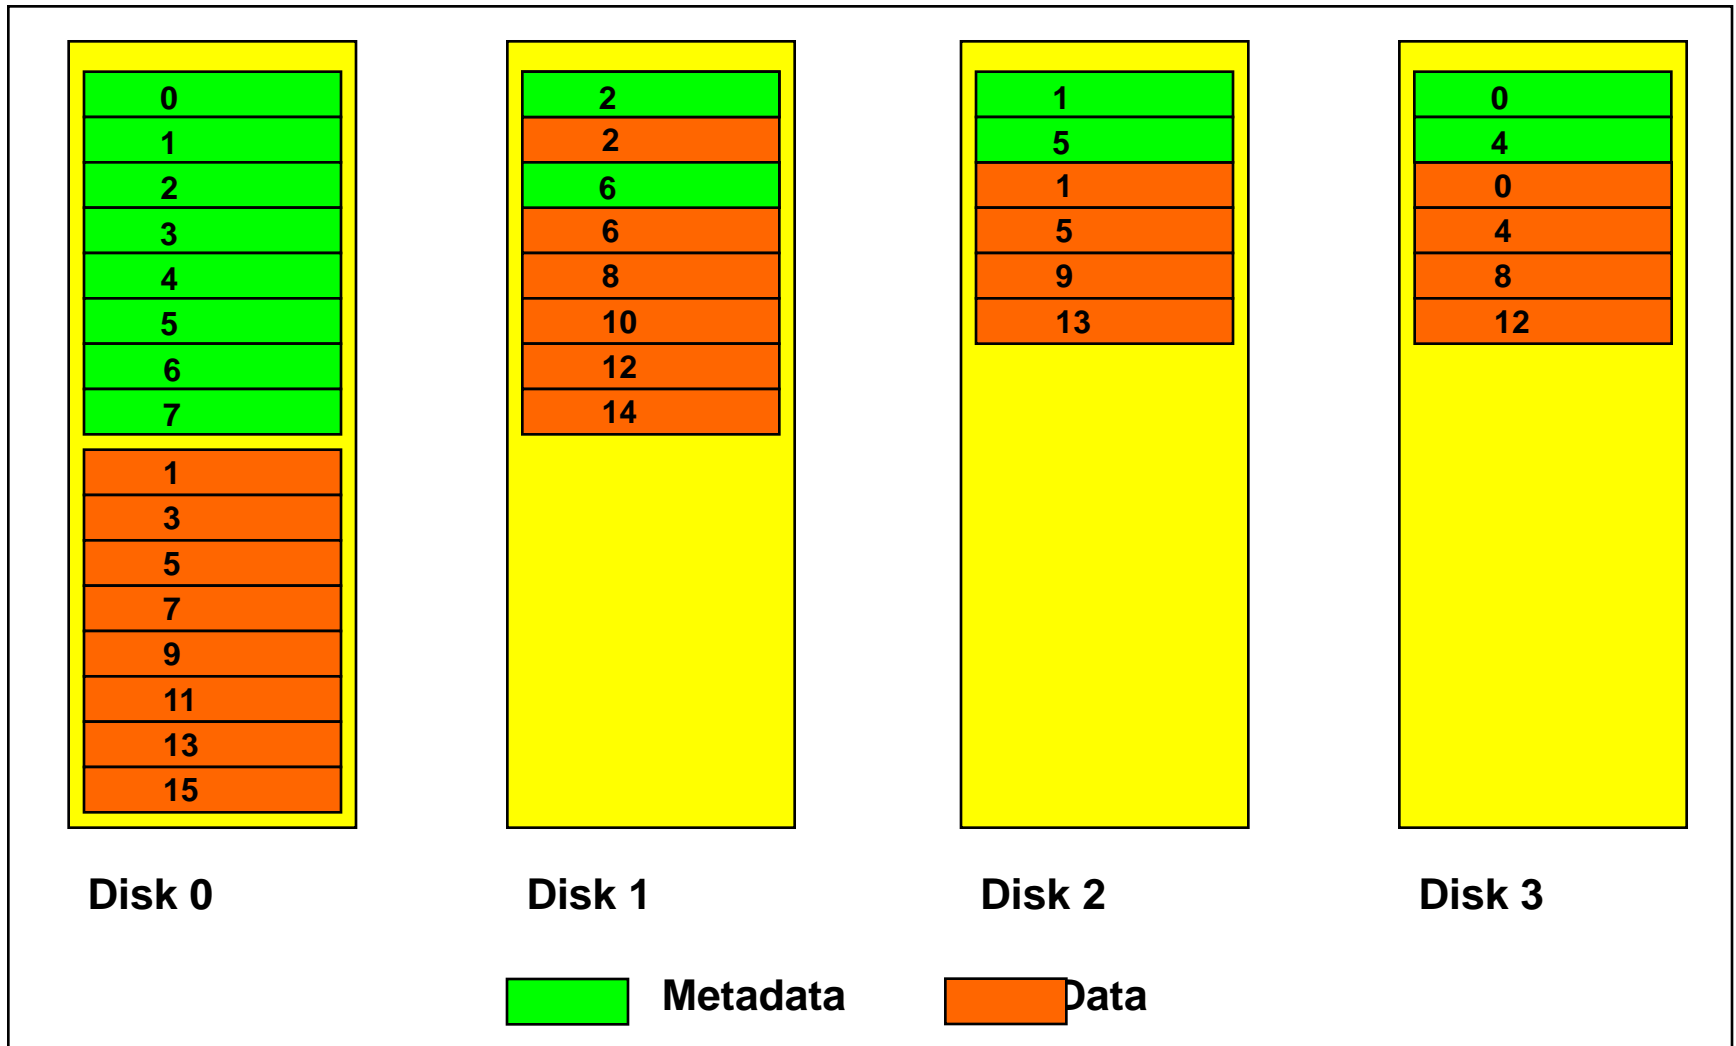
## Adding disks - 2 disks to 3 disks






# ASM重平衡 (Rebalancing)

## Adding disks - 2 disks to 4 disks




# ASM重平衡 (Rebalancing)

## Dropping disks - 3 disks to 1 disk



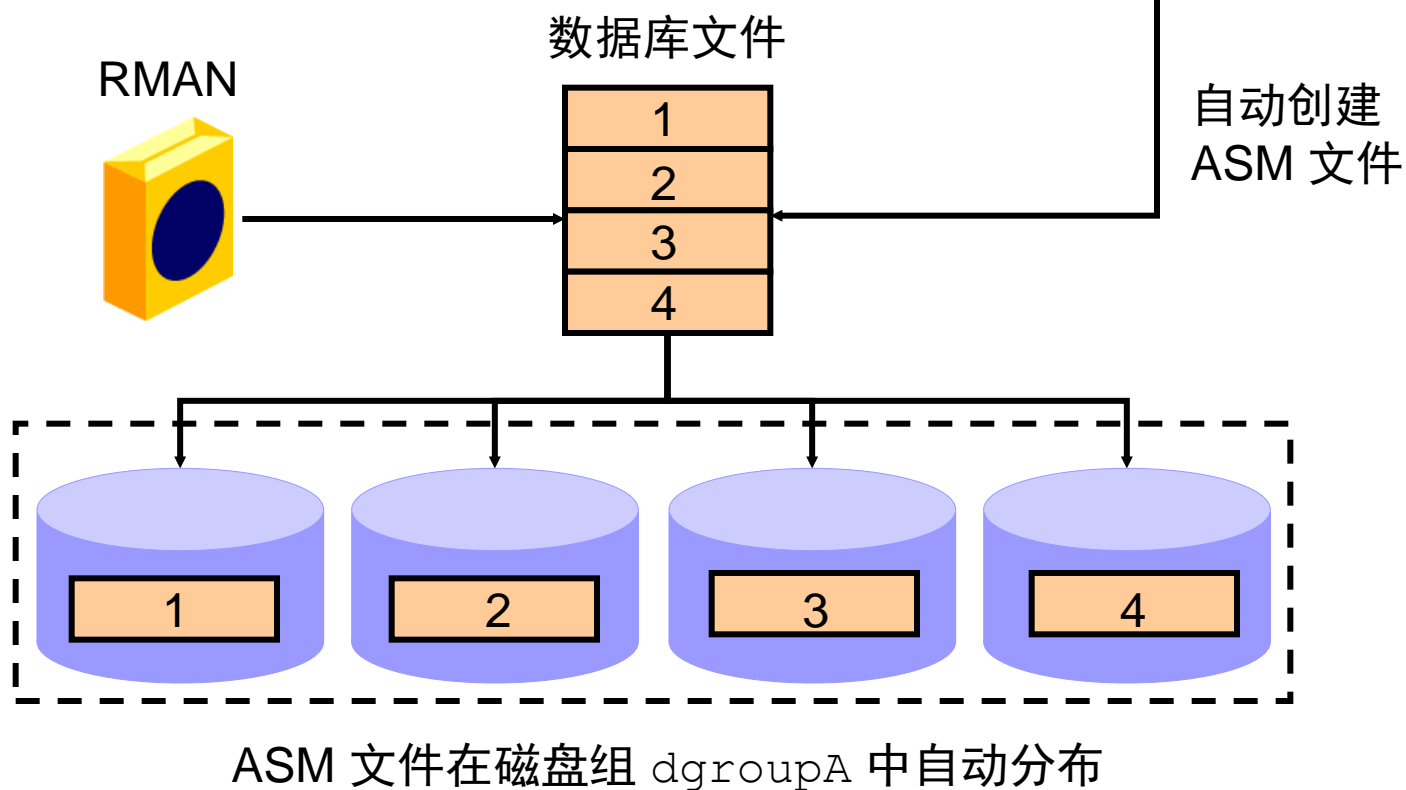
# ASM重平衡 (Rebalancing)

## Moving disks - 2 disks to 2 disks

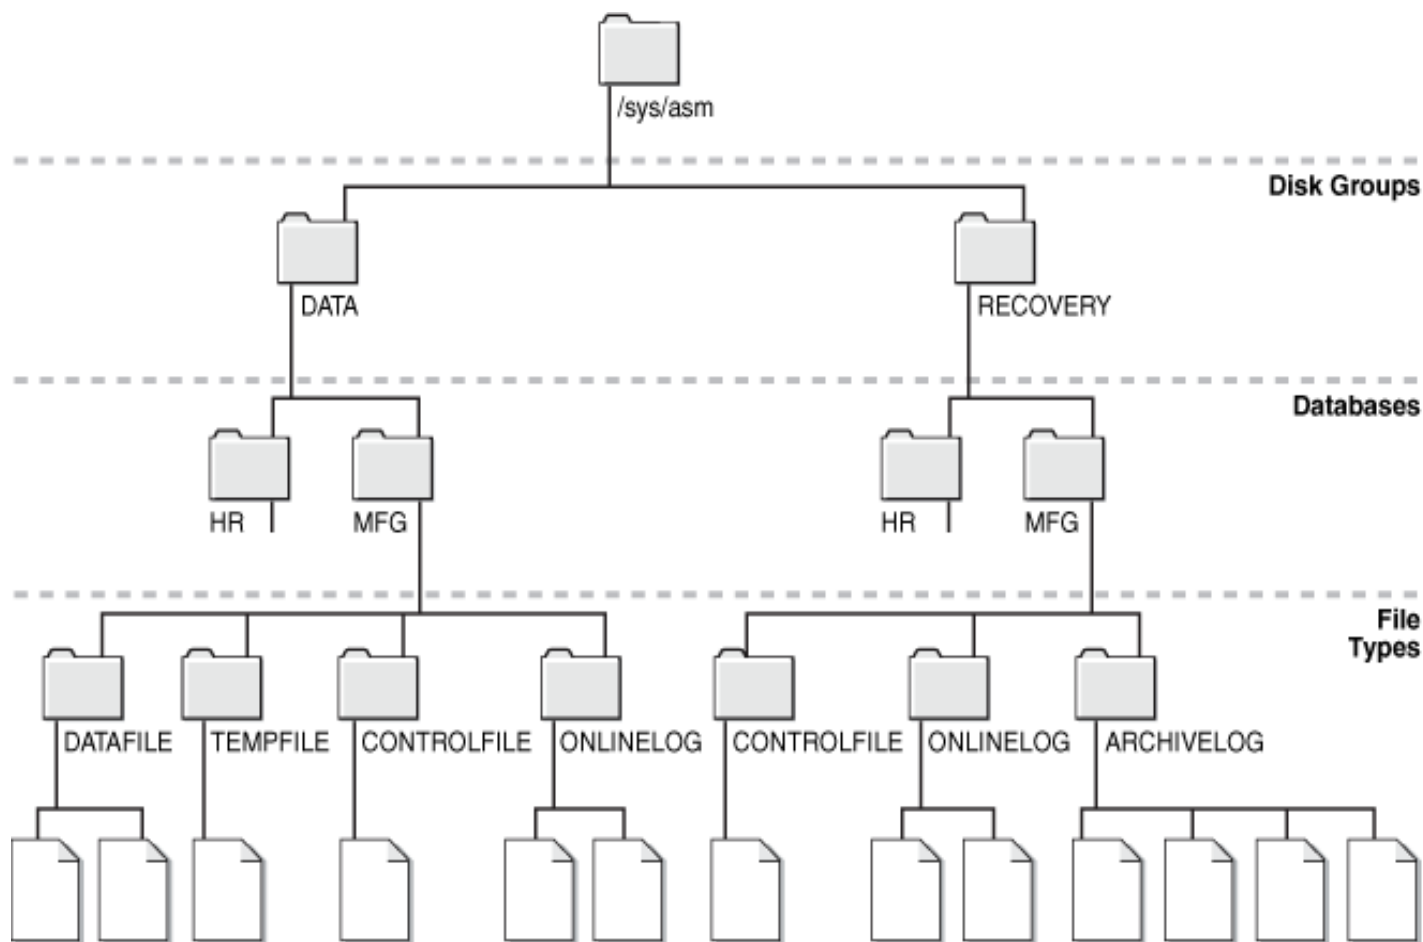


# ASM 文件

```
CREATE TABLESPACE sample DATAFILE '+dgroupA';
```



# ASM Disk directory tree



# X\$KFFXP

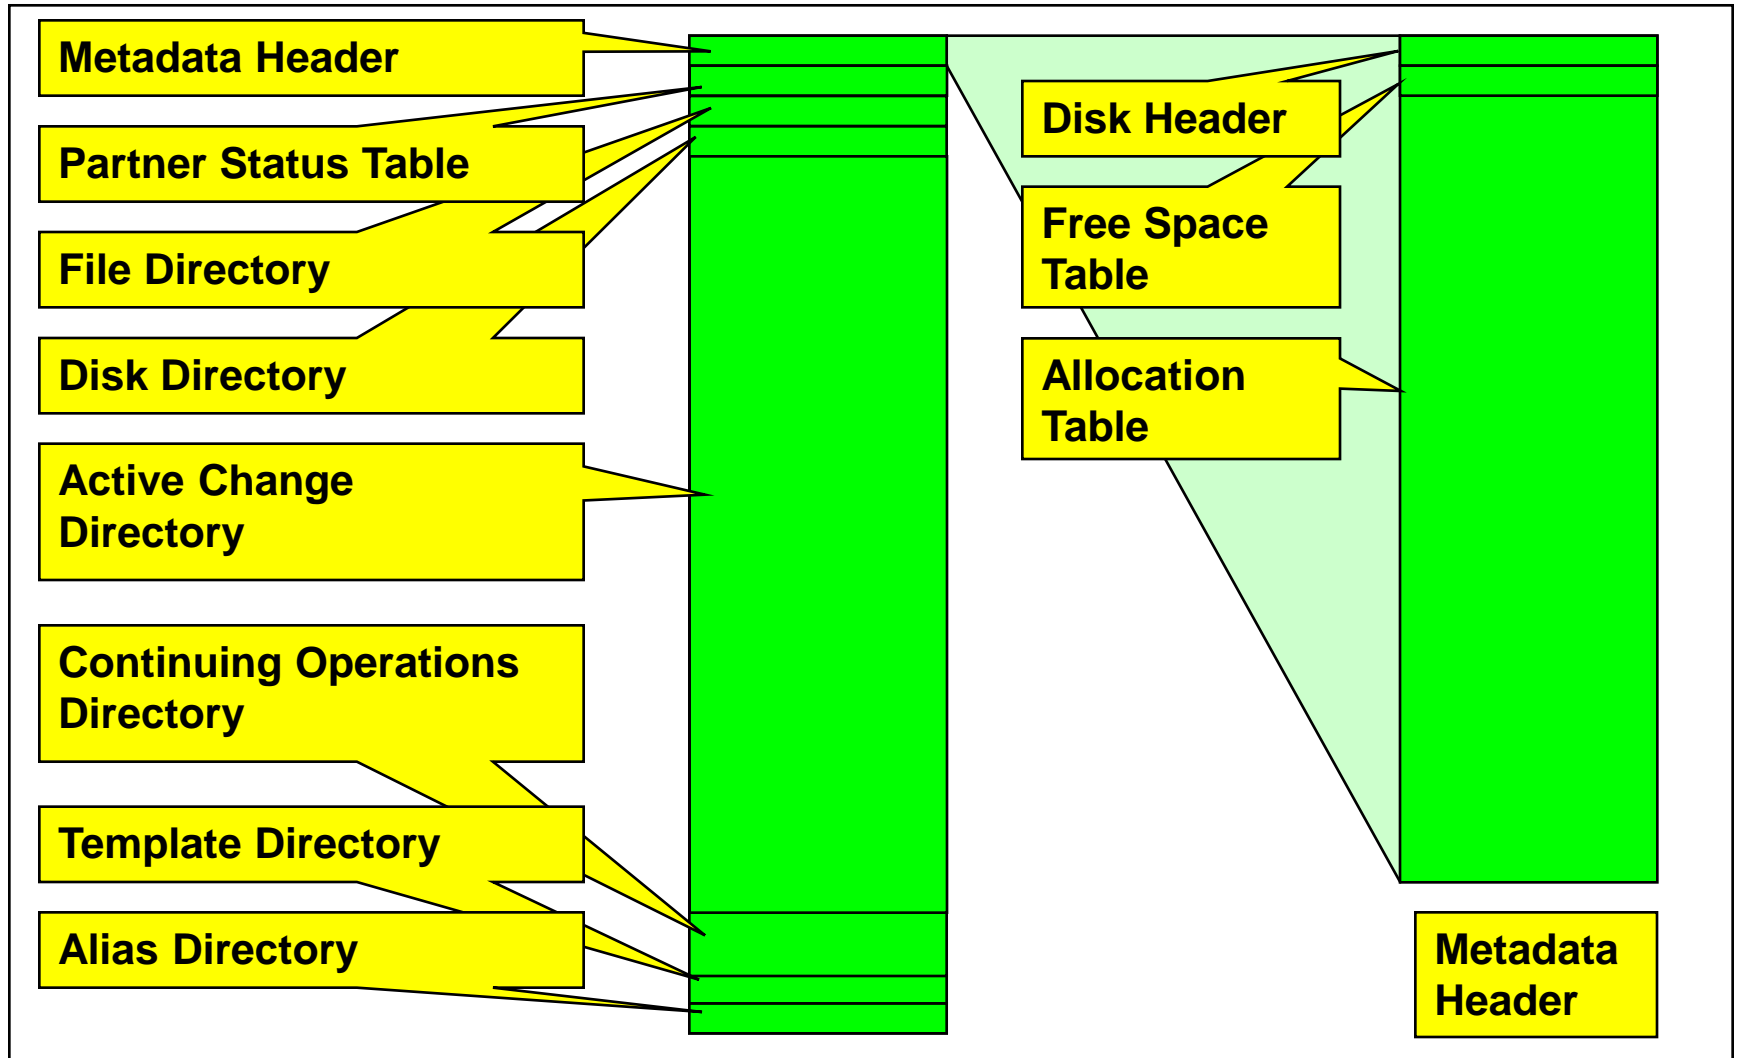
**Maps file extents to allocation units**

**Only populated in ASM instance**

**Columns include**

Column Name	Description
GROUP_KFFXP	Disk Group Number
NUMBER_KFFXP	File Number
COMPOUND_KFFXP	Disk Group Number    File Number
INCARN_KFFXP	Incarnation Number
PXN_KFFXP	Physical Extent Number (within file)
XNUM_KFFXP	Logical Extent Number (within file)
LXN_KFFXP	0=primary, 1=first mirror, 2=second mirror
DISK_KFFXP	Disk Number
AU_KFFXP	Allocation Unit Number (within disk)
SIZE_KFFXP	Size (# allocation units)

# ASM Metadata



# ASM Metadata

**Metadata is stored in first 256 files in ASM disk group**

- Space is initially allocated when disk group is created
- Can be subsequently extended

**Metadata allocation units are divided into blocks**

- Each block is 4096 bytes
- Block size specified using `_asm_blksize`

**Metadata files include**

File#	Description
0	Metadata Header
1	File Directory
2	Disk Directory
3	Active Change Directory
4	Continuing Operations Directory
5	Template Directory
6	Alias Directory
9	Attribute directory (optional)
12	Staleness registry (optional)



Block	0	1	2	3	.....				255	
AU0	1	2	3	3	3	3	3	3	3	1MB
AU1	17	17	13	18	13	13	13	13	19	
AU2	5	4	4	4	4	4	4	4	4	asmfile# 1
AU3	6	6	6	6	6	6	6	6	6	asmfile# 2
AU4	7	8	8	8	8	8	8	8	8	asmfile# 3
...	8	8	8	8	8	8	8	8	8	
AU24	12	12	12	12	12	12	12	12	12	Extend maps
...	8	8	8	8	8	8	8	8	8	asmfile# 3
AU50	D	A	T	A	F	I	L	E	S	~50MB
				...						

# kfed - Kernel File metadata Editor

11.1以前:

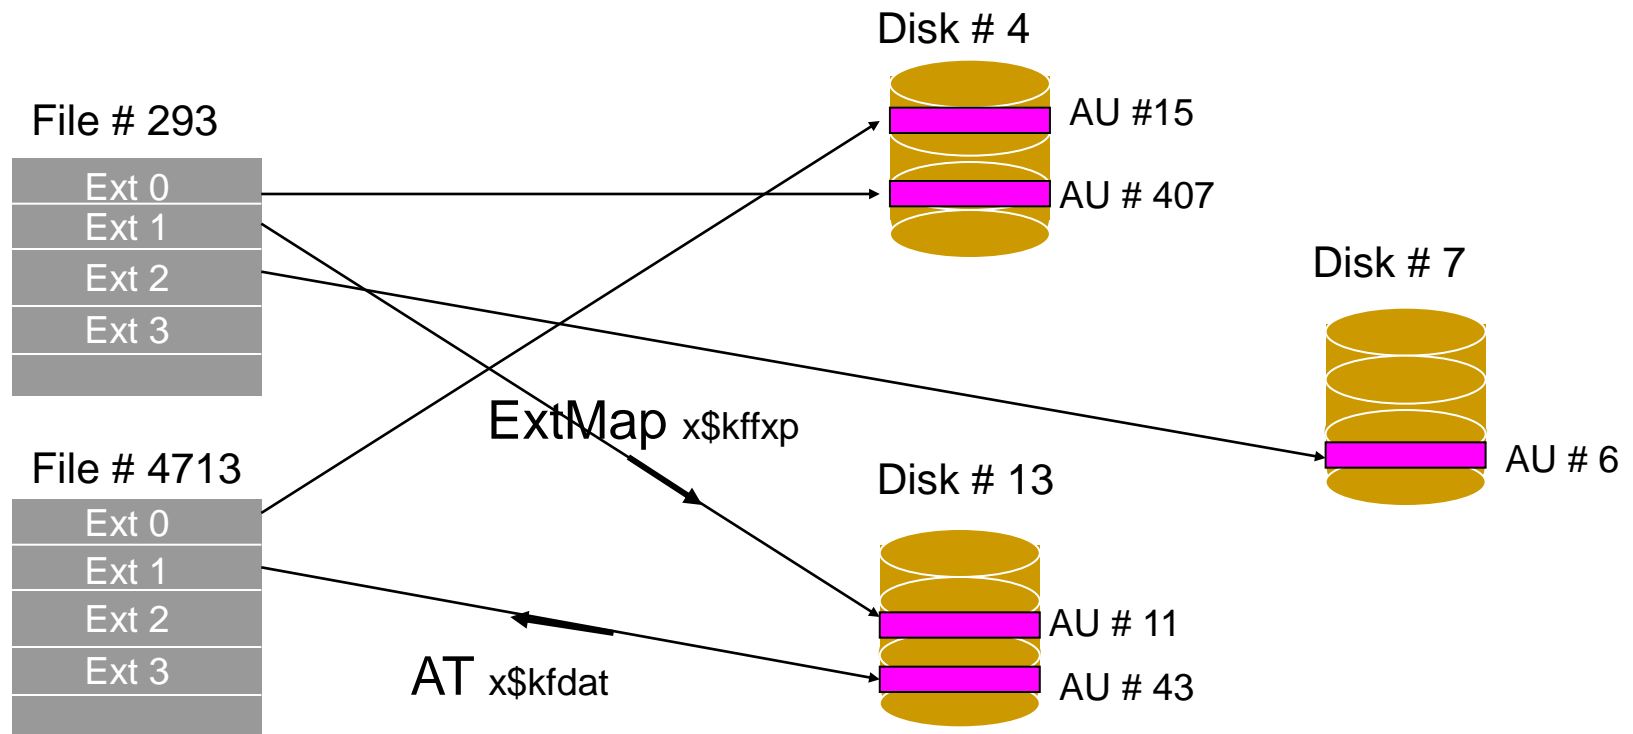
```
cd $ORACLE_HOME/rdbms/lib
modify ins_rdbms.mk
make -f ins_rdbms.mk ikfed
```

```
$ORACLE_HOME/bin/kfed -h
as/mlib          ASM Library [asmlib='lib']
aun/um           AU number to examine or update [AUNUM=number]
aus/z            Allocation Unit size in bytes [AUSZ=number]
blkn/um          Block number to examine or update [BLKNUM=number]
blks/z           Metadata block size in bytes [BLKSZ=number]
ch/ksum          Update checksum before each write [CHKSUM=YES/NO]
cn/t             Count of AUs to process [CNT=number]
d/ev             ASM device to examine or update [DEV=string]
o/p             KFED operation type
                [OP=READ/WRITE/MERGE/NEW/FORM/FIND/STRUCT]
p/rovnm          Name for provisioning purposes [PROVNM=string]
s/seek          AU number to seek to [SEEK=number]
te/xt            File name for translated block text [TEXT=string]
ty/pe           ASM metadata block type number [TYPE=number]
```

```
for (( f = 12544 ; f < 12544 + 256 ; f++ ))
do
    kfed op=read blkn=$f dev='/dev/oracleasm/disks/VOL1' > blk${f}
done
```

# ASM File Allocation

- ASM files分为多个小片，每一片叫做一个DE(Data extents)
- 每个ASM files都跨越在同一个磁盘组的多个磁盘上
- 将extents存储到磁盘上的位置叫做AU(Allocation Unit)



# X\$KFFXP Extent Maps

## ASMFILE# 1

one row per extent over all files

GROUP_KFFXP	diskgroup number (1 - 63)
NUMBER_KFFXP	file number for the extent
COMPOUND_KFFXP	$(\text{group\_kffxp} \ll 24) + \text{file \#}$
INCARN_KFFXP	file incarnation number
PXN_KFFXP	physical extent number
XNUM_KFFXP	extent number
<ul style="list-style-type: none"><li>• bit 31 set if indirect</li></ul>	
LXN_KFFXP	logical extent number
DISK_KFFXP	disk on which AU is located
AU_KFFXP	AU number on disk of AU

disk0

AU 86	E	X	T	0		
AU 87	E	X	T	3		

disk1

AU 86	E	X	T	1		

disk2

AU 85	E	X	T	2		

Datafile size 3MB => 4 AU's (4MB)

# X\$KFDAT – Allocation Tables

One row per AU over all disks

**GROUP\_KFDAT**      diskgroup # (1 - 63)

**NUMBER\_KFDAT**    disk # for the AU

**COMPOUND\_KFDAT** (group\_kfdat << 24) + number\_kfdat

**AUNUM\_KFDAT**      AU # within disk

**V\_KFDAT**            Valid: Y if allocated, N if unused

**FNUM\_KFDAT**        Meaningless unless V\_KFDAT set

- File number using this AU
- 0 for physical/PST AUs
- -1 if past end of disk

**I\_KFDAT**            Indirect: Meaningless unless V\_KFDAT set

- Y if indirect extent using the AU
- N if direct extent using this AU

**XNUM\_KFDAT**        Physical extent number using AU

**RAW\_KFDAT**         Raw 8-byte contents of extent pointer

- useful when V\_KFDAT not set to check for corruptions

# Metadata Block Layout

## e.g. Diskheader

```
kfbh.endian:      1 ; 0x000: 0x01
kfbh.hard:        130 ; 0x001: 0x82
kfbh.type:      1 ; 0x002: KFBTYP_DISKHEAD
kfbh.datfmt:      1 ; 0x003: 0x01
kfbh.block.blk:   0 ; 0x004: T=0 NUMB=0x0
kfbh.block.obj:   2147483648 ; 0x008: TYPE=0x8 NUMB=0x0
kfbh.check:       3768521839 ; 0x00c: 0xe09f146f
kfbh.fcn.base:    1696 ; 0x010: 0x0000006a0
kfbh.fcn.wrap:    0 ; 0x014: 0x000000000
kfbh.spare1:      0 ; 0x018: 0x000000000
kfbh.spare2:      0 ; 0x01c: 0x000000000
```

**32 byte  
BlockHeader**

**kfbh**

**kfbh**

## e.g. PST

```
kfbh.endian:      1 ; 0x000: 0x01
kfbh.hard:        130 ; 0x001: 0x82
kfbh.type:      17 ; 0x002: KFBTYP_PST_META
kfbh.datfmt:      1 ; 0x003: 0x01
kfbh.block.blk:   256 ; 0x004: T=0 NUMB=0x100
kfbh.block.obj:   2147483648 ; 0x008: TYPE=0x8 NUMB=0x0
kfbh.check:       2396453971 ; 0x00c: 0x8ed6fc53
kfbh.fcn.base:    0 ; 0x010: 0x000000000
kfbh.fcn.wrap:    0 ; 0x014: 0x000000000
kfbh.spare1:      0 ; 0x018: 0x000000000
kfbh.spare2:      0 ; 0x01c: 0x000000000
```

```
kfdhdb.driver.provstr: ORCLDISK ; 0x000: length=8
kfdhdb.driver.reserved[0]: 0 ; 0x008: 0x000000000
kfdhdb.driver.reserved[1]: 0 ; 0x00c: 0x000000000
kfdhdb.driver.reserved[2]: 0 ; 0x010: 0x000000000
kfdhdb.driver.reserved[3]: 0 ; 0x014: 0x000000000
kfdhdb.driver.reserved[4]: 0 ; 0x018: 0x000000000
kfdhdb.driver.reserved[5]: 0 ; 0x01c: 0x000000000
kfdhdb.compat:      168820736 ; 0x020: 0x0a100000
kfdhdb.dsknum:      0 ; 0x024: 0x0000
kfdhdb.grptyp:     2 ; 0x026: KFDGTP_NORMAL
kfdhdb.hdrsts:    3 ; 0x027: KFDHDR_MEMBER
kfdhdb.dskname:   ASMDG3_0000 ; 0x028: length=11
kfdhdb.grpname:   ASMDG3 ; 0x048: length=6
kfdhdb.fgname:   ASMDG3_0000 ; 0x068: length=11
kfdhdb.capname:     ; 0x088: length=0
```

**kfdhdb**

```
kfdpHdrB.time.hi:   32870224 ; 0x000: HOUR=0x10 DA
kfdpHdrB.time.lo:   238088192 ; 0x004: USEC=0x0 MSE
kfdpHdrB.last:      5 ; 0x008: 0x000000005
kfdpHdrB.next:      5 ; 0x00c: 0x000000005
kfdpHdrB.copyCnt:   2 ; 0x010: 0x02
kfdpHdrB.ub1spare:  0 ; 0x011: 0x00
kfdpHdrB.ub2spare:  0 ; 0x012: 0x0000
kfdpHdrB.incarn:    0 ; 0x014: 0x000000000
kfdpHdrB.copy[0]:   0 ; 0x018: 0x0000
kfdpHdrB.copy[1]:   1 ; 0x01a: 0x0001
kfdpHdrB.copy[2]:   2 ; 0x01c: 0x0002
kfdpHdrB.copy[3]:   0 ; 0x01e: 0x0000
kfdpHdrB.copy[4]:   0 ; 0x020: 0x0000
kfdpHdrB.dtaSz:     3 ; 0x022: 0x0003
ub1[0]:             2 ; 0x024: 0x02
```

# Agend

- ASM的历史和简介
- ASM体系架构
- ASM在不同版本的新特性
- ASM存储和分配机制
- **ASM日常使用注意事项**



# ASM实例管理的注意事项

- 在10.2的环境中，该ASM实例使用ASMM的方式管理，设置为100M足够了
- 从11.2开始，Oracle强烈推荐使用AMM方式管理ASM实例，并且最小值是MEMORY\_TARGET (256 MB)，最佳实践是1.5GB。大部分平台的缺省值是MEMORY\_TARGET (272 MB)。也可以采用传统模式：ALTER SYSTEM SET SGA\_TARGET=0
- 12c ( 12.1 ) 的文档中除了11.2的上述描述外，增加了一个内容：  
In an Oracle Exadata environment, the recommended settings for managing memory are SGA\_TARGET = 1250MB, PGA\_AGGREGATE\_TARGET = 400MB, MEMORY\_TARGET = 0, and MEMORY\_MAX\_TARGET = 0.

# 其他注意事项

- 1, 存储上的条带大小
- 2, ASM磁盘大小
- 3, 同一磁盘组中ASM磁盘大小和数目
- 4, 在ASM中添加ASMDISK (LUN或者其他)
- 5, 数据库维护 (RMAN备份, mv数据库文件等) 时, 如何避免ASM文件或者目录自动被删除的情况
- 6, 双存储使用ASM镜像

**Q & A**