

锻造凝练IT服务，助推用户事业发展
Perfecting IT service and favoring clients' success



CHINA
ETEK
SERVICE
PERFECTING IT SERVICE
AND FAVORING
CLIENTS' SUCCESS

技术人生系列背后的故事



北京·上海·西安·广州·南京·郑州·杭州·青岛·深圳·沈阳·昆明

About Me

- 周永康（老K）
- 11g OCP, 11gOCM
- 多年的总行数据中心Oracle运维经验
- 擅长SQL优化
- 擅长Oracle故障诊断
- 擅长Oracle性能优化
- 《技术人生系列微信文章》主编
- 中亦安图科技股份有限公司

从技术人生系列看ORACLE技术进阶之路

- ORACLE的视图进阶
- ORACLE的EVENT的进阶
- ORACLE 的dump 方法使用进阶
- ORACLE相关工具的进阶
- ORACLE的trace分析的进阶

ORACLE的视图使用进阶

- 基本的视图

常设置的事件：v\$session/v\$lock/v\$transaction

- 视图查询的进阶

v\$sesstat/v\$sysstat

v\$active_session_history

dba_hist_active_sess_history

x\$kgllk

x\$bh

ORACLE的event使用进阶

- 基本的event

常设置的事件：10046/10053/10231

- **Event**事件的进阶

这些event你会设置吗：

errorstack

10949

28041

ORACLE的dump使用进阶

- 基本的**dump**方法

常用的dump方法：

Dump block

Dump systemstate

Dump processstate

- **Dump**方法的进阶

Dump heapdump 536870917---dump pga

Dump heapdump 536870914---dump shared pool

其他内存dump的方法

ORACLE的分析工具使用进阶

- 基本的工具方法

常用的工具方法：

Awrrpt

Awrsqrpt

Nmon

Oswatcher

- 分析工具的进阶

gdb/xdb

truss/Strace

svmon

ORACLE的trace分析的进阶

- 基本的**trace**分析
10053/10046的**trace**分析
Alert日志/lgwr的**trace**分析
- **Trace**分析的进阶
600错误的**trace**分析
7445错误的**trace**分析
4031产生的**trace**分析
各种手动产生的**trace**的分析

两个简单的案例分析

- 一次数据库实例性能分析
- 一条高**SQL**的优化

一次数据库实例性能分析

- 背景：
 - 一套**ORACLE RAC**实例性能出现严重问题
 - 客户现场分析，主要等待事件与**GC**相关
 - 客户现场解决方案，关闭另一节点，留下其中一个节点，问题解决
- 客户的问题
 - 为什么
 - 我要怎么做

一次数据库实例性能分析

- 客户给的信息
 - 早上节点一的网卡好像出了一些问题（是不是网络有了问题）
 - 节点二的内存比节点一的内存要小一些（所以选择停了节点二）
- 自己收集信息
 - 动态性能视图
 - **v\$active_session_history/dba_hist_active_sess_history**
 - **Alert**日志/及其他**trace**文件

- 多维度视图分析

```
select sample_time,event,count(*) from v$active_session_history where
SAMPLE_TIME sample_time>to_date('20170606 11:08','yyyymmdd hh24:mi') and
sample_time<to_date('20170606 11:09','yyyymmdd hh24:mi')
-----
06-JUN-17 11.08.17.359 group by sample_time,event order by 1,3;
```

北京·上海·西安·广州·南京·郑州·杭州·青岛·深圳·沈阳·昆明

一次数据库实例性能分析

- Alert日志的信息
 - 二节点性能出现了问题

```
Thread 2 advanced to log sequence 19574 (LGWR switch)
Current log# 11 seq# 19574 mem# 0: +ARCHIVE/online/redo11.log
Tue Jun 06 10:25:52 2017
LNS: Standby redo logfile selected for thread 2 sequence 19574 for destination LOG_ARCHIVE_DEST_2
Tue Jun 06 10:25:53 2017
LNS: Standby redo logfile selected for thread 2 sequence 19574 for destination LOG_ARCHIVE_DEST_3
Tue Jun 06 10:26:16 2017
Archived Log entry 122015 added for thread 2 sequence 19573 ID 0xffffffff0412add4 dest 1:
Tue Jun 06 10:53:16 2017

*****
Tue Jun 06 10:57:27 2017

Fatal NI connect error 12537, connecting to:
  (LOCAL=NO)
Tue Jun 06 10:59:19 2017
Tue Jun 06 10:59:38 2017
Errors in file /db/diag/rdbms/.../pd2/trace/..._15532226.trc (incident=5962100):
ORA-00445: background process "M000" not start after 120 seconds
Time drift detected. Please check VKTM trace file for more details.

VERSION INFORMATION:
  TNS for IBM/AIX RISC System/6000: Version 11.2.0.2.0 - Production
  TCP/IP NT Protocol Adapter for IBM/AIX RISC System/6000: Version 11.2.0.2.0 - Production
  Oracle Bequeath NT Protocol Adapter for IBM/AIX RISC System/6000: Version 11.2.0.2.0 - Production
Tue Jun 06 10:57:13 2017
Errors in file /db/diag/rdbms/.../pd2/trace/..._32964788.trc (incident=5961692):
ORA-00445: background process "M000" not start after 120 seconds
Tue Jun 06 11:00:23 2017
Tue Jun 06 11:00:21 2017
DIAG (ospid: 10682508) waits for event 'DIAG idle wait' for 0 secs. Time: 06-JUN-2017 11:00:18
Incident details in: /db/diag/rdbms/.../pd2/incident/incdir_5962100/..._d2_cjq0_15532226_i5962100.trc

Incident details in: /db/diag/rdbms/.../pd2/incident/incdir_5961692/..._d2_rmon_32964788_i5961692.trc
Tue Jun 06 11:01:33 2017
Tracing not turned on.
Tns error struct:
  ns main err code: 12537
```

一次数据库实例性能分析

- 隐藏的信息
 - 为啥要重启？
 - 网卡可能存在问题，在开门前赶紧重启排查一下
- 结合前面的特征大概知道原因了吗

```
Fri Jun 06 08:42:40 2017
Shutting down instance (abort)
License high water mark = 541
USER (ospid: 12779808): terminating the instance
Fri Jun 06 08:42:40 2017
opiodr aborting process unknown ospid (19661218) as a result of ORA-1092
Fri Jun 06 08:42:40 2017
ORA-1092 : opitsk aborting process
Fri Jun 06 08:42:40 2017
opiodr aborting process unknown ospid (56033574) as a result of ORA-1092
Fri Jun 06 08:42:40 2017
opiodr aborting process unknown ospid (18153536) as a result of ORA-1092
Instance terminated by USER, pid = 12779808
Fri Jun 06 08:42:56 2017
Instance shutdown complete
Tue Jun 06 09:27:30 2017
Starting ORACLE instance (normal)
sskgpgetexename failed to get name
LICENSE_MAX_SESSION = 0
LICENSE_SESSIONS_WARNING = 0
Private Interface 'en2' configured from GPnP for use as a private interconn
[Name='en2', type=1, ip=169.254.171.201, mac=5c-f3-fc-04-0d-d0, net=169.25
Public Interface 'en4' configured from GPnP for use as a public interface.
[Name='en4', type=1, ip=192.168.154.198, mac=5c-f3-fc-04-0d-c0, net=192.16
Public Interface 'en4' configured from GPnP for use as a public interface.
[Name='en4', type=1, ip=192.168.154.208, mac=5c-f3-fc-04-0d-c0, net=192.16
Picked latch-free SCN scheme 3
Autotune of undo retention is turned on.
LICENSE_MAX_USERS = 0
SYS auditing is disabled
Starting up:
Oracle Database 11g Enterprise Edition Release 11.2.0.2.0 - 64bit Production
With the Partitioning, Real Application Clusters, OLAP, Data Mining
and Real Application Testing options.
Using parameter settings in server-side pfile /db/product/11.2.0.2/dbhome_1/
System parameters with non-default values:
processes = 5000
sessions = 7552
event = "10298 trace name context forever, level 32"
sga_max_size = 16960M
```

一次数据库实例性能分析

- 验证之道
 - 没有oswatcher怎么办？

8196Time: 11:04:13 -----

8197

8198CONFIGCPUMEMORYPAGING

8199ModeDedKern3.4Sz,GB30.5Sz,GB32.5

8200LP32.0User11.5InU30.5InU9.1

8201SMTONWait27.5%Comp99.5Flt13652

8202Ent0.0Idle57.7%NonC0.4Pg-I423

8203Poolid-PhyBS118.8%Clnt0.4Pg-O3498

8204

8205

8206PHYPEVENTS/QUEUESNFS

8207Bdon0.0Cswth25039SrvV20

8208Idon0.0Syscl58709ClntV20

Find result - 155 hits

Line6878: SMTONWait9.1%Comp47.9Flt20084

Line6941: SMTONWait9.5%Comp81.3Flt12759

Line7004: SMTONWait11.4%Comp95.7Flt9898

Line7067: SMTONWait12.4%Comp98.9Flt11993

Line7130: SMTONWait13.9%Comp98.2Flt10908

Line7193: SMTONWait16.1%Comp98.4Flt9616

Line7256: SMTONWait14.8%Comp99.0Flt10056

Line7319: SMTONWait14.6%Comp98.5Flt12477

Line7382: SMTONWait16.8%Comp98.3Flt12783

Line7445: SMTONWait15.6%Comp98.7Flt11744

Line7508: SMTONWait14.7%Comp98.7Flt13559

Line7571: SMTONWait14.9%Comp98.8Flt13896

Line7634: SMTONWait16.2%Comp98.9Flt14434

Line7697: SMTONWait16.3%Comp98.8Flt13912

Line7760: SMTONWait15.9%Comp98.9Flt14382

Line7823: SMTONWait14.7%Comp99.0Flt13945

Line7886: SMTONWait15.5%Comp98.9Flt12915

Line7949: SMTONWait14.6%Comp98.8Flt19194

Line8012: SMTONWait16.1%Comp98.9Flt13066

Line8075: SMTONWait18.0%Comp98.9Flt14799

Line8138: SMTONWait19.2%Comp99.0Flt15048

Line8201: SMTONWait27.5%Comp99.5Flt13652

Line8264: SMTONWait26.5%Comp99.7Flt13190

Line8327: SMTONWait27.4%Comp99.8Flt15534

Line8390: SMTONWait23.0%Comp99.8Flt12570

一次数据库实例性能分析

- 验证之道

```
bash-3.00# svmon -P 9503010
```

Pid	Command	Inuse	Pin	Pgsp	Virtual	64-bit	Mthrd	16MB
9503010	oracle	4178483	27008	0	4155208	Y	N	N

	PageSize	Inuse	Pin	Pgsp	Virtual			
s	4 KB	36707	0	0	13432			
m	64 KB	258861	1688	0	258861			

50005	9ffffffd work shared library			sm	2818	0	0	2818
e988e9	80020014 work USLA heap			sm	1899	0	0	1899
e688e6	11 work text data BSS heap			sm	1793	0	0	1793

Bug 13443029 AIX: Excess "work USLA heap" process memory use in 11.2 on AIX

SYMPTOMS

Dedicated server processes using significantly more mem

Using svmon to monitor process memory segment after ini
7.0M in 11gR2.

svmon -P PID - where PID is an Oracle process id

Oracle Release -> (work USLA heap times 4k pages size)

11.2.0.1.0 -> 7M bytes

11.1.0.7.0 -> 60KB

10.2.0.4.0 -> 420KB

一次数据库实例性能分析

- 建议方案
 - 调整连接池
 - 打上相应的补丁
 - 调小**SGA**

Time: 09:21:39 -----

CONFIG		CPU		MEMORY		PAGING	
Mode	Ded	Kern	4.6	Sz,GB	46.4	Sz,GB	32.5
LP	32.0	User	13.0	InU	46.3	InU	1.4
SMT	ON	Wait	15.2	%Comp	88.5	Flt	12217
Ent	0.0	Idle	67.1	%NonC	11.3	Pg-I	0
Poolid	-	PhyB	141.4	%Clnt	11.3	Pg-O	68

Time: 09:21:55 -----

CONFIG		CPU		MEMORY		PAGING	
Mode	Ded	Kern	3.8	Sz,GB	30.5	Sz,GB	32.5
LP	32.0	User	11.4	InU	30.4	InU	4.7
SMT	ON	Wait	14.9	%Comp	94.8	Flt	10228
Ent	0.0	Idle	69.9	%NonC	4.7	Pg-I	0
Poolid	-	PhyB	121.7	%Clnt	4.7	Pg-O	2
		Entc	0.0				

一条SQL的优化

- 首先来简单的对比一下

```
SQL> select unit_pkgeid
2   from tmp_nodatelink_0613_2 vc
3  inner join tmp_calendar5 tc
4      on
5     vc.cf_dates <= tc.cf_dates
6    and vc.end_date > tc.cf_dates
7    and
8   tc.dates_type = vc.dates_type
9  order by unit_name,currency,tc.cf_dates;
```

304227 rows selected.

Elapsed: 00:00:04.32

Execution Plan

Plan hash value: 1923048139

Id	Operation	Name	Rows	Bytes	(TempSpc)	Cost	(%CPU)	Time
0	SELECT STATEMENT		10746	912K		447	(24)	00:00:06
1	SORT ORDER BY		10746	912K	1112K	447	(24)	00:00:06
* 2	HASH JOIN		10746	912K		227	(45)	00:00:03
3	TABLE ACCESS FULL	tmp_calendar5	2152	27976		5	(0)	00:00:01
4	TABLE ACCESS FULL	tmp_nodatelink_0613_2	19237	1390K		122	(1)	00:00:02

Predicate Information (identified by operation id):

```
2 - access("TC"."DATES_TYPE"="VC"."DATES_TYPE")
   filter("VC"."CF_DATES"<="TC"."CF_DATES" AND "VC"."END_DATE">"TC"."CF_DATES")
```

Statistics

```
25 recursive calls
0 db block gets
451 consistent gets
0 physical reads
0 redo size
```

```
SQL> select unit_pkgeid
2   from tmp_nodatelink_0613_2 vc
3  inner join tmp_calendar4 tc
4      on
5     vc.cf_dates <= tc.cf_dates
6    and vc.end_date > tc.cf_dates
7    and
8   tc.dates_type = vc.dates_type
9  order by unit_name,currency,tc.cf_dates;
```

304227 rows selected.

Elapsed: 00:00:01.32

Execution Plan

Plan hash value: 1255175667

Id	Operation	Name	Rows	Bytes	(TempSpc)	Cost	(%CPU)	Time
0	SELECT STATEMENT		10746	912K		447	(24)	00:00:06
1	SORT ORDER BY		10746	912K	1112K	447	(24)	00:00:06
* 2	HASH JOIN		10746	912K		227	(45)	00:00:03
3	TABLE ACCESS FULL	tmp_calendar4	2152	27976		5	(0)	00:00:01
4	TABLE ACCESS FULL	tmp_nodatelink_0613_2	19237	1390K		122	(1)	00:00:02

Predicate Information (identified by operation id):

```
2 - access("TC"."DATES_TYPE"="VC"."DATES_TYPE")
   filter("VC"."CF_DATES"<="TC"."CF_DATES" AND "VC"."END_DATE">"TC"."CF_DATES")
```

Statistics

```
8 recursive calls
0 db block gets
451 consistent gets
0 physical reads
0 redo size
```

一条SQL的优化

优化的起因

Host Name	Platform	CPU	Cores	Sockets	Memory (GB)
fb1	AIX-Based Systems (64-bit)	16	4		24.00

Snap Id	Snap Time	Sessions	Cursors/Session	Instances
Begin Snap: 1437	27-May-17 14:00:08	56	2.2	2
End Snap: 1438	27-May-17 14:52:05	106	2.5	2
Elapsed:	51.95 (mins)			
DB Time:	925.28 (mins)			

Top 10 Foreground Events by Total Wait Time

Event	Waits	Total Wait Time (sec)	Wait Avg(ms)	% DB time	Wait Class
resmgr:cpu quantum	393,104	30K	76	54.1	Scheduler
DB CPU		5017.1		9.0	

SQL ordered by CPU Time

- Resources reported for PL/SQL code includes the resources used by all SQL statements called by the code.
- %Total - CPU Time as a percentage of Total DB CPU
- %CPU - CPU Time as a percentage of Elapsed Time
- %IO - User I/O Time as a percentage of Elapsed Time
- Captured SQL account for 96.5% of Total CPU Time (s): 5,017
- Captured PL/SQL account for 0.2% of Total CPU Time (s): 5,017

CPU Time (s)	Executions	CPU per Exec (s)	%Total	Elapsed Time (s)	%CPU	%IO	SQL Id	SQL Module	SQL Text
2,294.42	417	5.50	45.73	27,088.64	8.47	0.07	ffk94w1wmj360	JDBC Thin Client	SELECT D2.* FROM (SELECT A.*, ...
2,265.46	442	5.13	45.15	26,505.12	8.55	0.07	fj634umh8g7a7	JDBC Thin Client	SELECT COUNT(*) FROM vrep_cash...
81.72	33	2.48	1.63	143.19	57.07	0.00	7b5au4wny3y9a	JDBC Thin Client	SELECT e.* FROM (SELECT s.*, ...
80.99	33	2.45	1.61	140.84	57.50	0.00	6ksgg2bs92yqg	JDBC Thin Client	select count(1) from (WITH PR...

- 原语句是这样的

北京·上海·西安·广州·南京·郑州·杭州·青岛·深圳·沈阳·昆明

一条SQL的优化

- 执行计划是这样的

Execution Plan

Plan hash value: 1329878397

Id	Operation	Name	
0	SELECT STATEMENT	*190	CONNECT BY WITHOUT FILTERING
* 1	VIEW	191	FAST DUAL
* 2	WINDOW NOSORT STOPKEY	192	LOAD AS SELECT
* 3	FILTER	193	HASH GROUP BY
* 4	VIEW	194	TABLE ACCESS FULL
5	SORT ORDER BY	195	LOAD AS SELECT
6	WINDOW SORT	196	HASH GROUP BY
7	MERGE JOIN	197	TABLE ACCESS FULL
8	VIEW	198	LOAD AS SELECT
9	WINDOW SORT	199	HASH GROUP BY
* 10	HASH JOIN RIGHT OUTER	200	TABLE ACCESS FULL
11	VIEW	201	VIEW
12	HASH GROUP BY	*202	TABLE ACCESS FULL
* 13	VIEW	203	NESTED LOOPS
* 14	HASH JOIN	204	NESTED LOOPS
15	TABLE ACCESS FULL	205	MERGE JOIN CARTESIAN
* 16	TABLE ACCESS FULL	*206	TABLE ACCESS BY INDEX ROWID
17	TABLE ACCESS FULL	*207	INDEX RANGE SCAN
18	VIEW	208	BUFFER SORT
* 19	FILTER	*209	INDEX RANGE SCAN
20	SORT GROUP BY ROLLUP	*210	INDEX RANGE SCAN
21	VIEW	*211	TABLE ACCESS BY INDEX ROWID

一条SQL的优化

- 手动执行统计信息是这样的

Statistics

```
-----
      4  recursive calls
     46  db block gets
69947  consistent gets
    2968  physical reads
    2128  redo size
   1514  bytes sent via SQL*Net to client
    513  bytes received via SQL*Net from client
      1  SQL*Net roundtrips to/from client
      6  sorts (memory)
      0  sorts (disk)
      0  rows processed
```

- 特征是什么？
 - 使用了视图/分页排序
 - 逻辑读并不大（**69947**个逻辑读）
 - 执行时间较长

```
> oradebug short_stack
ksedsts()+360<-ksdxfstk()+44<-ksdxcdb()+3384<-sspuser()+116<-47dc<-expepr()+100<-evaor()+88<-expepr()+100<-evacssr()+168<-qerghRow
-----
> oradebug short_stack
ksedsts()+360<-ksdxfstk()+44<-ksdxcdb()+3384<-sspuser()+116<-47dc<-qerghRowP()+520<-qerghWalkHashBucket()+520<-qerghIncrHash
> oradebug short_stack
ksedsts()+360<-ksdxfstk()+44<-ksdxcdb()+3384<-sspuser()+116<-47dc<-qerghAggregateRecords()+528<-qeshLoadRowForGBY()+3020<-qerghRow
```

一条SQL的优化

- SQL执行过程中最重要的那个hash join

```
with tmp_calendar as
(select to_char((select to_date(curr_date, 'yyyy-mm-dd') from ttrd_currdate) +
               level - 1,
               'yyyy-mm-dd') as pay_date
 from dual
 connect by level < 366 * 5),
tmp_calendar2 as
(select pay_date,
       to_char(trunc(to_date(pay_date, 'YYYY-MM-DD'), 'MONTH'),
               'YYYY-MM-DD') as year_mon,
       to_char(trunc(to_date(pay_date, 'YYYY-MM-DD'), 'iw'), 'YY
 from tmp_calendar),
tmp_calendar3 as
(select year_mon,
       year_week,
       pay_date,
       case
         when grouping(pay_date) = 0 then
           'D'
         else
           case
             when grouping(year_week) = 0 then
               'W'
             else
               'M'
           end
       end dates_type,
       decode(case
         when grouping(pay_date) = 0 then
           'D'
         else
           case
             when grouping(year_week) = 0 then
               'W'
             else
               'M'
           end
       end,
       rcf.year_mon,
       'W',
       rcf.year_week,
       tmp_nodatelink1 as
(select ..... ----省略若干行
       nvl(lead(tcf.cf_dates, 1) over(partition by tcf.unit_pkgeid,
                                     tcf.currency,
                                     tcf.dates_type order by tcf.cf_dates),
          '2050-12-31') as end_date,
       sum(nvl(tcf.cf_amount, 0)) over(partition by tcf.unit_pkgeid, tcf.currency,
 from (select rcf.unit_pkgeid,
              max(rcf.unit_name) as unit_name,
              rcf.currency,
              case
                when grouping(rcf.pay_date) = 0 then
                  'D'
                else
                  case
                    when grouping(rcf.year_week) = 0 then
                      'W'
                    else
                      'M'
                  end
              end dates_type,
              decode(case
                when grouping(rcf.pay_date) = 0 then
                  'D'
                else
                  case
                    when grouping(rcf.year_week) = 0 then
                      'W'
                    else
                      'M'
                  end
              end,
              'M',
              rcf.year_mon,
              'W',
              rcf.year_week,
```

一条SQL的优化

- SQL执行过程中最重要的那个hash join

```
end_date,  
plus_cf_amount  
  from tmp_nodatelink1 where end_date>cf_dates and to_date(cf_dat  
select ... --此处省略若干行,  
      plus_cf_amount  
  from tmp_nodatelink_vc  
 inner join tmp_calendar3 tc  
    on vc.cf_dates <= tc.cf_dates  
   and vc.end_date > tc.cf_dates  
   and tc.dates_type = vc.dates_type  
 order by unit_name, currency, cf_dates;
```

- 最重要的几点：
 - 理解这里的**dates_type**
 - 一个等值关联
 - 两个非等值关联

一条SQL的优化

• 回顾这个对比

```
SQL> select unit_pkgeid
2   from tmp_nodatelink_0613_2 vc
3   inner join tmp_calendar5 tc
4     on
5     vc.cf_dates <= tc.cf_dates
6     and vc.end_date > tc.cf_dates
7     and tc.dates_type = vc.dates_type
9   order by unit_name, currency, tc.cf_dates;
```

304227 rows selected.

Elapsed: 00:00:04.32

Execution Plan

Plan hash value: 1923048139

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		10746	912K		447 (24)	00:00:06
1	SORT ORDER BY		10746	912K	1112K	447 (24)	00:00:06
* 2	HASH JOIN		10746	912K		227 (45)	00:00:03
3	TABLE ACCESS FULL	TMP_CALENDAR5	2152	27976		5 (0)	00:00:01
4	TABLE ACCESS FULL	TMP_NODATELINK_0613_2	19237	1390K		122 (1)	00:00:02

Predicate Information (identified by operation id):

```
2 - access("TC"."DATES_TYPE"="VC"."DATES_TYPE")
   filter("VC"."CF_DATES"<="TC"."CF_DATES" AND "VC"."END_DATE">"TC"."CF_DATES")
```

Statistics

```
25 recursive calls
0 db block gets
451 consistent gets
0 physical reads
0 redo size
```

```
SQL> select unit_pkgeid
2   from tmp_nodatelink_0613 vc
3   inner join tmp_calendar4 tc
4     on
5     vc.cf_dates <= tc.cf_dates
6     and vc.end_date > tc.cf_dates
7     and
8     tc.dates_type = vc.dates_type
9   order by unit_name, currency, tc.cf_dates;
```

304227 rows selected.

Elapsed: 00:00:01.92

Execution Plan

Plan hash value: 2255175667

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		233K	19M		4924 (1)	00:01:00
1	SORT ORDER BY		233K	19M	24M	4924 (1)	00:01:00
* 2	HASH JOIN		233K	19M		155 (13)	00:00:02
3	TABLE ACCESS FULL	TMP_CALENDAR4	2152	30128		5 (0)	00:00:01
4	TABLE ACCESS FULL	TMP_NODATELINK_0613	20803	1523K		132 (1)	00:00:02

Predicate Information (identified by operation id):


```
2 - access("TC"."DATES_TYPE"="VC"."DATES_TYPE")
   filter("VC"."CF_DATES"<="TC"."CF_DATES" AND "VC"."END_DATE">"TC"."CF_DATES")
```

Statistics

```
8 recursive calls
0 db block gets
481 consistent gets
0 physical reads
```

一条SQL的优化

- Hash join的原理？
 - 分桶
 - join



假设驱动表 记录分布为

D值 1000

W值 100

M值 10

被驱动表记录分布为

D值10000

W值1000

M值100

那么dates_type等值关联后的记录数
是：



$1000 * 10000 + 100 * 1000 + 10 * 100$

一条SQL的优化

• hash的优化


- 如果我们将D/W/M分的更细呢？
 - 比如将D细分为D1、D2、D3...
 - 为什么更细会更好

细分前




等值关联返回结果：
 1000×10000

细分后



等值关联返回结果：
 $200 \times 2000 + 200 \times 2000 + 200 \times 2000 + \dots$



一条SQL的优化

- 真的有那么理想吗
 - 改写的前提是什么
 - 不能改变结果集
 - 带来了什么问题
 - 驱动表可以像前面那样简单的分组（分桶）
 - 被驱动表需要保证落到同样的桶中
 - 意味着被驱动表的 $D1+D2+D3+……>原D$
 - 最理想的情况&最不理想的情况
 - 如果只有等值关联条件，我们还能细分吗？
 - 分组时，应该还与另外的关联条件有关

一条SQL的优化

```
end_date,  
plus_cf_amount  
  from tmp_nodatelink1 where end_date>cf_dates and to_date(cf_dat  
select ... ---此处省略若干行,  
      plus_cf_amount  
  from tmp_nodatelink vc  
 inner join tmp_calendar3 tc  
   on vc.cf_dates <= tc.cf_dates  
   and vc.end_date > tc.cf_dates  
   and tc.dates_type = vc.dates_type  
 order by unit_name, currency, cf_dates;
```

- 最重要的几点：
 - 理解这里的**dates_type**
 - 一个等值关联
 - 两个非等值关联

一条SQL的优化

- 有优化空间的前提
 - 如果只有相等条件，结果集很大
 - 加上另一个非等值连接，
返回的结果集就没有那么大了，
过滤了大部分数据
 - 如果我们在hash join的时候就
做到这一步呢

```
SQL> select count(*)
2   from tmp_nodatelink_0613_2 vc
3   inner join tmp_calendar5 tc
4       on
5       vc.cf_dates <= tc.cf_dates
6       and vc.end_date > tc.cf_dates
7       and tc.dates_type = vc.dates_type;
      COUNT(*)
-----
      304227

SQL> select count(*)
2   from tmp_nodatelink_0613_2 vc
3   inner join tmp_calendar5 tc
4       on
5       and tc.dates_type = vc.dates_type;
      COUNT(*)
-----
      23004227
```

一条SQL的优化

- 实现它（驱动表的改写，很简单）

```
--  
tmp_calendar3 as  
(select  
    year_mon,  
    year_week,  
    pay_date,  
    case when to_date(cf_dates,'yyyy-mm-dd')<sysdate+366 then dates_type||'1' else  
        case when to_date(cf_dates,'yyyy-mm-dd')<sysdate+366*2 then dates_type||'2' else  
            case when to_date(cf_dates,'yyyy-mm-dd')<sysdate+366*3 then dates_type||'3' else  
                case when to_date(cf_dates,'yyyy-mm-dd')<sysdate+366*4 then dates_type||'4' else  
                    case when to_date(cf_dates,'yyyy-mm-dd')<sysdate+366*5 then dates_type||'5' else  
                        dates_type||'6'  
                    end  
                end  
            end  
        end  
    end  
    dates_type,  
    cf_dates
```

- 原with部分不变，外层再套一个select来实现分组

一条SQL的优化

• 被驱动表的改写分析

vc.cf_dates <= tc.cf_dates and vc.end_date > tc.cf_dates and tc.dates_type = vc.dates_type				join和filter
分组前				
cf_dates	dates_type	dates_type	cf_dates	end_date
2016/7/1	D	D	2016/1/1	2018/12/31
2017/7/1	D			
2018/7/1	D			
2019/7/1	D			
2020/7/1	D			
2021/7/1	D			
分组后				
cf_dates	dates_type	dates_type	cf_dates	end_date
2016/7/1	D1	D1	2016/1/1	2018/12/31
2017/7/1	D2	D2	2016/1/1	2018/12/31
2018/7/1	D3	D3	2016/1/1	2018/12/31
2019/7/1	D4			
2020/7/1	D5			
2021/7/1	D6			

- 最好的情况&最差的情况&为何这里可以改
- 原with部分不变，外层再套一个select并使用union来实现

一条SQL的优化

- 实现它（被驱动表的改写，不那么简单）

```
dates_type||'1' as dates_type,
end_date,
plus_cf_amount
from
tmp_nodatelink1 where end_date>cf_dates and to_date(cf_dates,'yyyy-mm-dd')<(sysdate+366)
...
tmp_nodatelink1 where end_date>cf_dates and to_date(end_date,'yyyy-mm-dd')>(sysdate+366) and to_date(cf_dates,'yyyy-mm-dd')<(sysdate+366*2)
...
tmp_nodatelink1 where end_date>cf_dates and to_date(end_date,'yyyy-mm-dd')>sysdate+366*2 and to_date(cf_dates,'yyyy-mm-dd')<sysdate+366*3
...
tmp_nodatelink1 where end_date>cf_dates and to_date(end_date,'yyyy-mm-dd')>sysdate+366*3 and to_date(cf_dates,'yyyy-mm-dd')<sysdate+366*4
...
tmp_nodatelink1 where end_date>cf_dates and to_date(end_date,'yyyy-mm-dd')>sysdate+366*4 and to_date(cf_dates,'yyyy-mm-dd')<sysdate+366*5
...
tmp_nodatelink1 where end_date>cf_dates and to_date(cf_dates,'yyyy-mm-dd')>=sysdate+366*5)
```

一条SQL的优化

- 最后的结果（逻辑读变大，执行时间变短）

```
SQL> SELECT D2.*
2   FROM (SELECT A.*,
3              ROW_NUMBER() over(order by A.UNIT_NAME, A.CURRENCY, A.CF_DATES) AS NUM
4   FROM (select *
5         from vrep_cashflow_total vc
6         where vc.cf_dates_end >= '2017-05-27'
7         and vc.cf_dates <= '2017-06-10'
8         and vc.dates_type = 'D'
9         and exists (select 1
10                  from V_ACCAUTH_USER_NODE_EXT_MAP accauth
11                  where accauth.node_id = VC.UNIT_PKGEID
12                  and accauth.user_id = 1227
13                  and accauth.SUPER_NODE_ID IN ('a', 'b')
14                  and accauth.v_popedom >= 1)) A) D2
15 WHERE D2.NUM > 1
16       AND D2.NUM <= 30;
```

no rows selected

Elapsed: 00:00:04.11

Statistics

```
-----
4   recursive calls
46  db block gets
123947 consistent gets
278  physical reads
```

完

- 欢迎继续关注技术人生系列